



What the HELM !?

A HELM INTRODUCTION OR WHY EVERYONE SHOULD BE USING HELM



[Get Helm](#)

[Blog](#)

[Docs](#)



The package manager for Kubernetes

Helm is the best way to find, share, and use software built for Kubernetes.



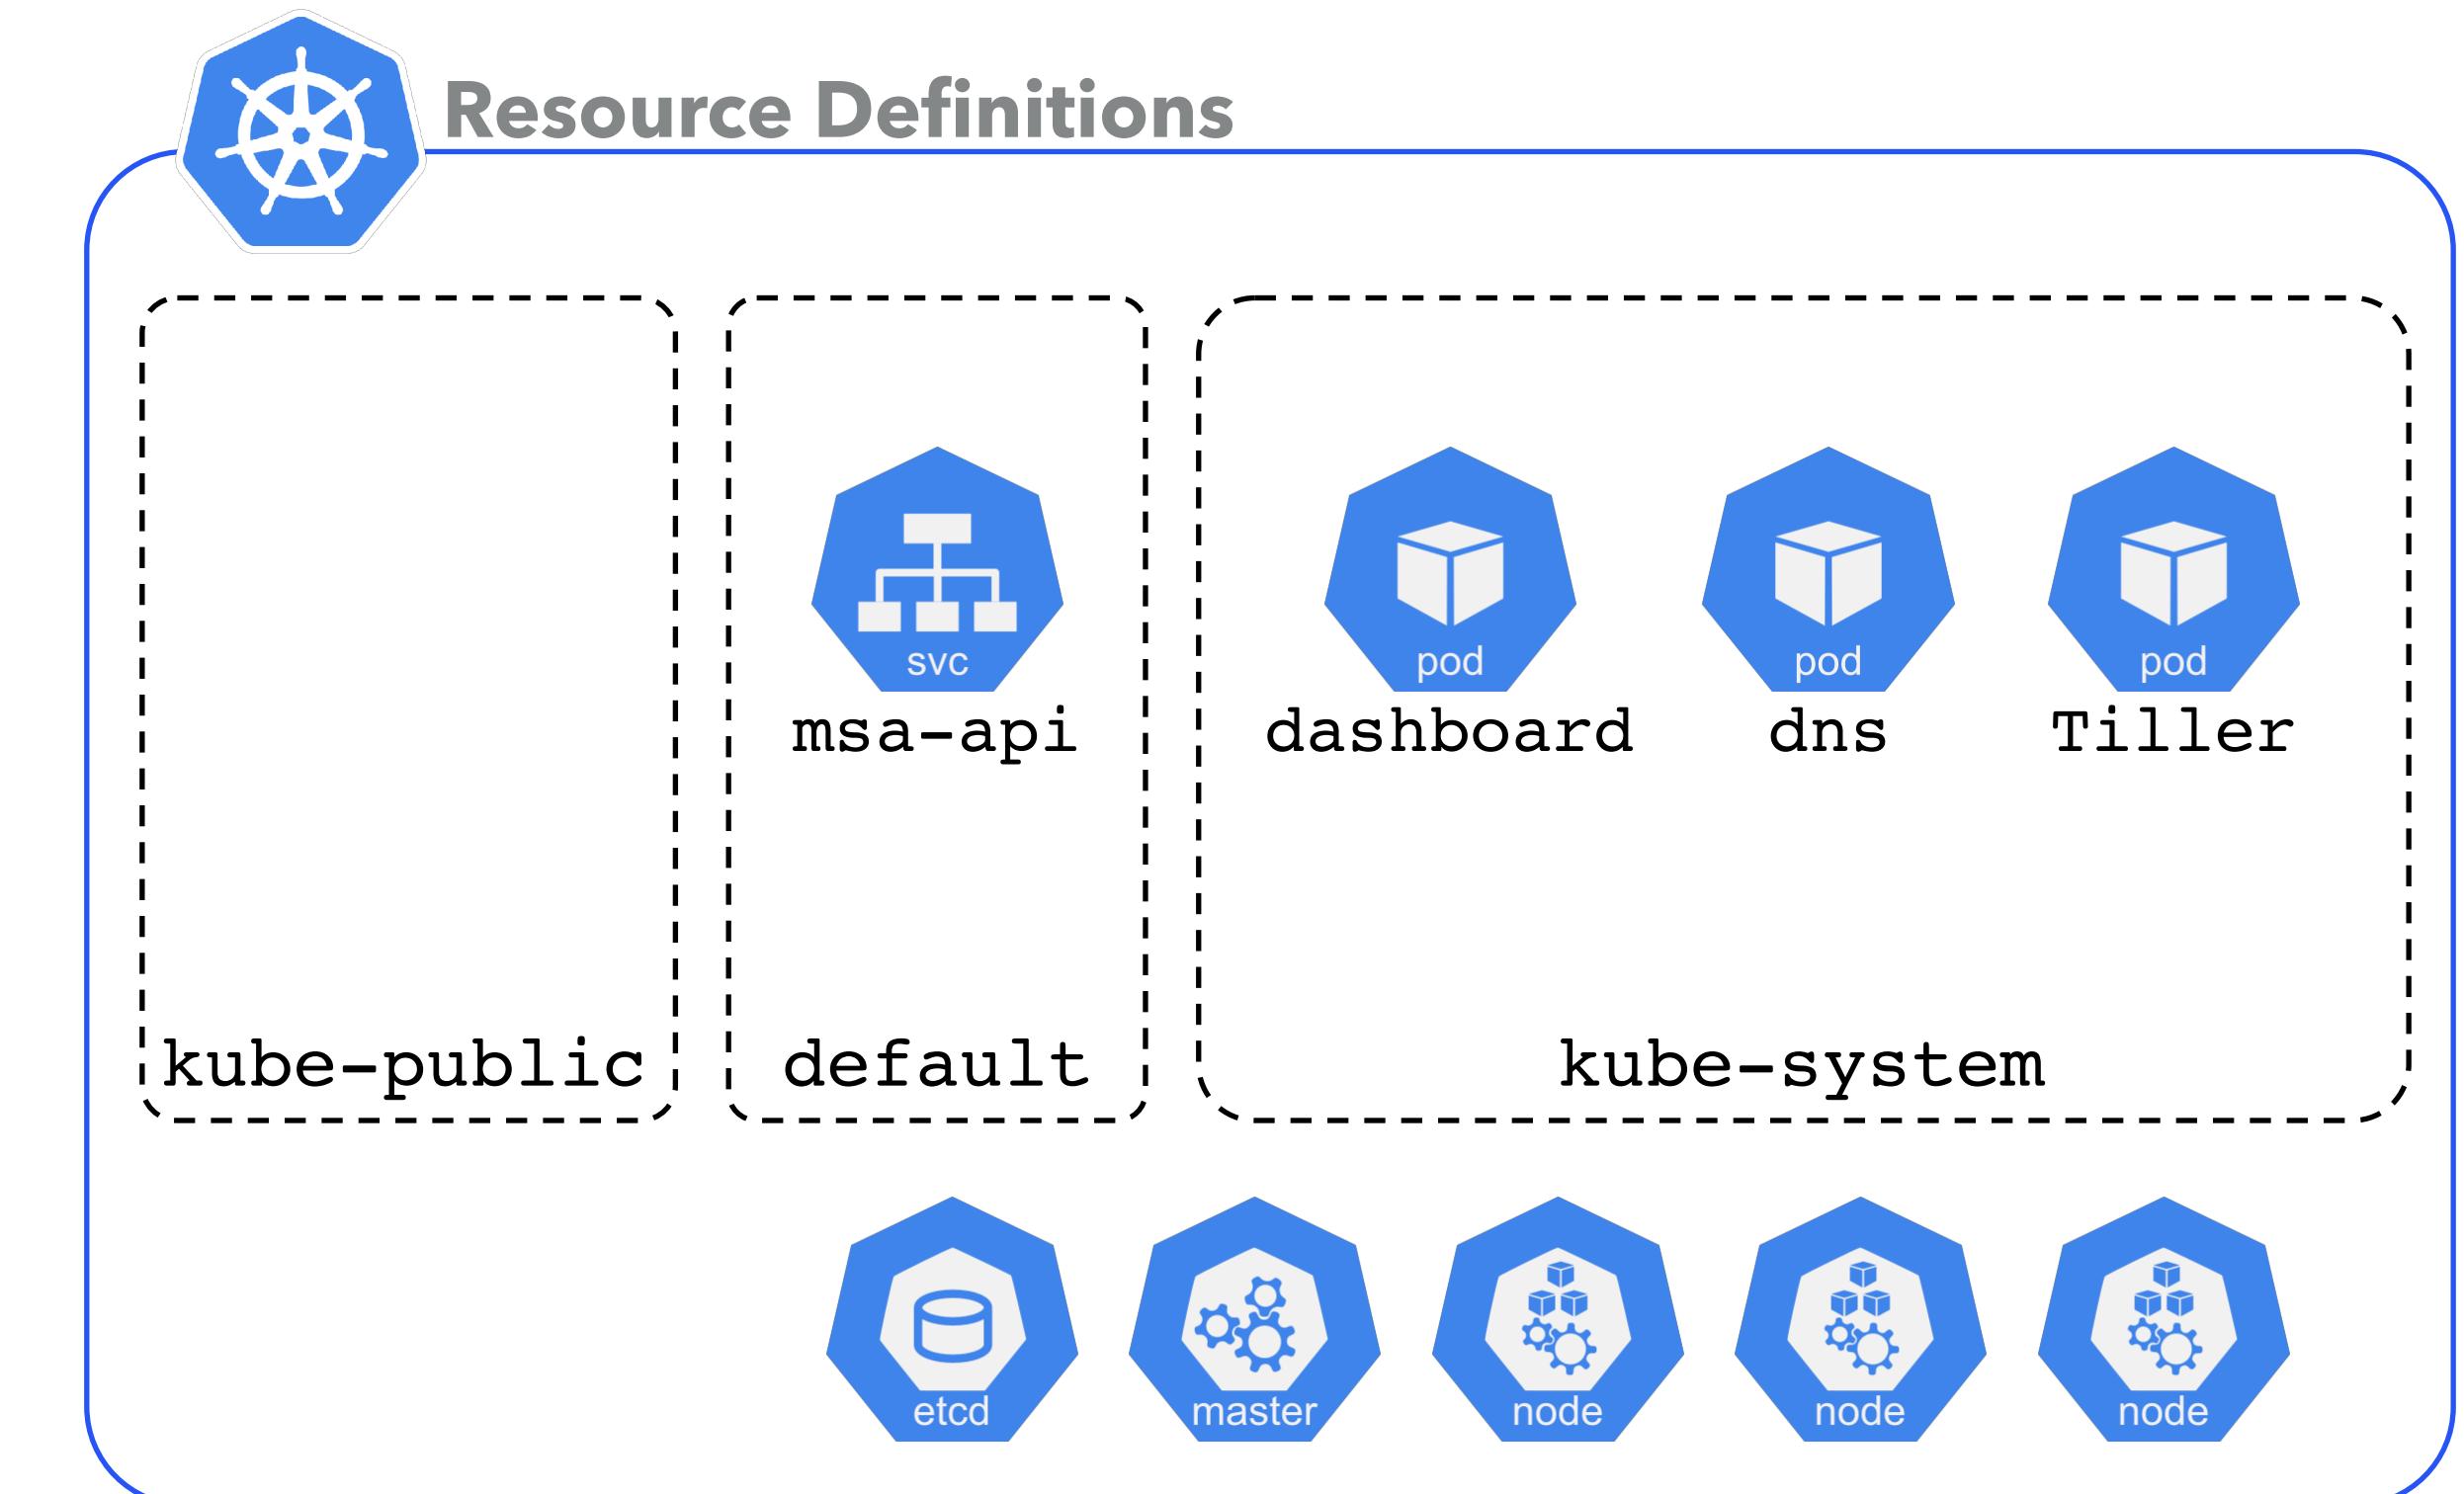
AGENDA

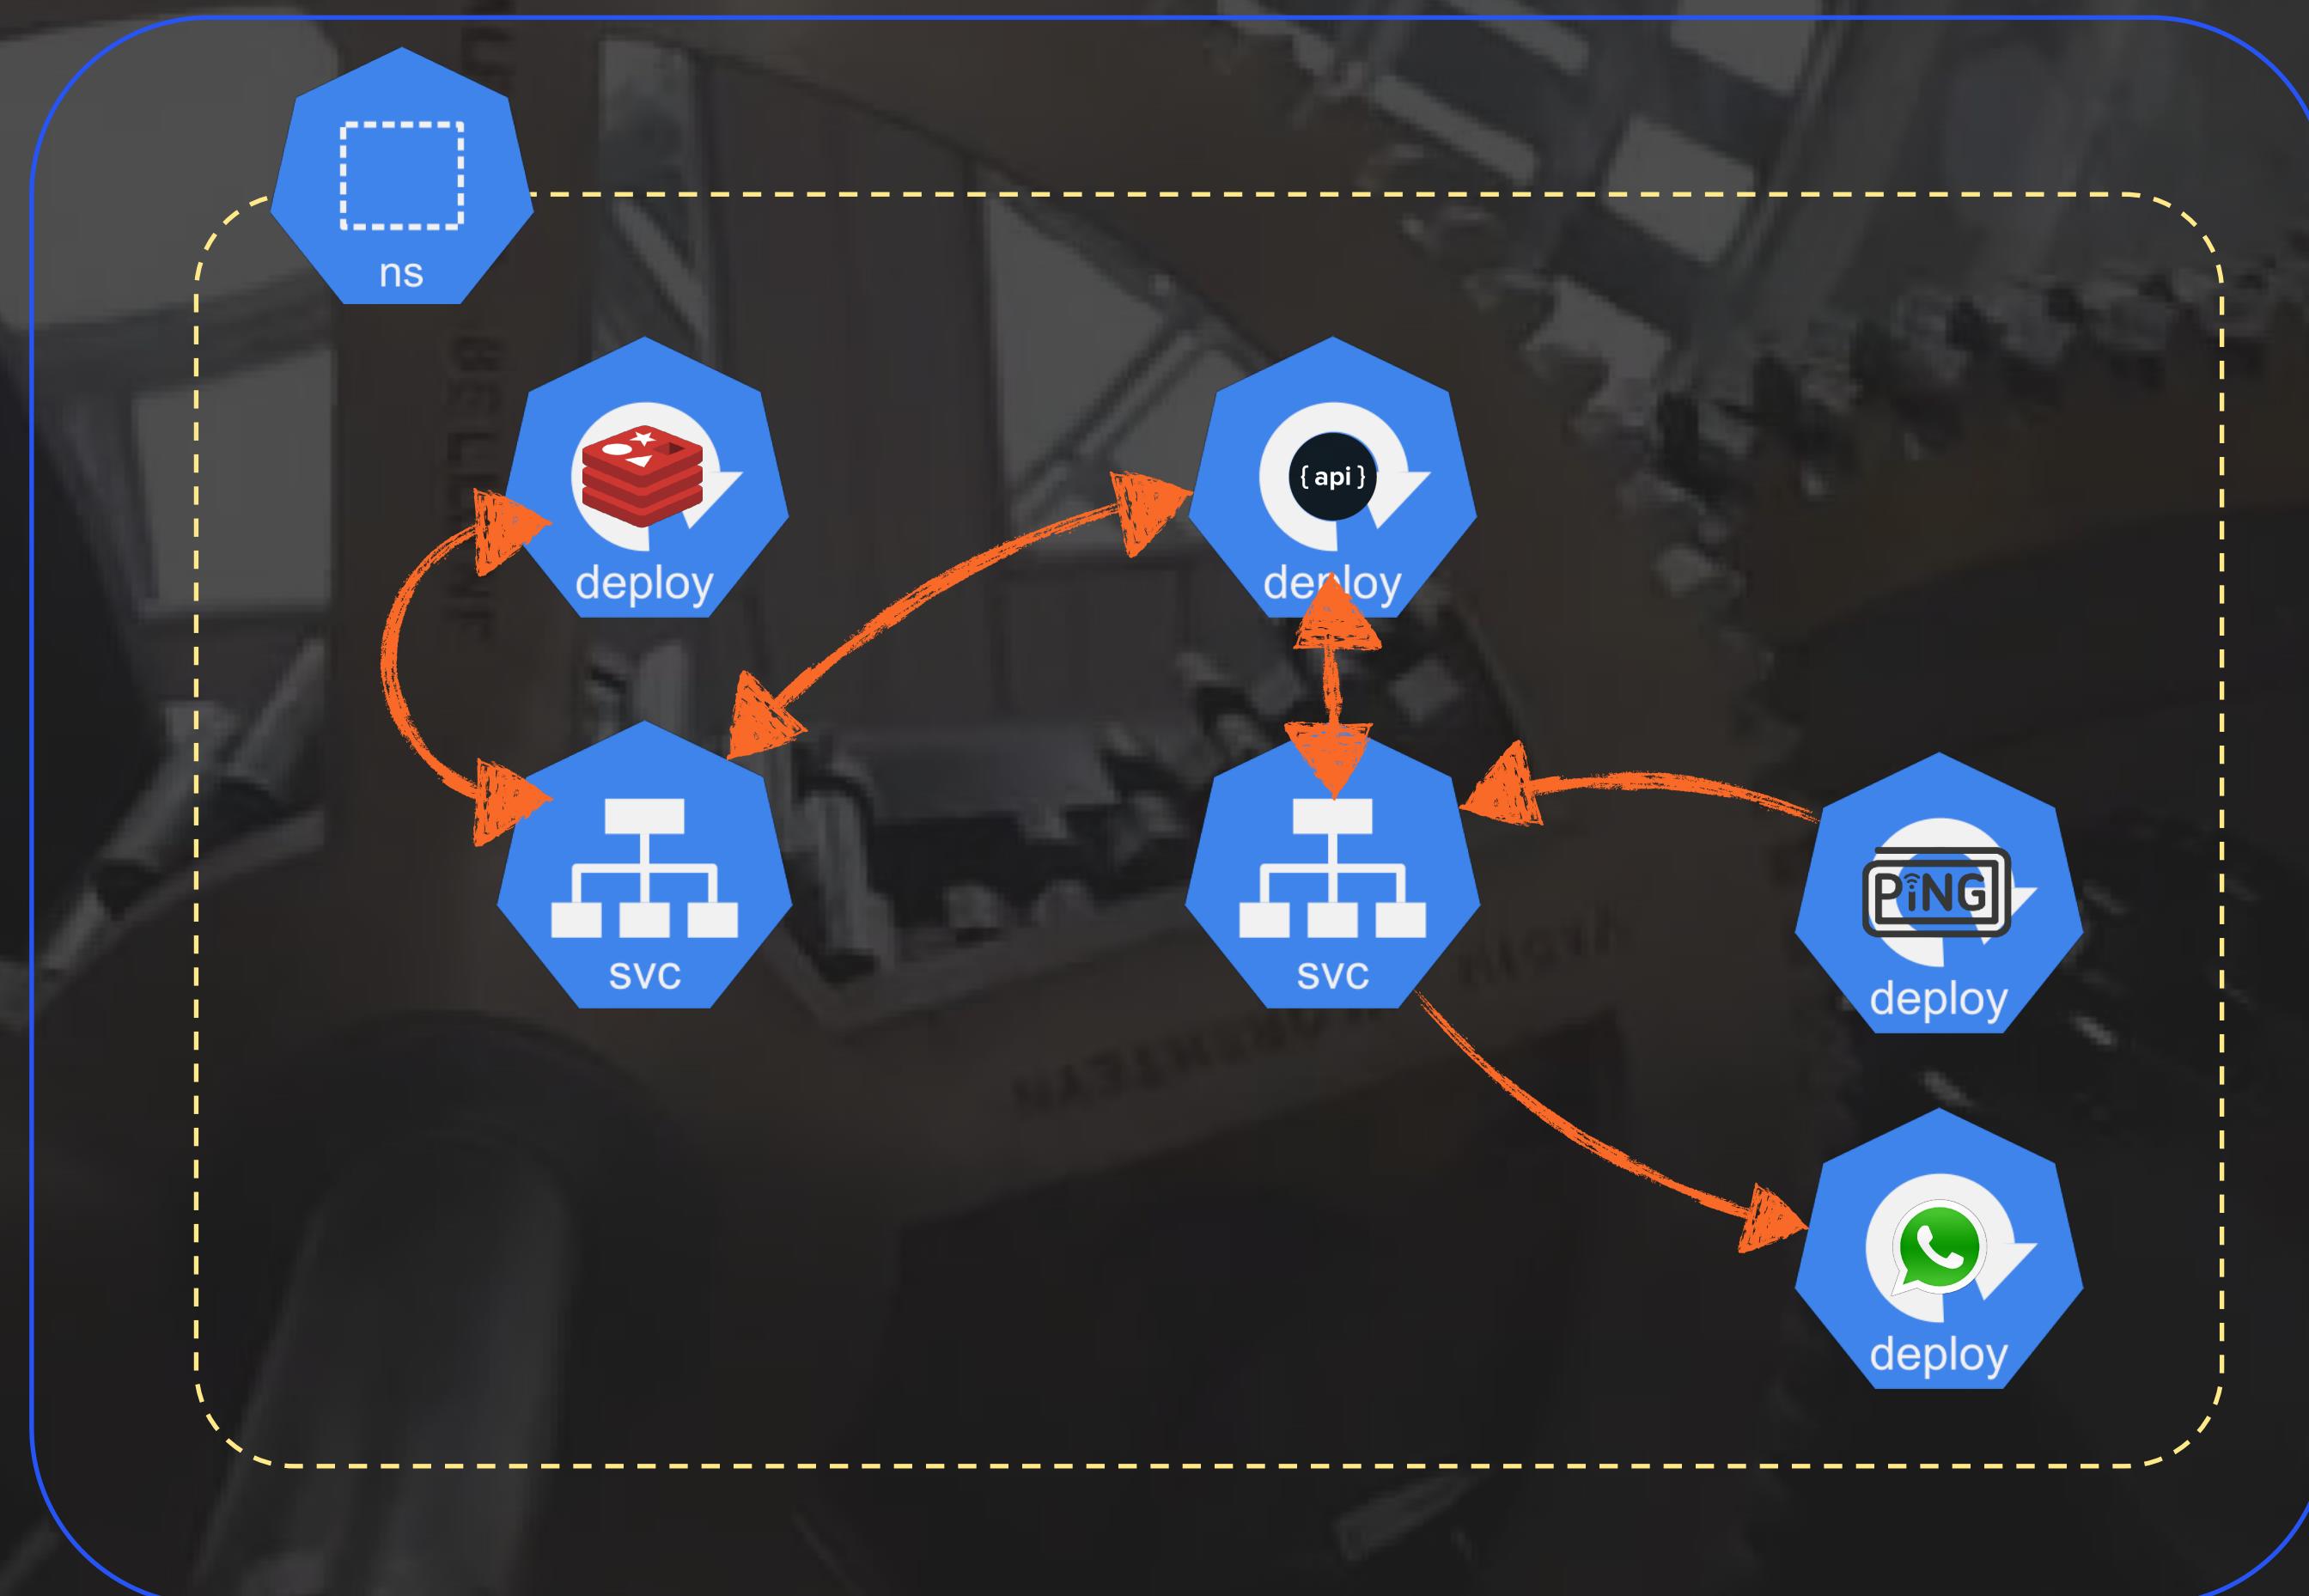
- ▶ Introduction to what helm is
- ▶ Life before HELM
 - ▶ Intro - sample app
 - ▶ demo
- ▶ Life with HELM
- ▶ HELM Architecture
- ▶ Real life examples



ASSUMPTIONS | PERQUISITES

- ▶ You know what kubernetes is:
 - ▶ Namespaces [[link](#)]
 - ▶ Labels & Selectors [[link](#)]
 - ▶ Deployments [[link](#)]
 - ▶ Pods [[link](#)]
 - ▶ Services [[link](#)]





LIFE BEFORE HELM



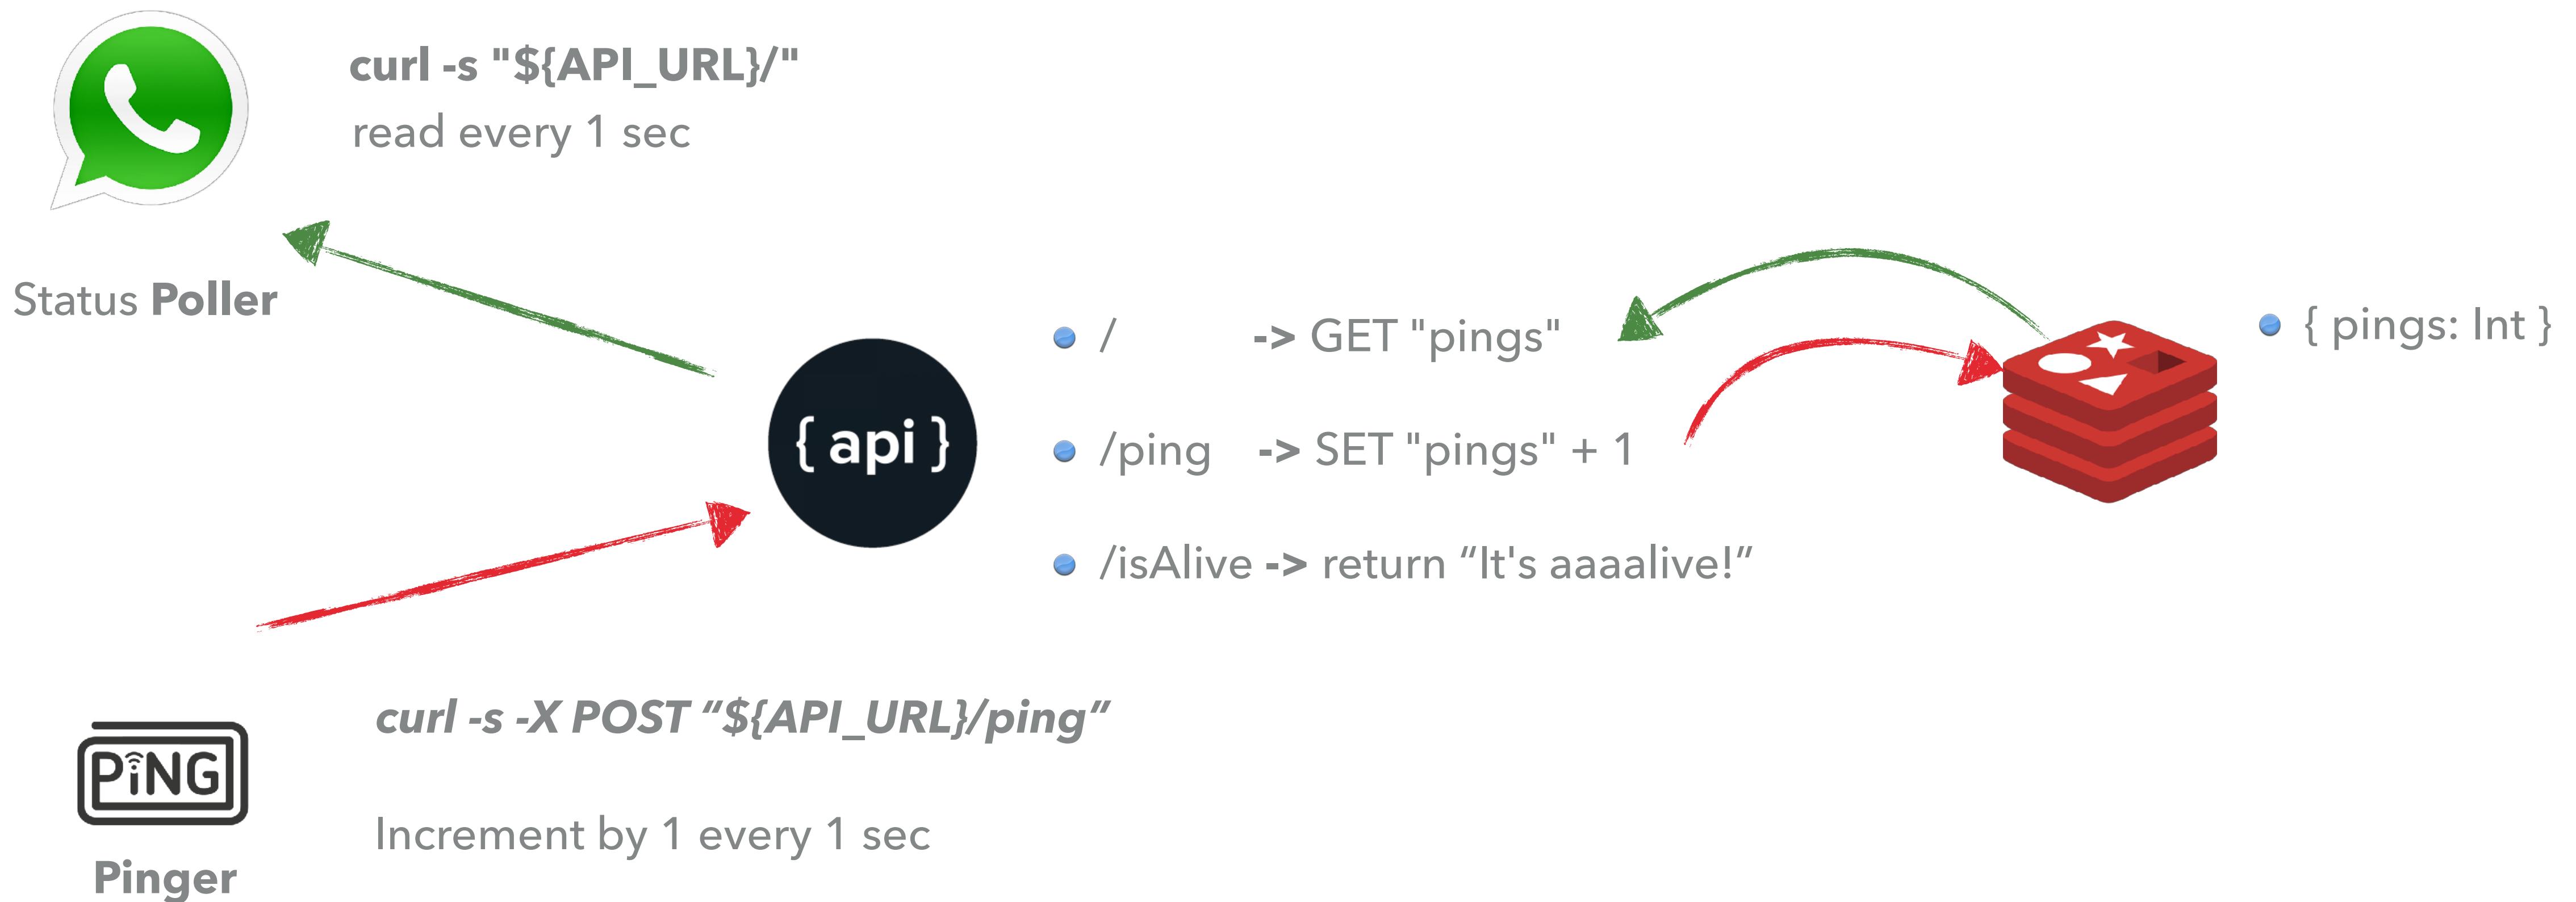


LAND

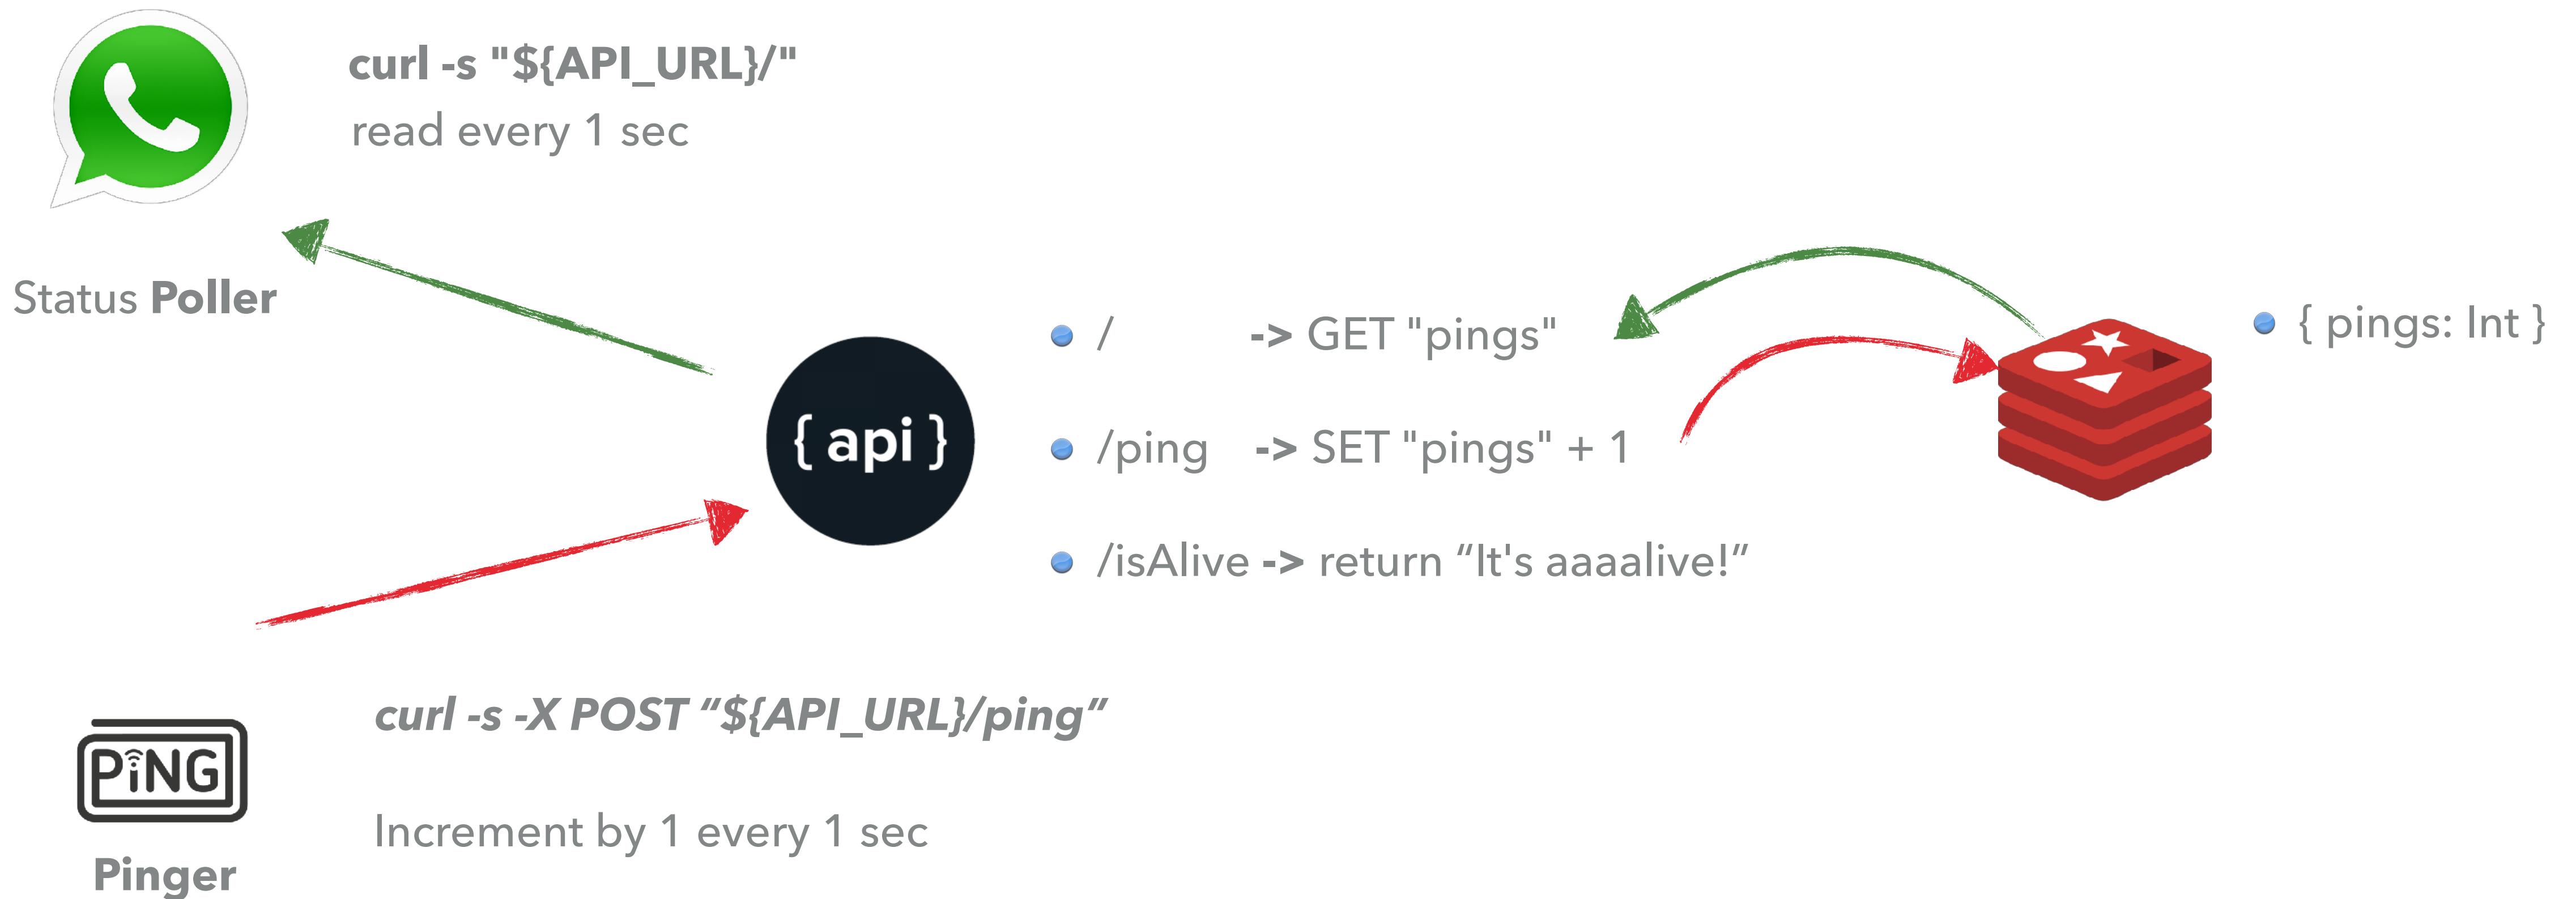
OUR “COMPLICATED” MICRO SERVICE APPLICATION



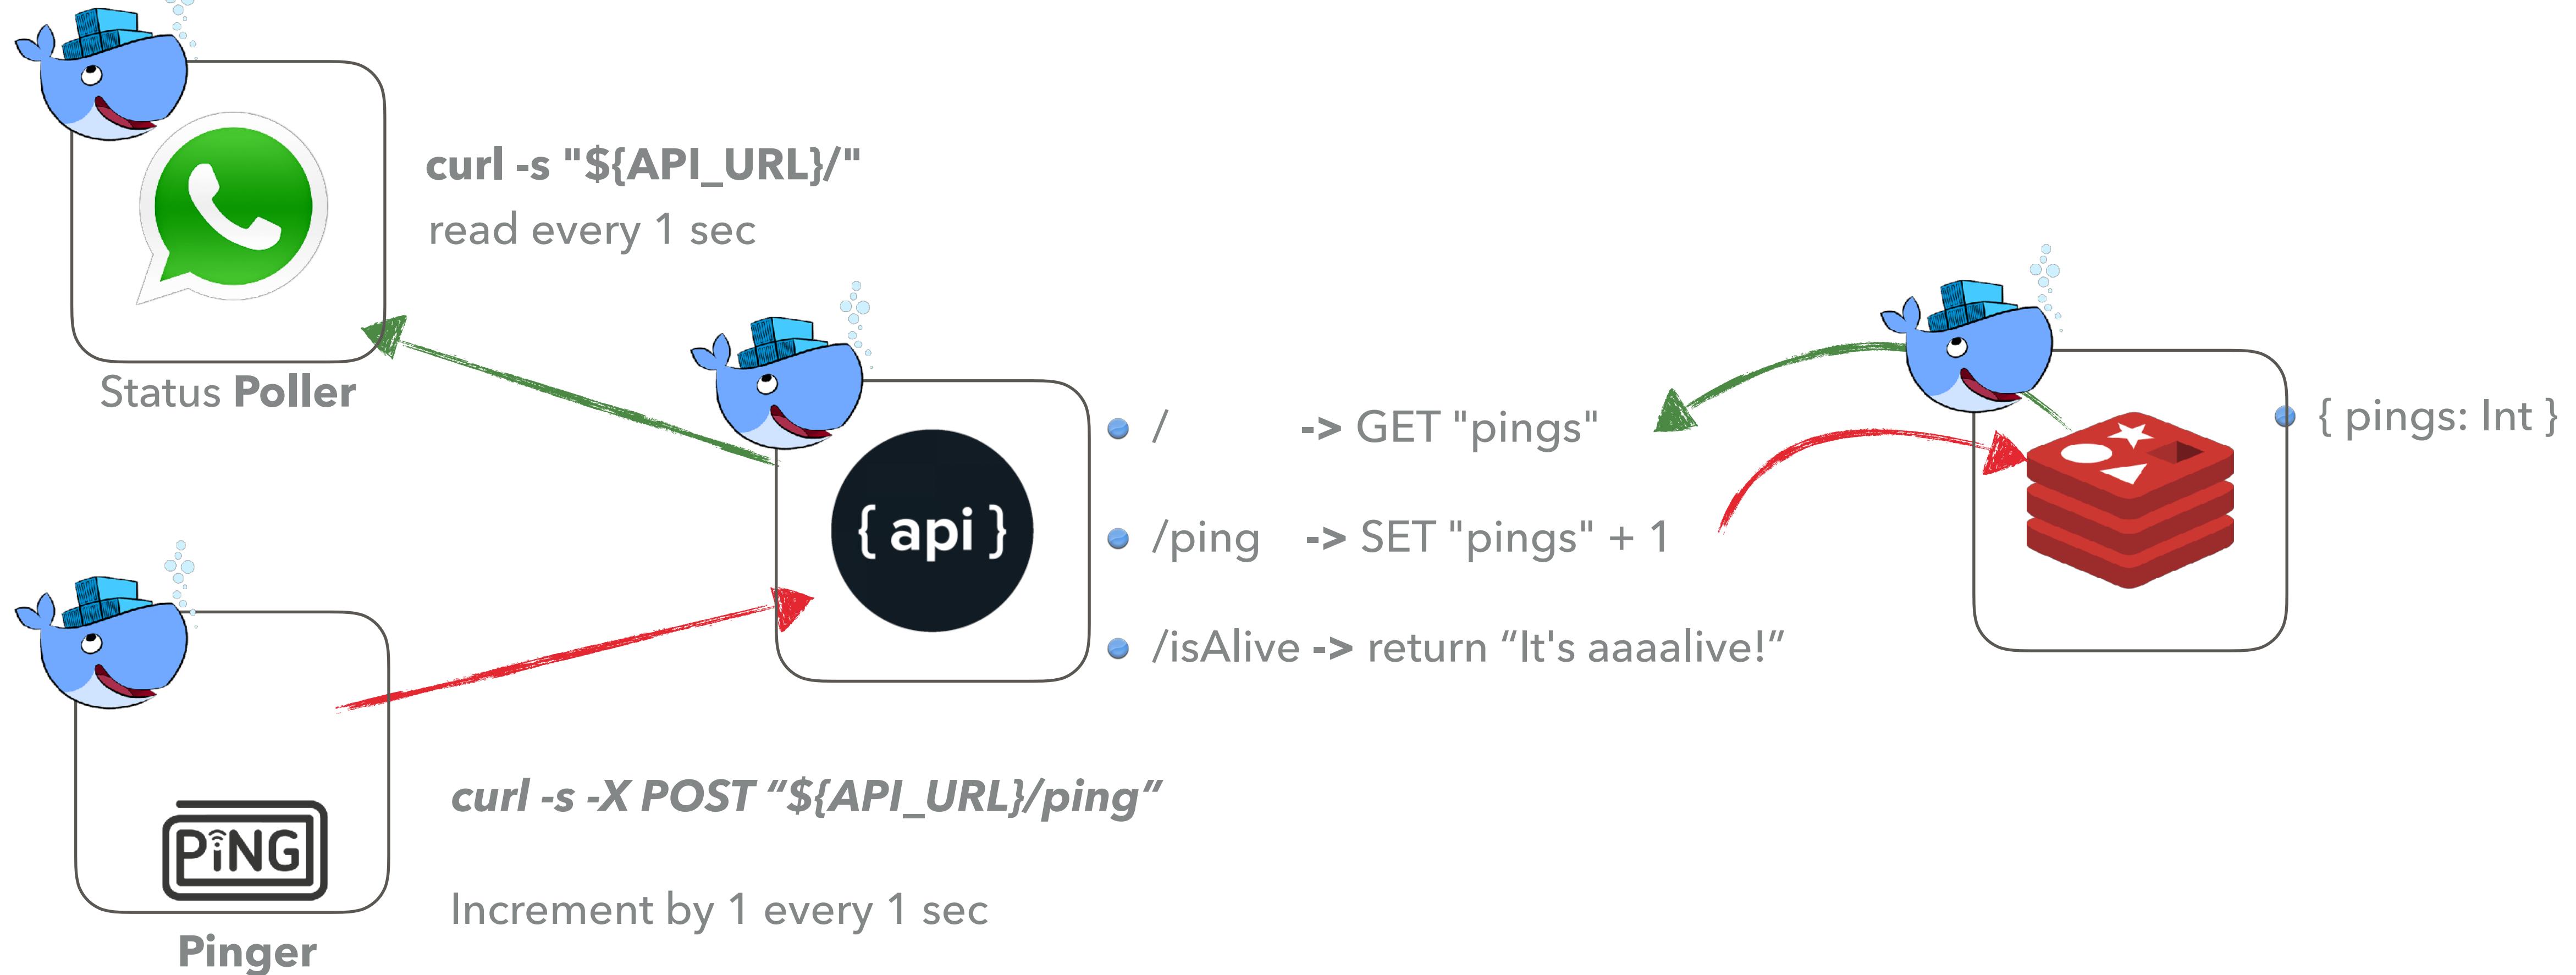
OUR “COMPLICATED” MICRO SERVICE APPLICATION



OUR “COMPLICATED” MICRO SERVICE APPLICATION



OUR “COMPLICATED” MICRO SERVICE APPLICATION - DOCKERIZED



INTRODUCTION TO HELM



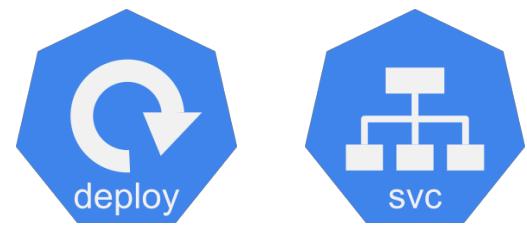
PLANNING OUR DEPLOYMENT

- ▶ Each micro service has a build process which pushes the container to a docker registry.
- ▶ Docker registry is available to the k8s cluster during deployment (with or without authentication depending on the deployment)



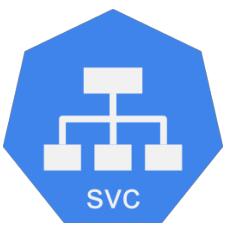
DEPLOYMENT SEQUENCE

DEPLOYMENT SEQUENCE



► ***redis*** deployment + service for redis

DEPLOYMENT SEQUENCE



▶ ***redis*** deployment + service for redis



▶ ***msa-api*** deployment + service for the ***msa-api***

DEPLOYMENT SEQUENCE



▶ ***redis*** deployment + service for redis



▶ ***msa-api*** deployment + service for the ***msa-api***

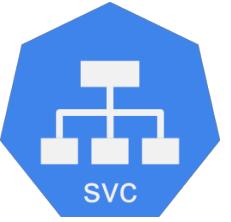


▶ ***msa-pinger*** deployment

DEPLOYMENT SEQUENCE



▶ ***redis*** deployment + service for redis



▶ ***msa-api*** deployment + service for the ***msa-api***

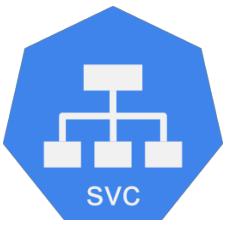


▶ ***msa-pinger*** deployment



▶ ***msa-poller*** deployment

DEPLOYMENT SEQUENCE



▶ ***redis*** deployment + service for redis



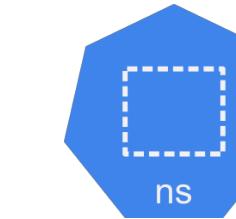
▶ ***msa-api*** deployment + service for the ***msa-api***



▶ ***msa-pinger*** deployment



▶ ***msa-poller*** deployment



▶ ***msa-demo*** namespace

DEPLOYMENT SEQUENCE



▶ ***redis*** deployment + service for redis



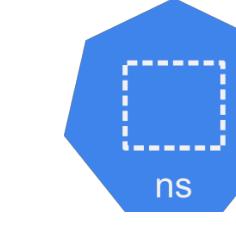
▶ ***msa-api*** deployment + service for the ***msa-api***



▶ ***msa-pinger*** deployment



▶ ***msa-poller*** deployment

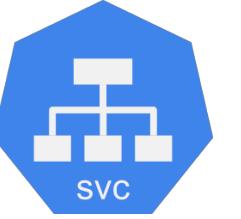
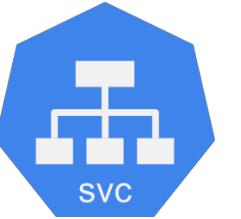


▶ ***msa-demo*** namespace



▶ ***persistent volume*** for redis data

DEPLOYMENT SEQUENCE

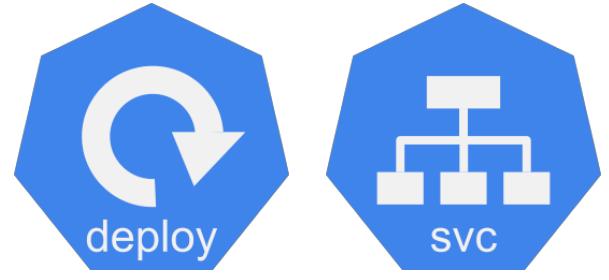
-   ➤ ***redis*** deployment + service for redis
-   ➤ ***msa-api*** deployment + service for the ***msa-api***
-  ➤ ***msa-pinger*** deployment
-  ➤ ***msa-poller*** deployment  ➤ ***msa-demo*** namespace
-  ➤ ***persistent volume*** for redis data
-   ➤ ***configmaps / secrets*** for all components

KUBERNETES “NATIVE” WAY



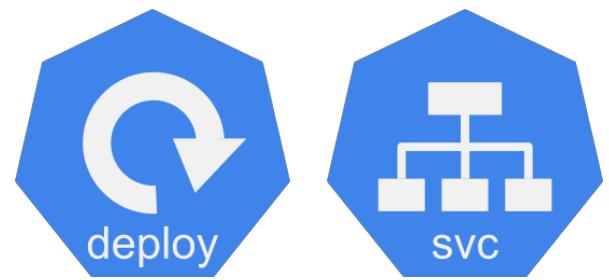
- ▶ Create a namespace:

`kubectl create namespace msa-demo --dry-run -o yaml > msa-demo-ns.yml`



- ▶ Create a deployment + service for redis:

`kubectl run redis --image=redis --port=6379 --expose --dry-run -o yaml > redis.yml`



- ▶ Create a deployment + service for msa-api:

`kubectl run msa-api --image=shelleg/msa-api:config --port=8080 --expose --image-pull-policy=Always --dry-run -o yaml > msa-api.yml`

KUBERNETES “NATIVE” WAY



- ▶ Create a deployment for msa-pinger:

```
kubectl run msa-pinger --image=shelleg/msa-pinger:latest --env="API_URL=msa-api:8080" --env="DEBUG=true" --dry-run -o yaml > msa-pinger.yml
```



- ▶ Create a deployment for msa-poller:

```
kubectl run msa-poller --image=shelleg/msa-poller:latest --env="API_URL=msa-api:8080" --dry-run -o yaml > msa-poller.yml
```

HANDS ON

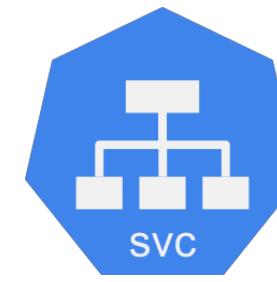


WE NOW HAVE A BUNCH OF FILES DESCRIBING OUR APPLICATION



-
- | -- msa-api.yml
- | -- msa-demo-ns.yml
- | -- msa-pinger.yml
- | -- msa-poller.yml
- | -- redis.yml

0 directories, 5 files

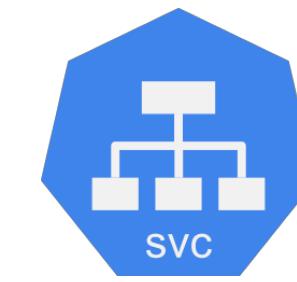


```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  name: msa-api
spec:
  ports:
    - port: 8080
      protocol: TCP
      targetPort: 8080
  selector:
    run: msa-api
status:
  loadBalancer: {}
```



```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    run: msa-api
    name: msa-api
spec:
  replicas: 1
  selector:
    matchLabels:
      run: msa-api
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        run: msa-api
    spec:
      containers:
        - image: shelleg/msa-api:config
          imagePullPolicy: Always
```

SO WHY DO WE NEED HELM ?!



- ▶ We want to better control “what goes where” e.g ports
- ▶ Container image versions ?! (is the only change between releases)
- ▶ Manage Dependencies
- ▶ Update components individually updating things like image versions e.g ***msa-api:latest -> msa-api:config***
- ▶ Write **once** run on any cluster
- ▶ **Reproducible & Shareable resource definition specs - TEMPLATES**



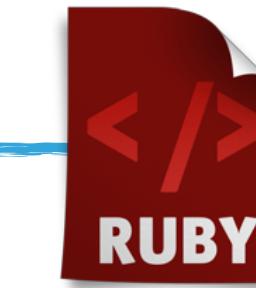
```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  name: msa-api
spec:
  ports:
    - port: {{ variable }}
      protocol: TCP
      targetPort: {{ variable }}
  selector:
    run: msa-api
status:
  loadBalancer: {}
```

TEMPLATES

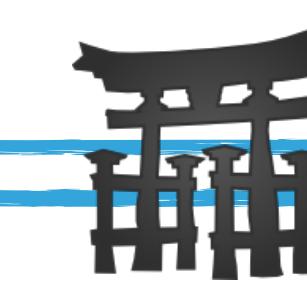


- ▶ Not a foreign concept ... Jinja / erb / jsonnet

```
{% for key, svc in services.items() %}  
upstream {{ key }} {  
  least_conn;  
  {% for server in svc.delegate %}  
    server {{ server }}:{{ svc.ports[0].split(':')[0] }} weight=10  
  max_fails=3 fail_timeout=30s;  
  {% endfor %}  
}  
{% endfor %}
```



```
<% @sudoers_users.each do |user| -%>  
<%= user %> ALL=(ALL) <%= "NOPASSWD:" if @passwordless %>ALL  
<% end -%>  
  
# Members of the sysadmin group may gain root privileges  
%sysadmin      ALL=(ALL) <%= "NOPASSWD:" if @passwordless %>ALL
```



```
// A function that returns an object.  
local Person(name='Alice') = {  
  name: name,  
  welcome: 'Hello ' + name + '!',  
};  
{  
  person1: Person(),  
  person2: Person('Bob'),  
}
```

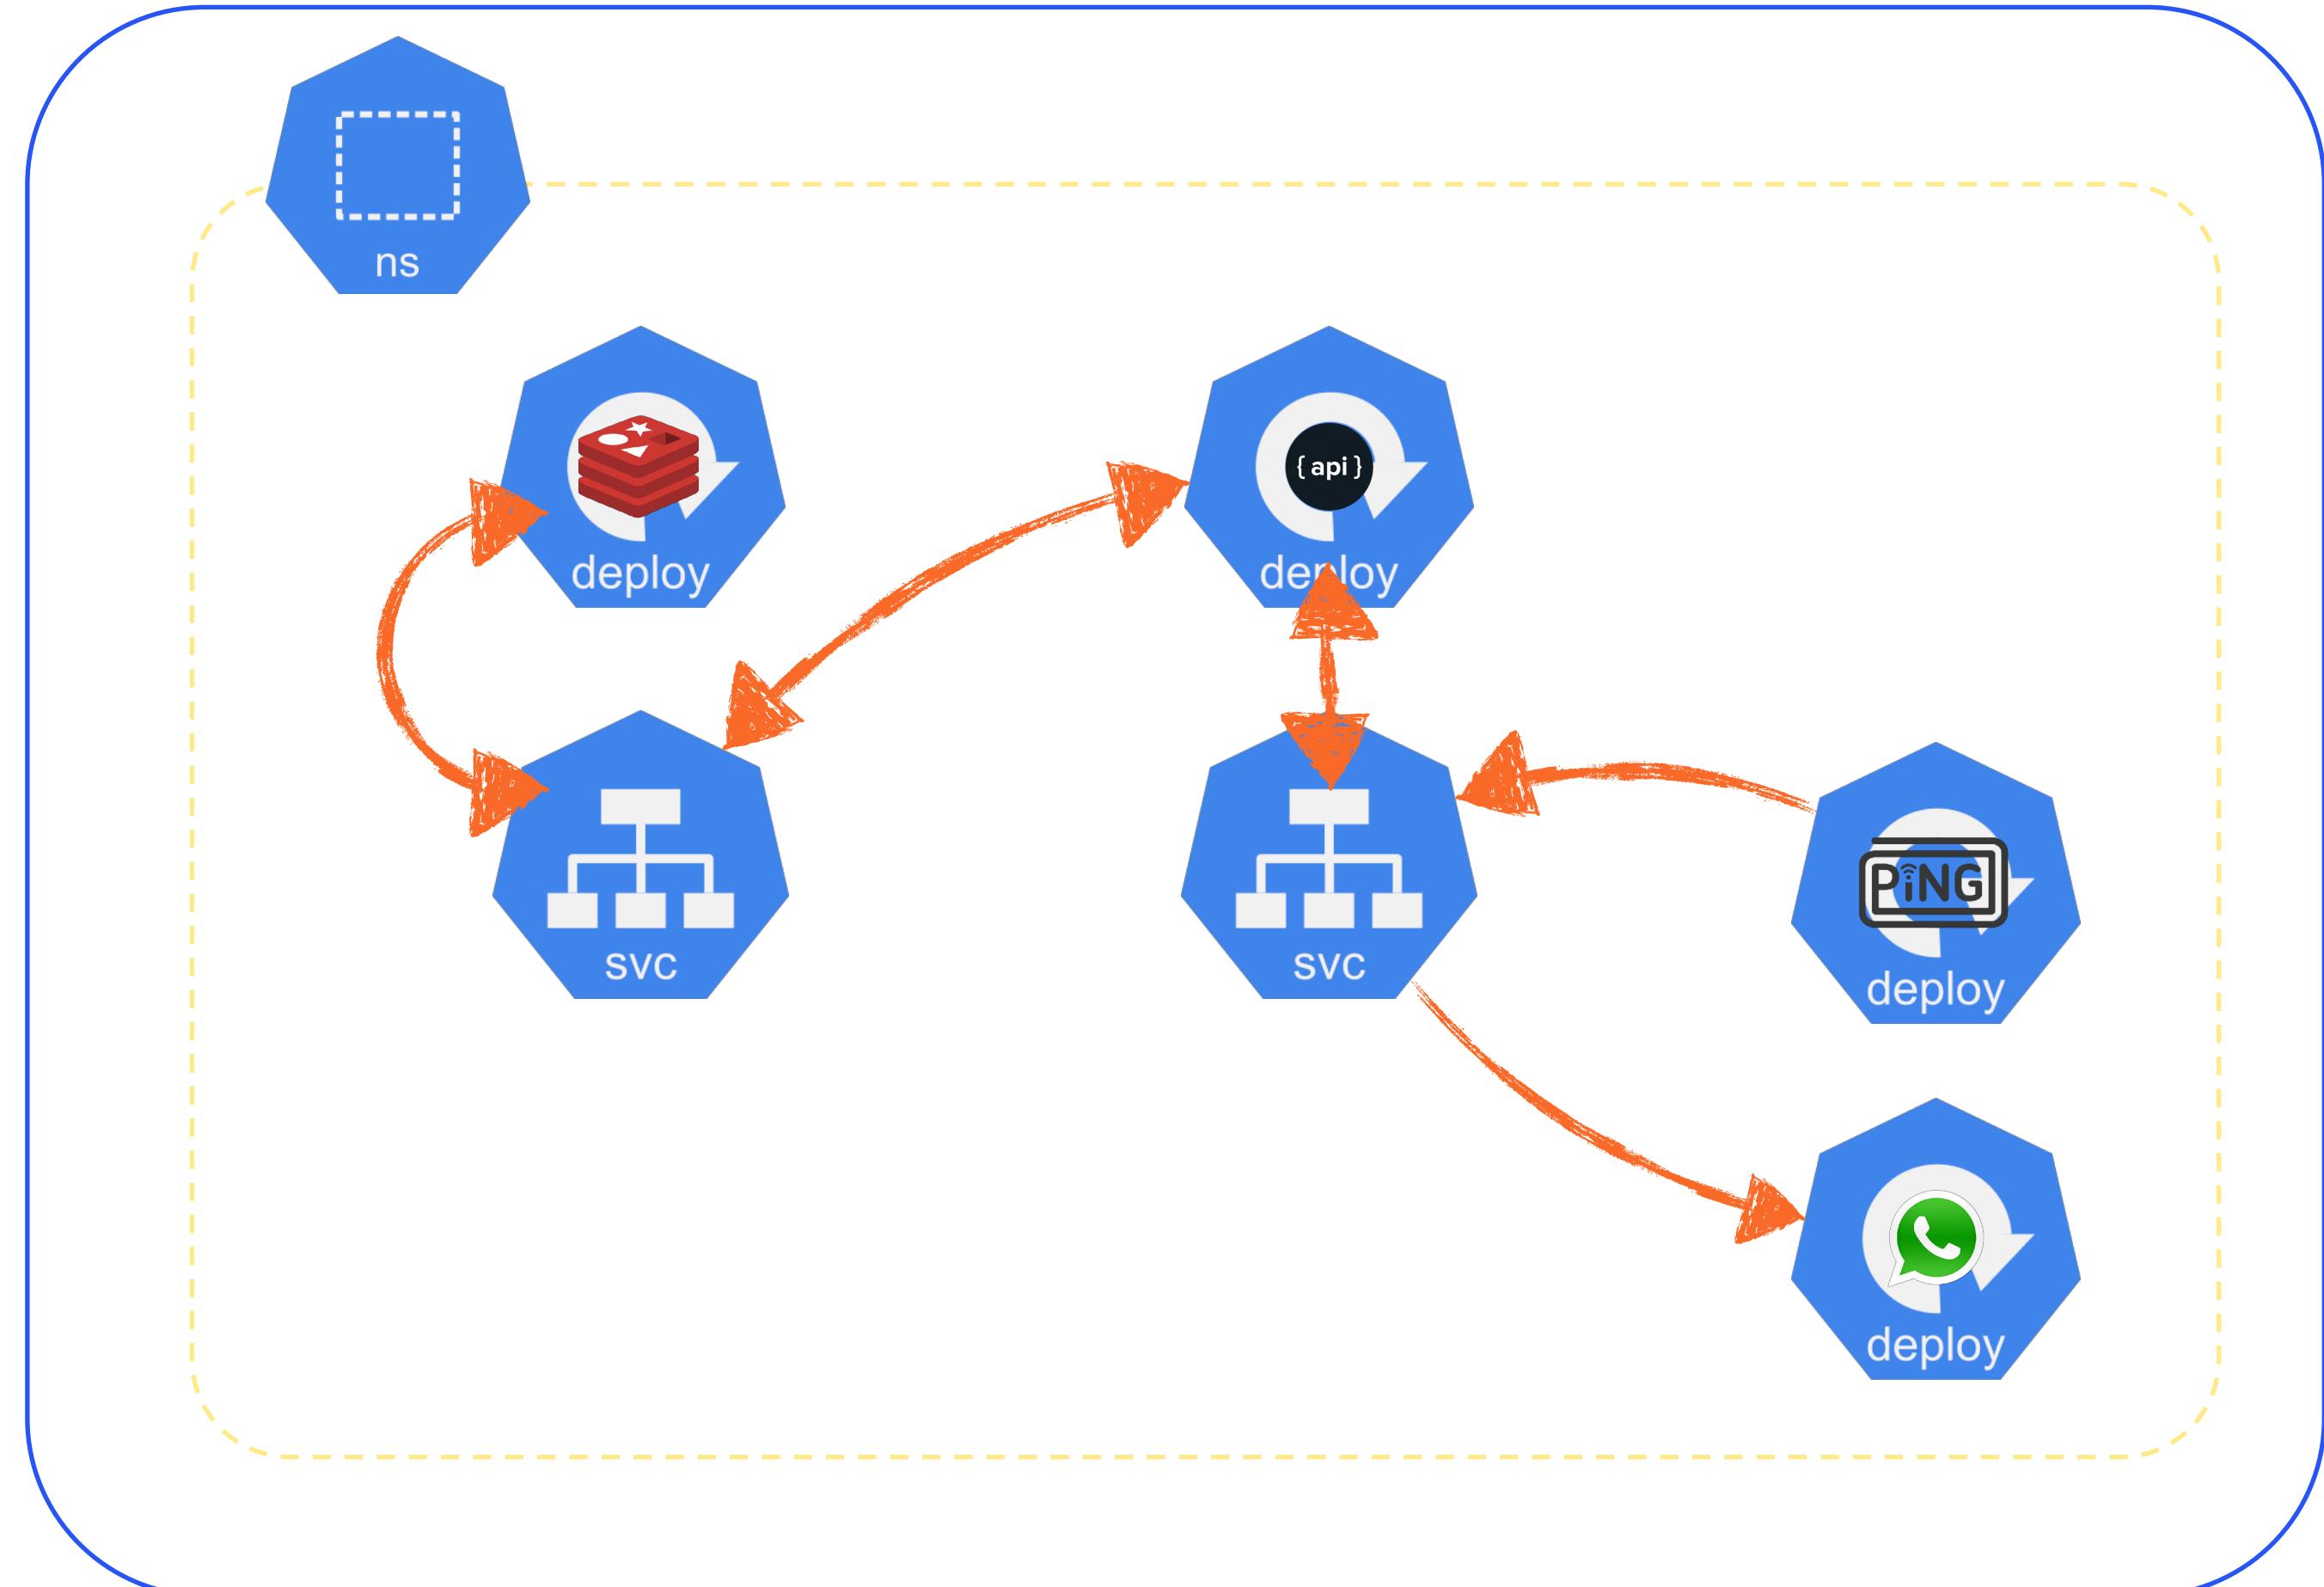


**BEEN THERE
DONE THAT**

SO WHY DO WE NEED HELM ?!

- ▶ Kubernetes native discovery mechanism is focused around the **Service** / deployment
- ▶ Helm Charts are **Application** focused





LIFE WITH HELM



IN A NUT SHELL -> HELM MANAGES CHARTS, RELEASES & VERSIONS



```
.  
|--- Chart.yaml  
|--- charts  
|--- templates  
|   |--- NOTES.txt  
|   |--- _helpers.tpl  
|   |--- deployment.yaml  
|   \--- service.yaml  
`--- values.yaml
```

2 directories, 6 files

- ▶ Describe your deployment + service / any other kubernetes Resource Definition

```
apiVersion: v1  
description: My Cool Chart  
engine: gotpl  
name: myCoolChart  
version: 0.1.0
```

IN A NUT SHELL -> HELM MANAGES CHARTS



- .
- |-- Chart.yaml
- |-- charts
- |-- templates
 - |-- NOTES.txt
 - |-- _helpers.tpl
 - |-- deployment.yaml
 - |-- service.yaml
- values.yaml

2 directories, 6 files

▶ Use **values.yaml**

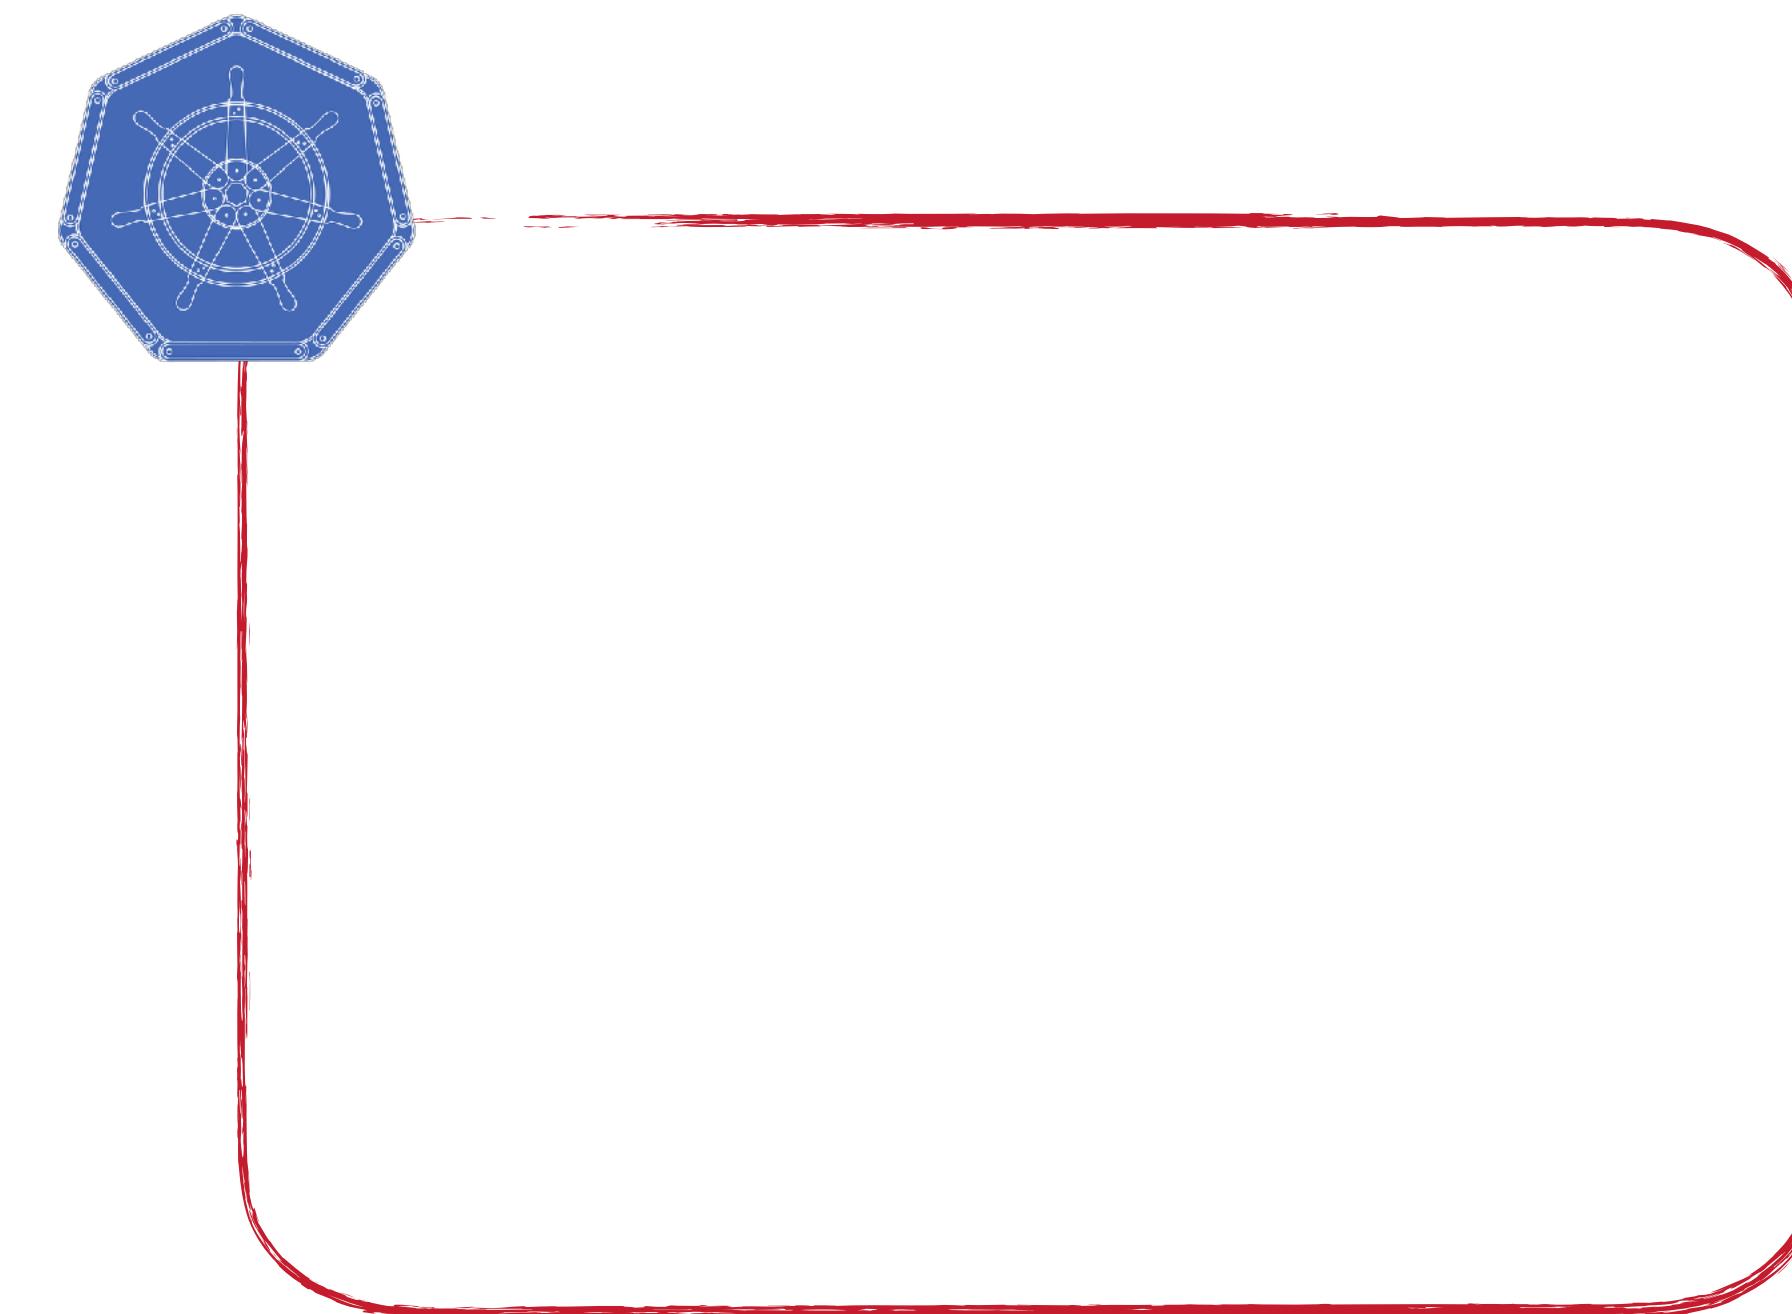
```
image:  
  repository: shelleg/msa-api  
  tag: config  
  pullPolicy: IfNotPresent  
  
service:  
  type: ClusterIP  
  port: 8080  
...
```

▶ A rich templating system for making charts generic yet highly customisable

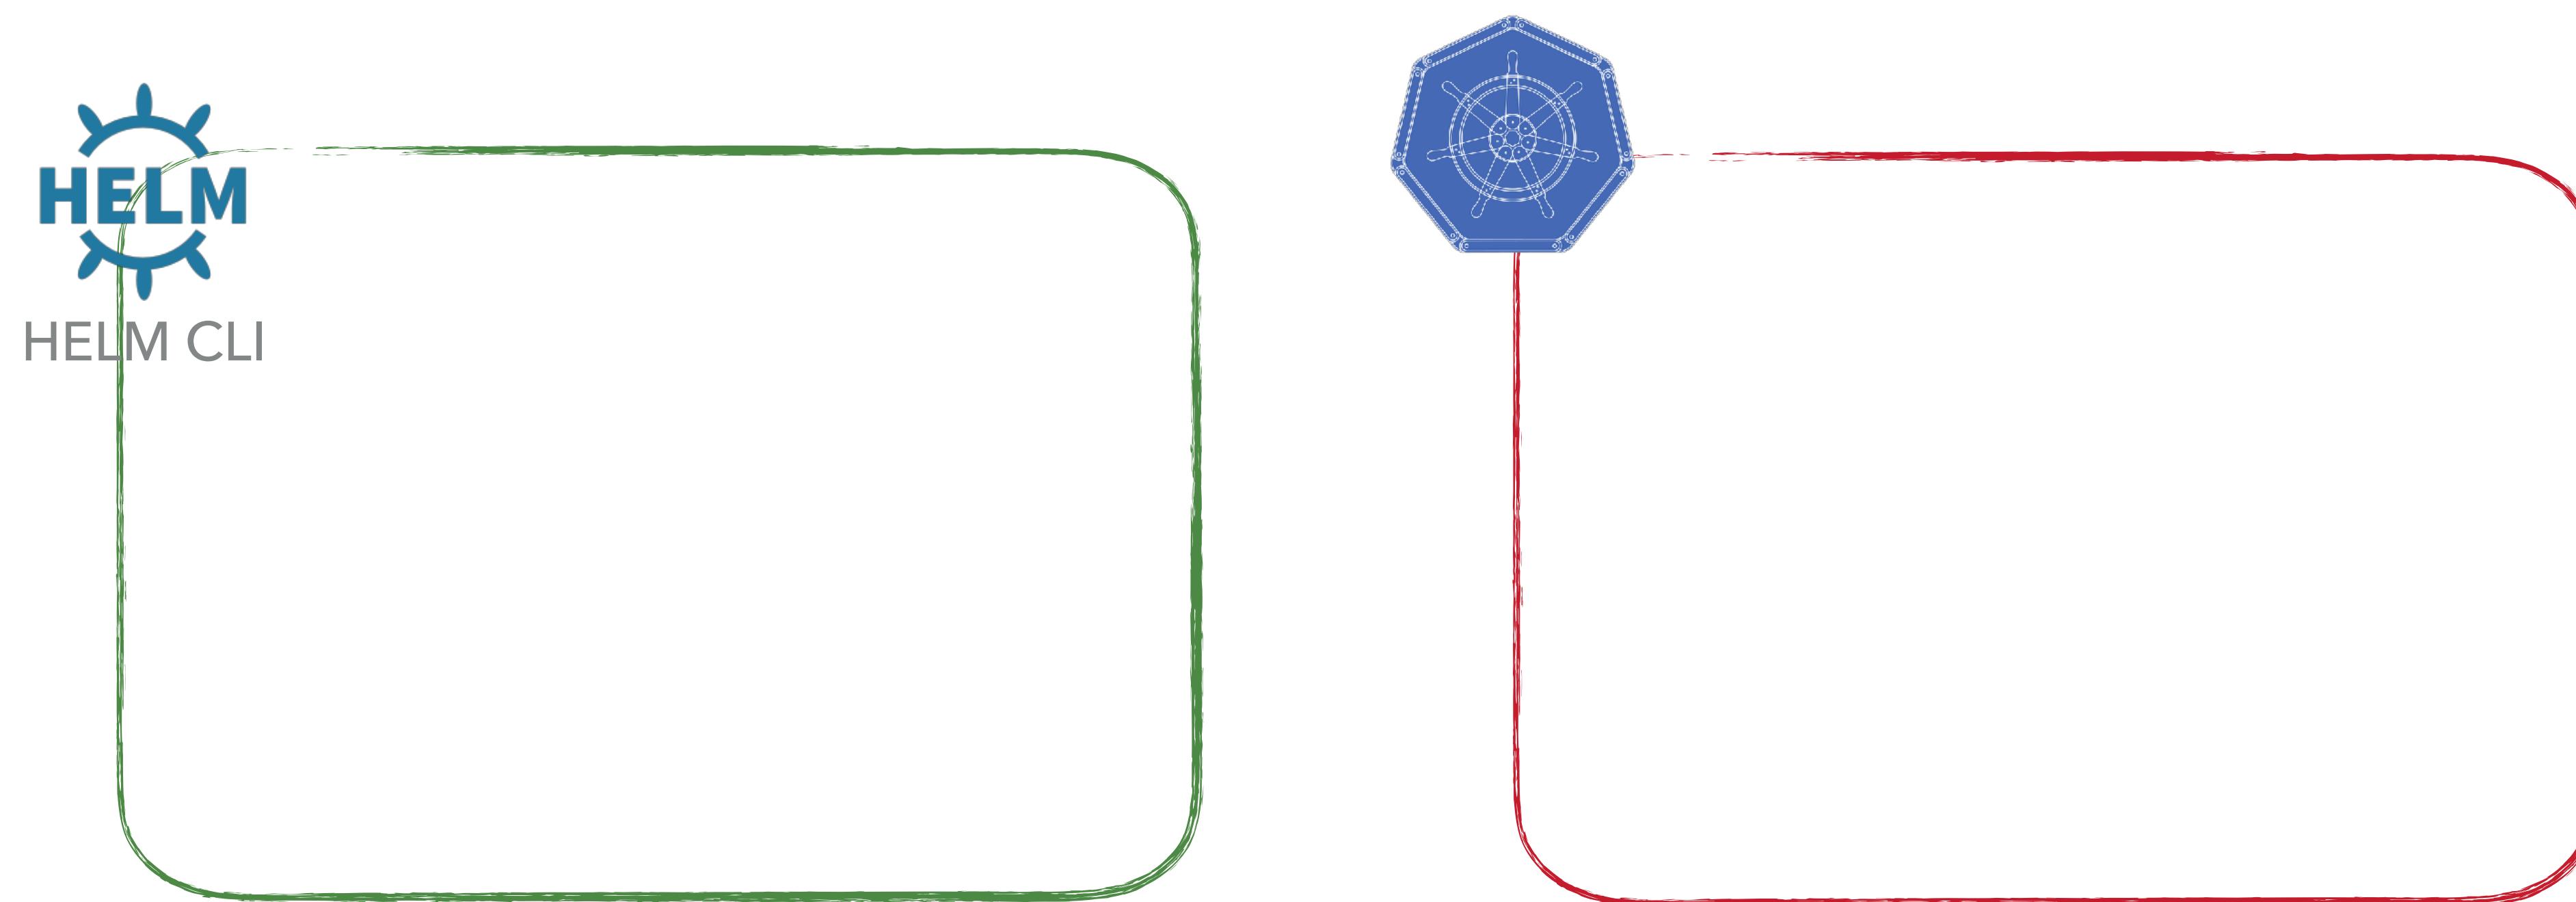
HOW DOES HELM WORK?



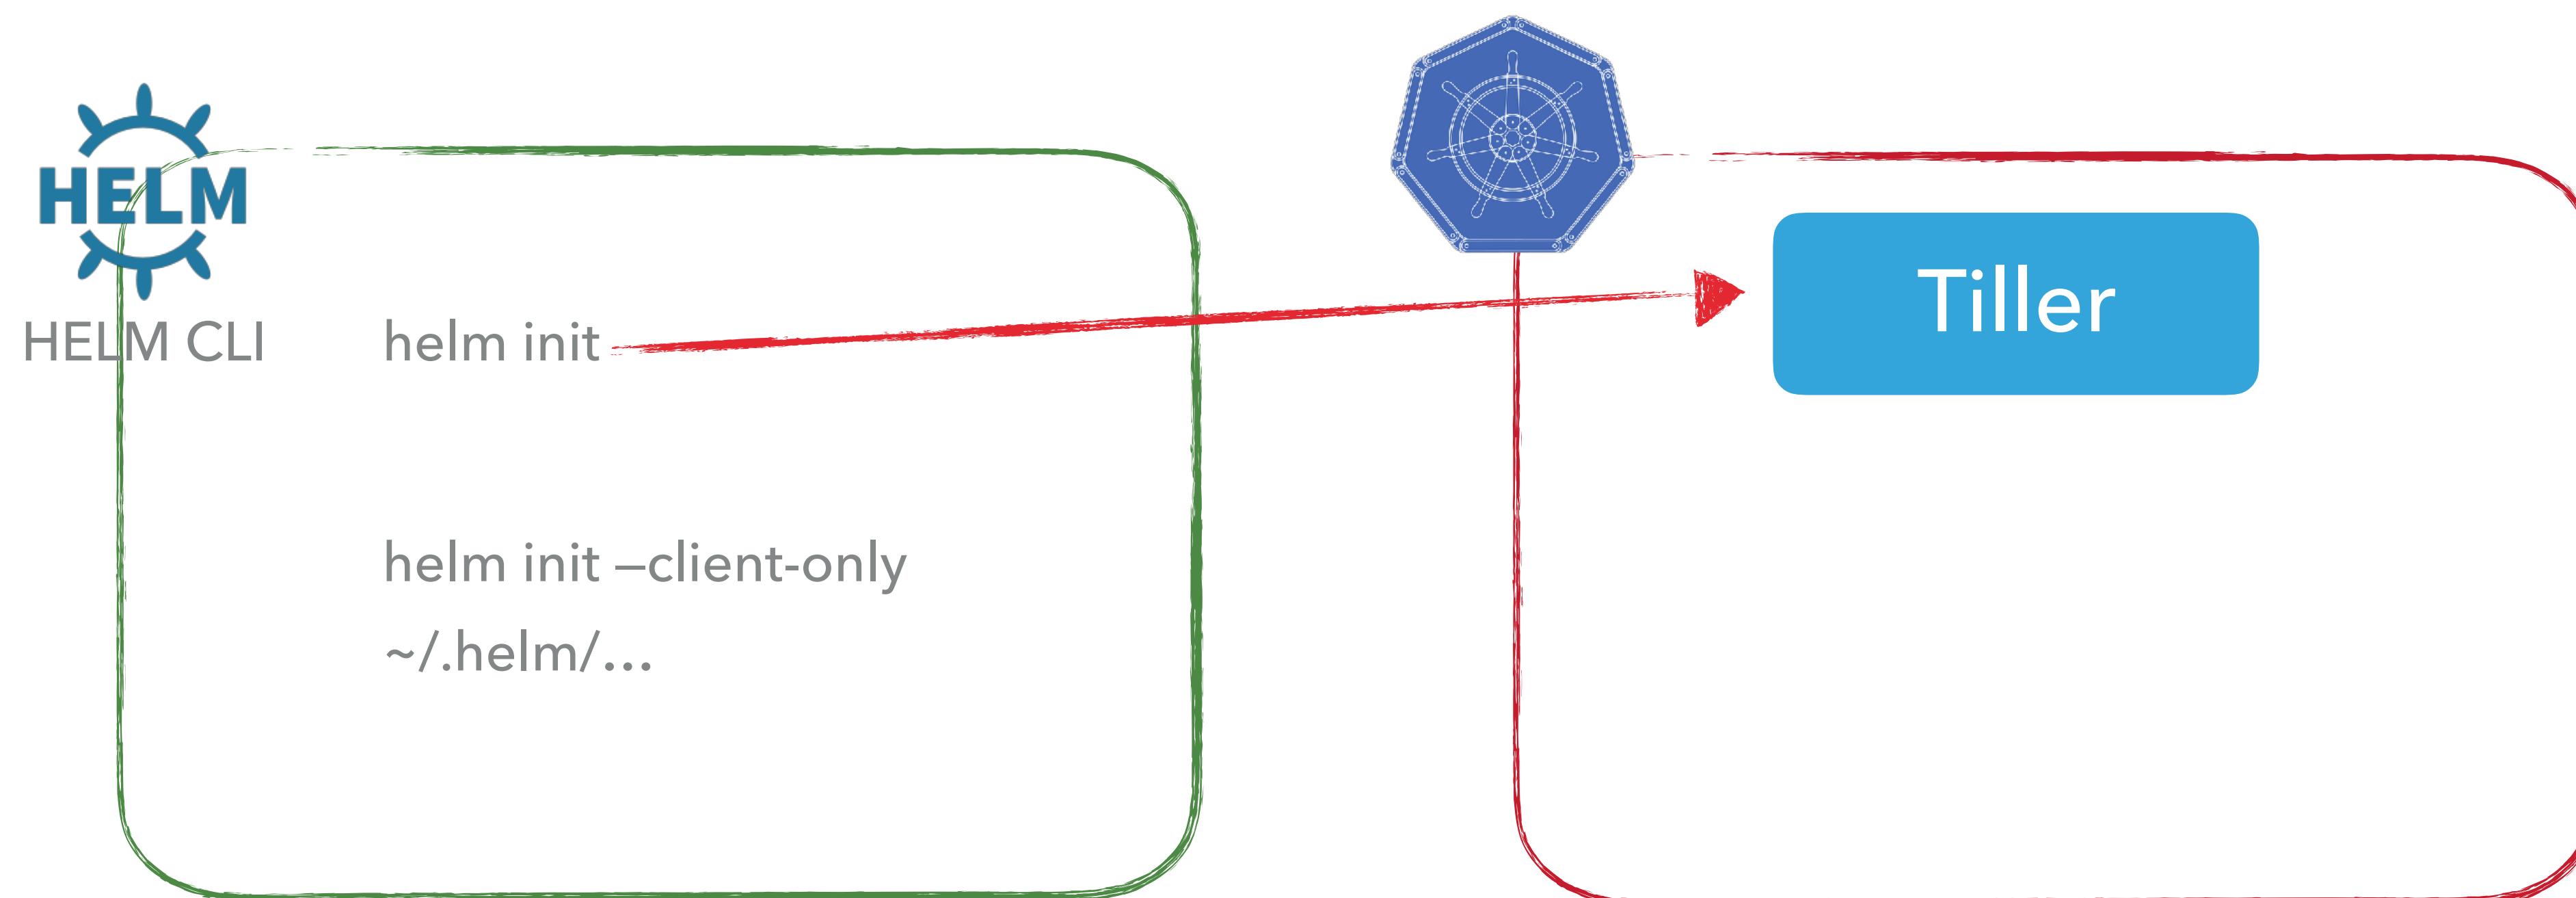
HELM ARCHITECTURE - A KUBERNETES CLUSTER



HELM ARCHITECTURE - A HELMSMAN WORKSTATION



HELM ARCHITECTURE - HELM INIT



HELM CLIENT



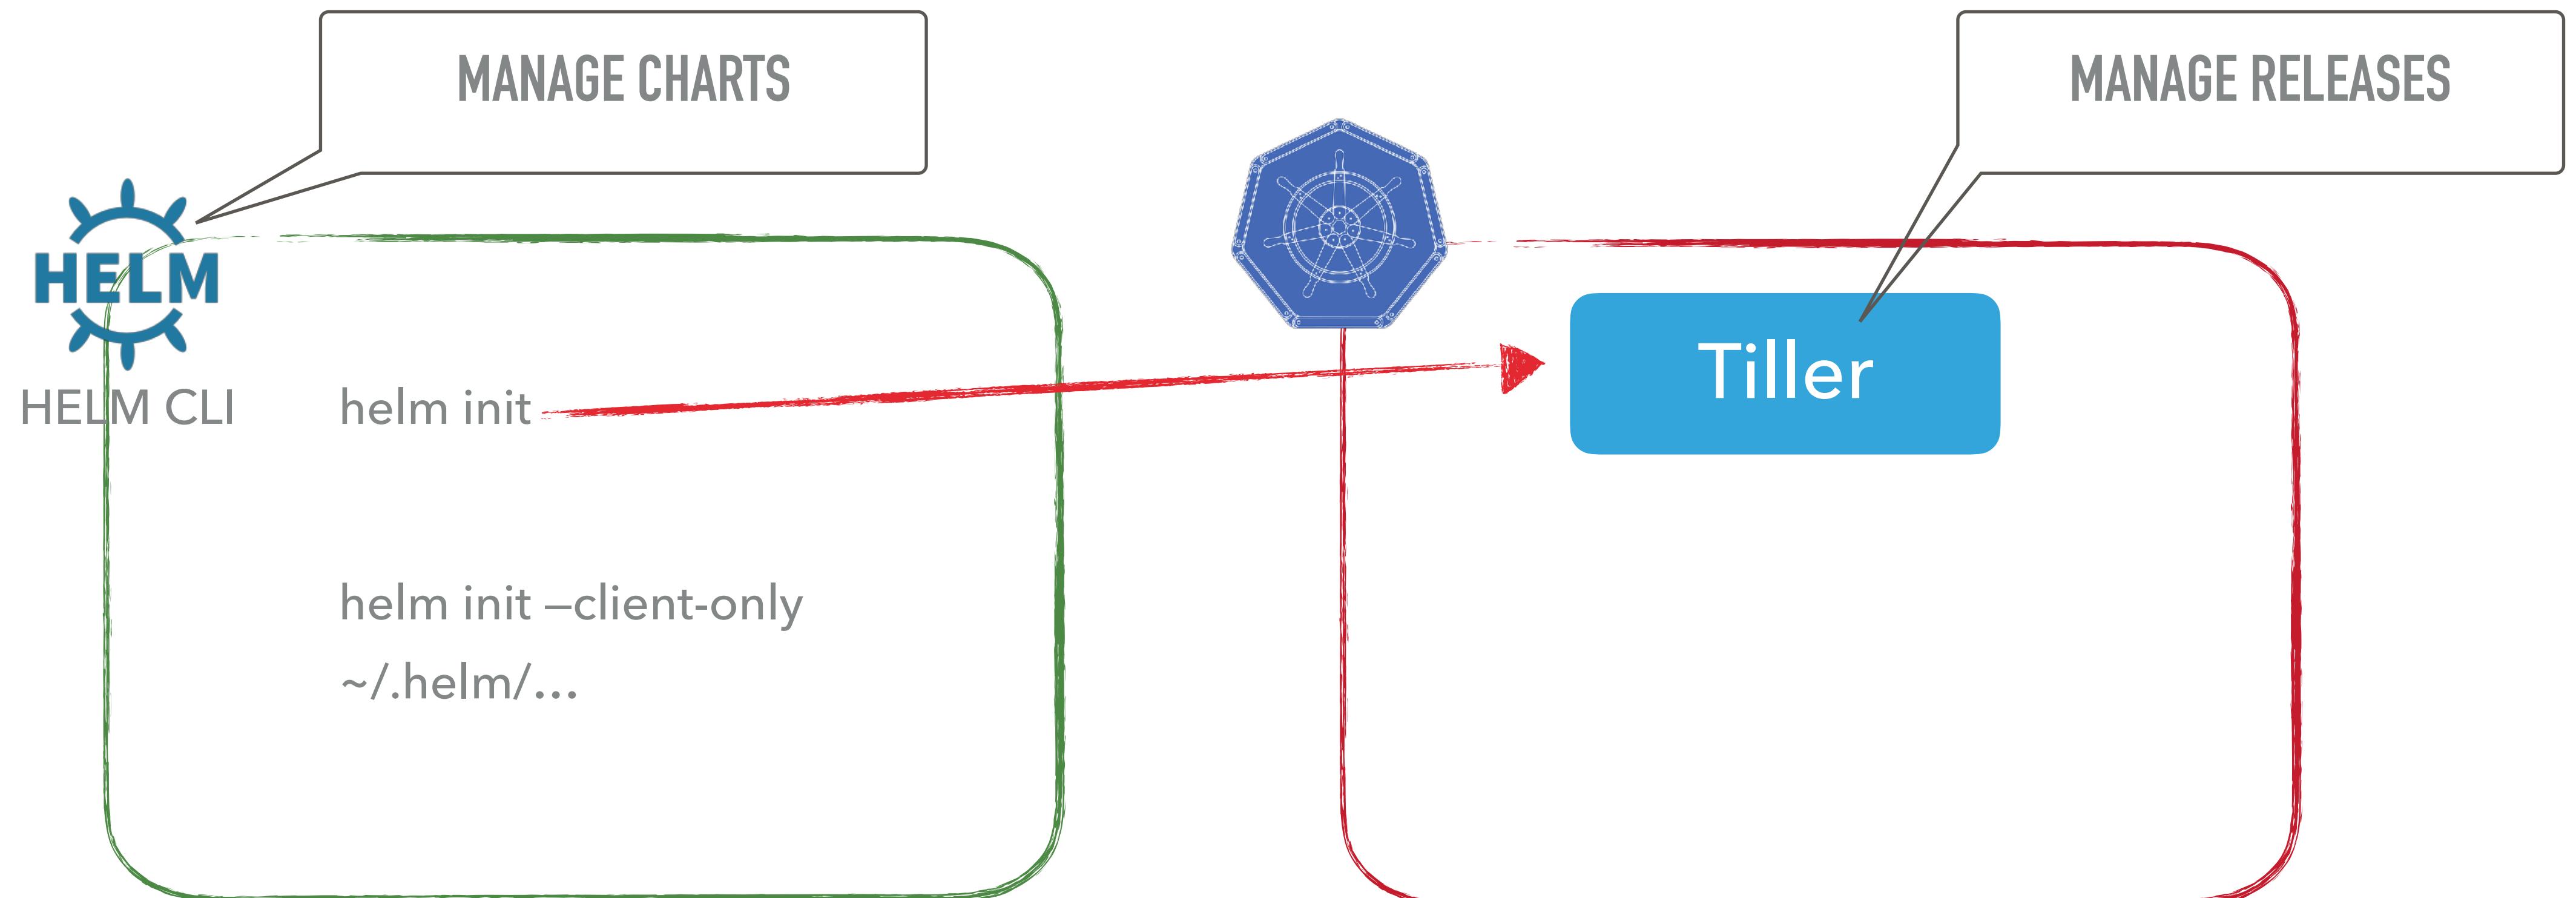
- ▶ Local chart development
- ▶ Managing repositories
- ▶ Interacting with the Tiller server
 - ▶ Sending charts to be installed
 - ▶ Asking for information about releases
 - ▶ Requesting upgrading or uninstalling of existing releases
- ▶ It can be used to serve your local charts via http ...

TILLER - “HELM SERVER”

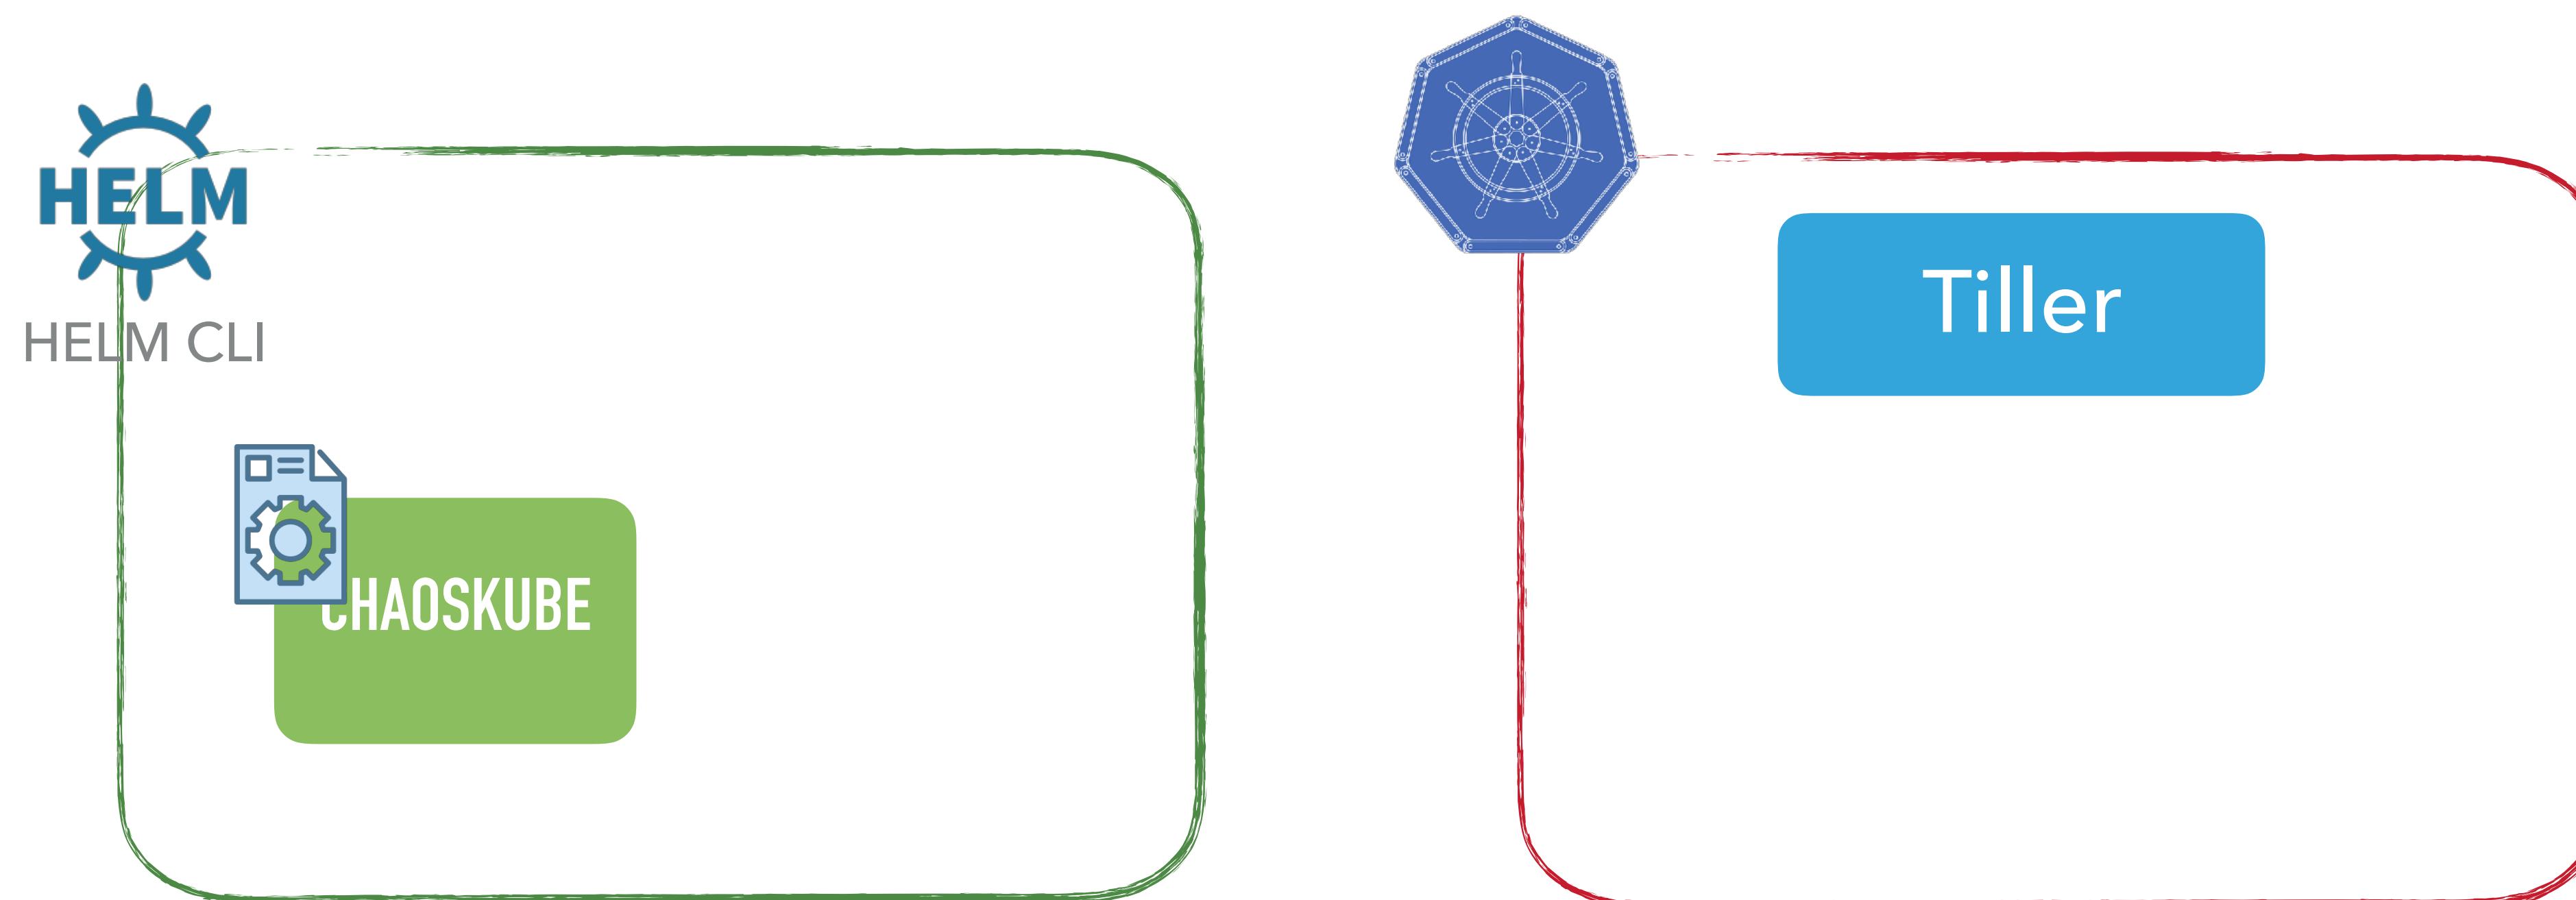


- ▶ Listening for incoming requests from the Helm client
- ▶ Combining a chart and configuration to build a release
- ▶ Installing charts into Kubernetes, and then tracking the subsequent release
- ▶ Upgrading and uninstalling charts by interacting with Kubernetes

HELM ARCHITECTURE - HELM INIT



HELM ARCHITECTURE - INSTALL AN EXISTING CHART





USE EXISTING CHARTS

CHAOSKUBE



- ▶ chaoskube is a “**chaos-monkey lite**” it basically takes down pod based on a schedule to test your resilience (and there are some tweaks via configuration)

<https://github.com/linki/chaoskube>

chaoskube



build passing coverage 85% release v0.10.0 container ready godoc reference

chaoskube periodically kills random pods in your Kubernetes cluster.

Why

Test how your system behaves under arbitrary pod failures.

Example

Running it will kill a pod in any namespace every 10 minutes by default.

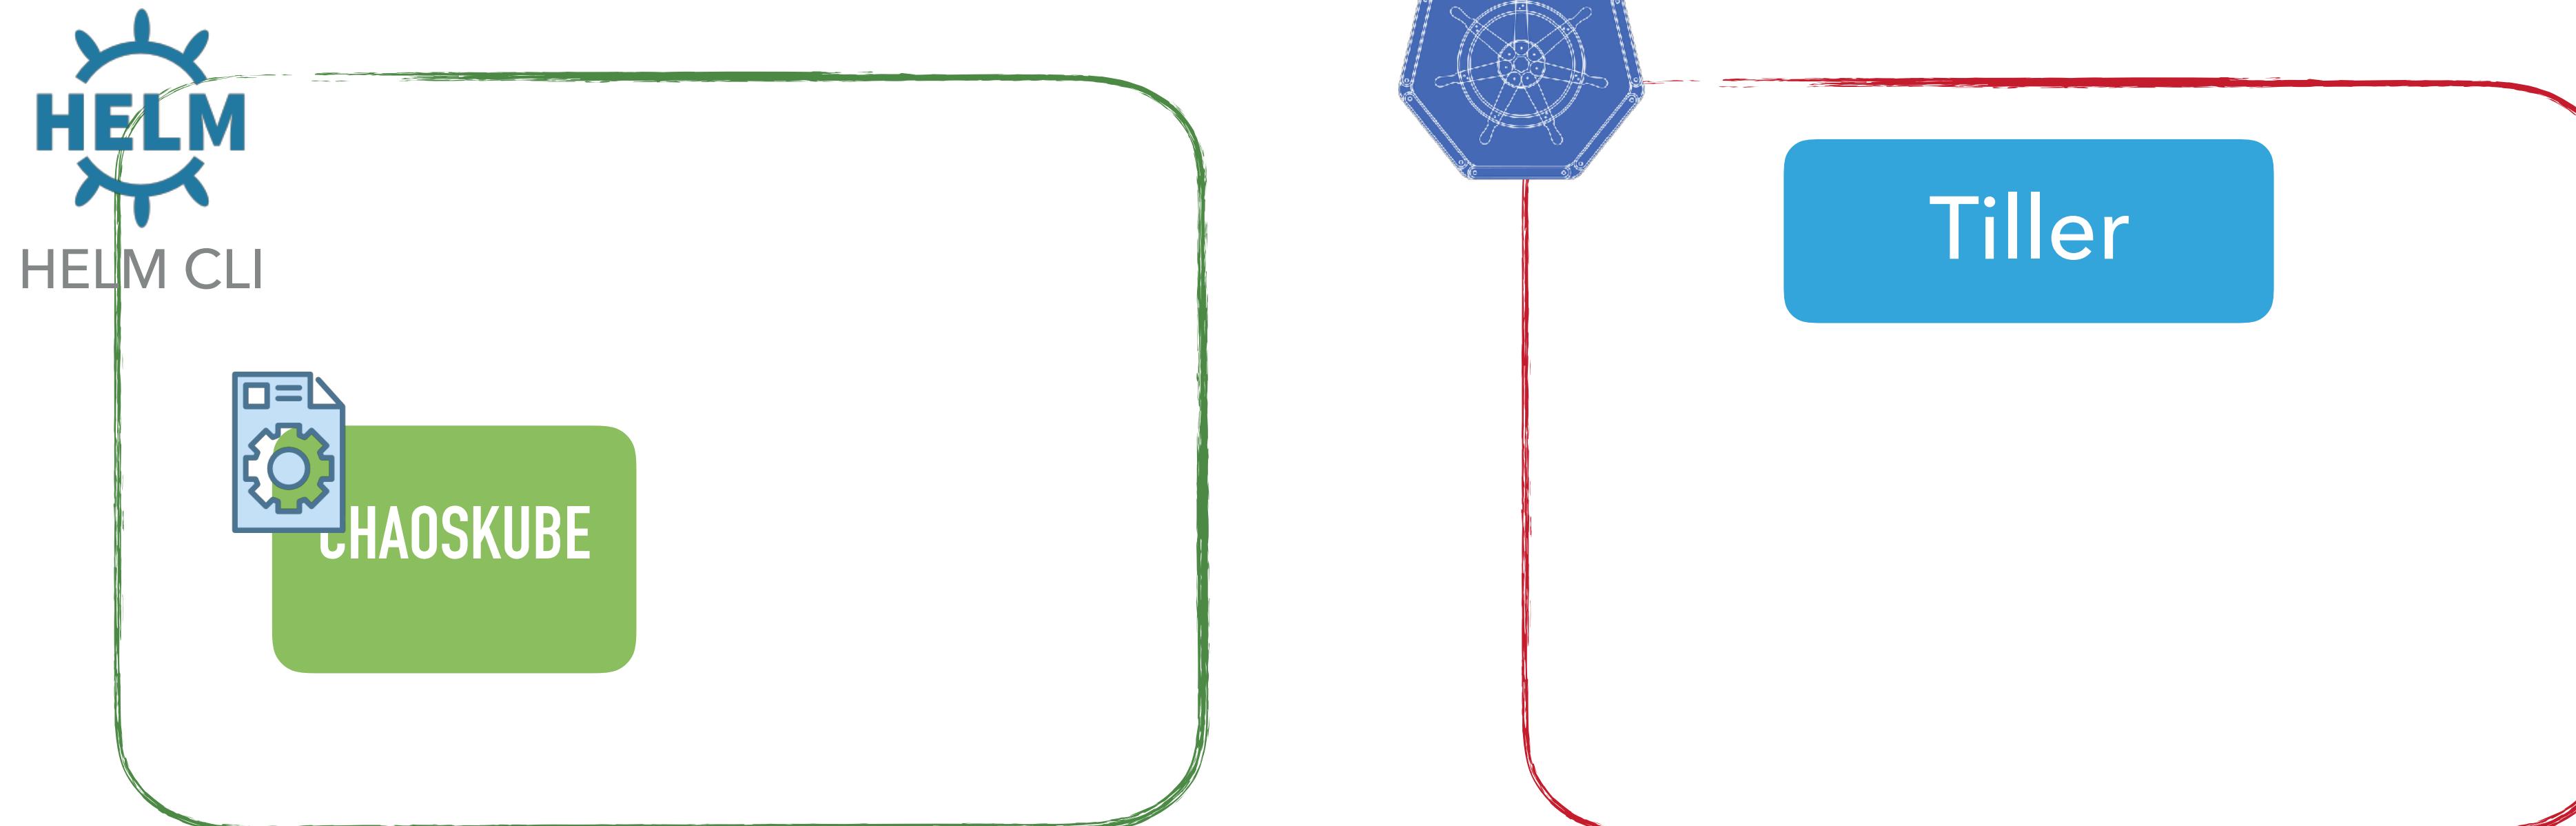
```
$ chaoskube
INFO[0000] starting up
INFO[0000] connecting to cluster
INFO[0000] setting pod filter
dryRun=true interval=10m0s version=v0.9.0
serverVersion=v1.9.3+coreos.0 master="https://kube.you.me"
annotations= labels= namespaces=
```

HELM ARCHITECTURE - WORK WITH REPOS

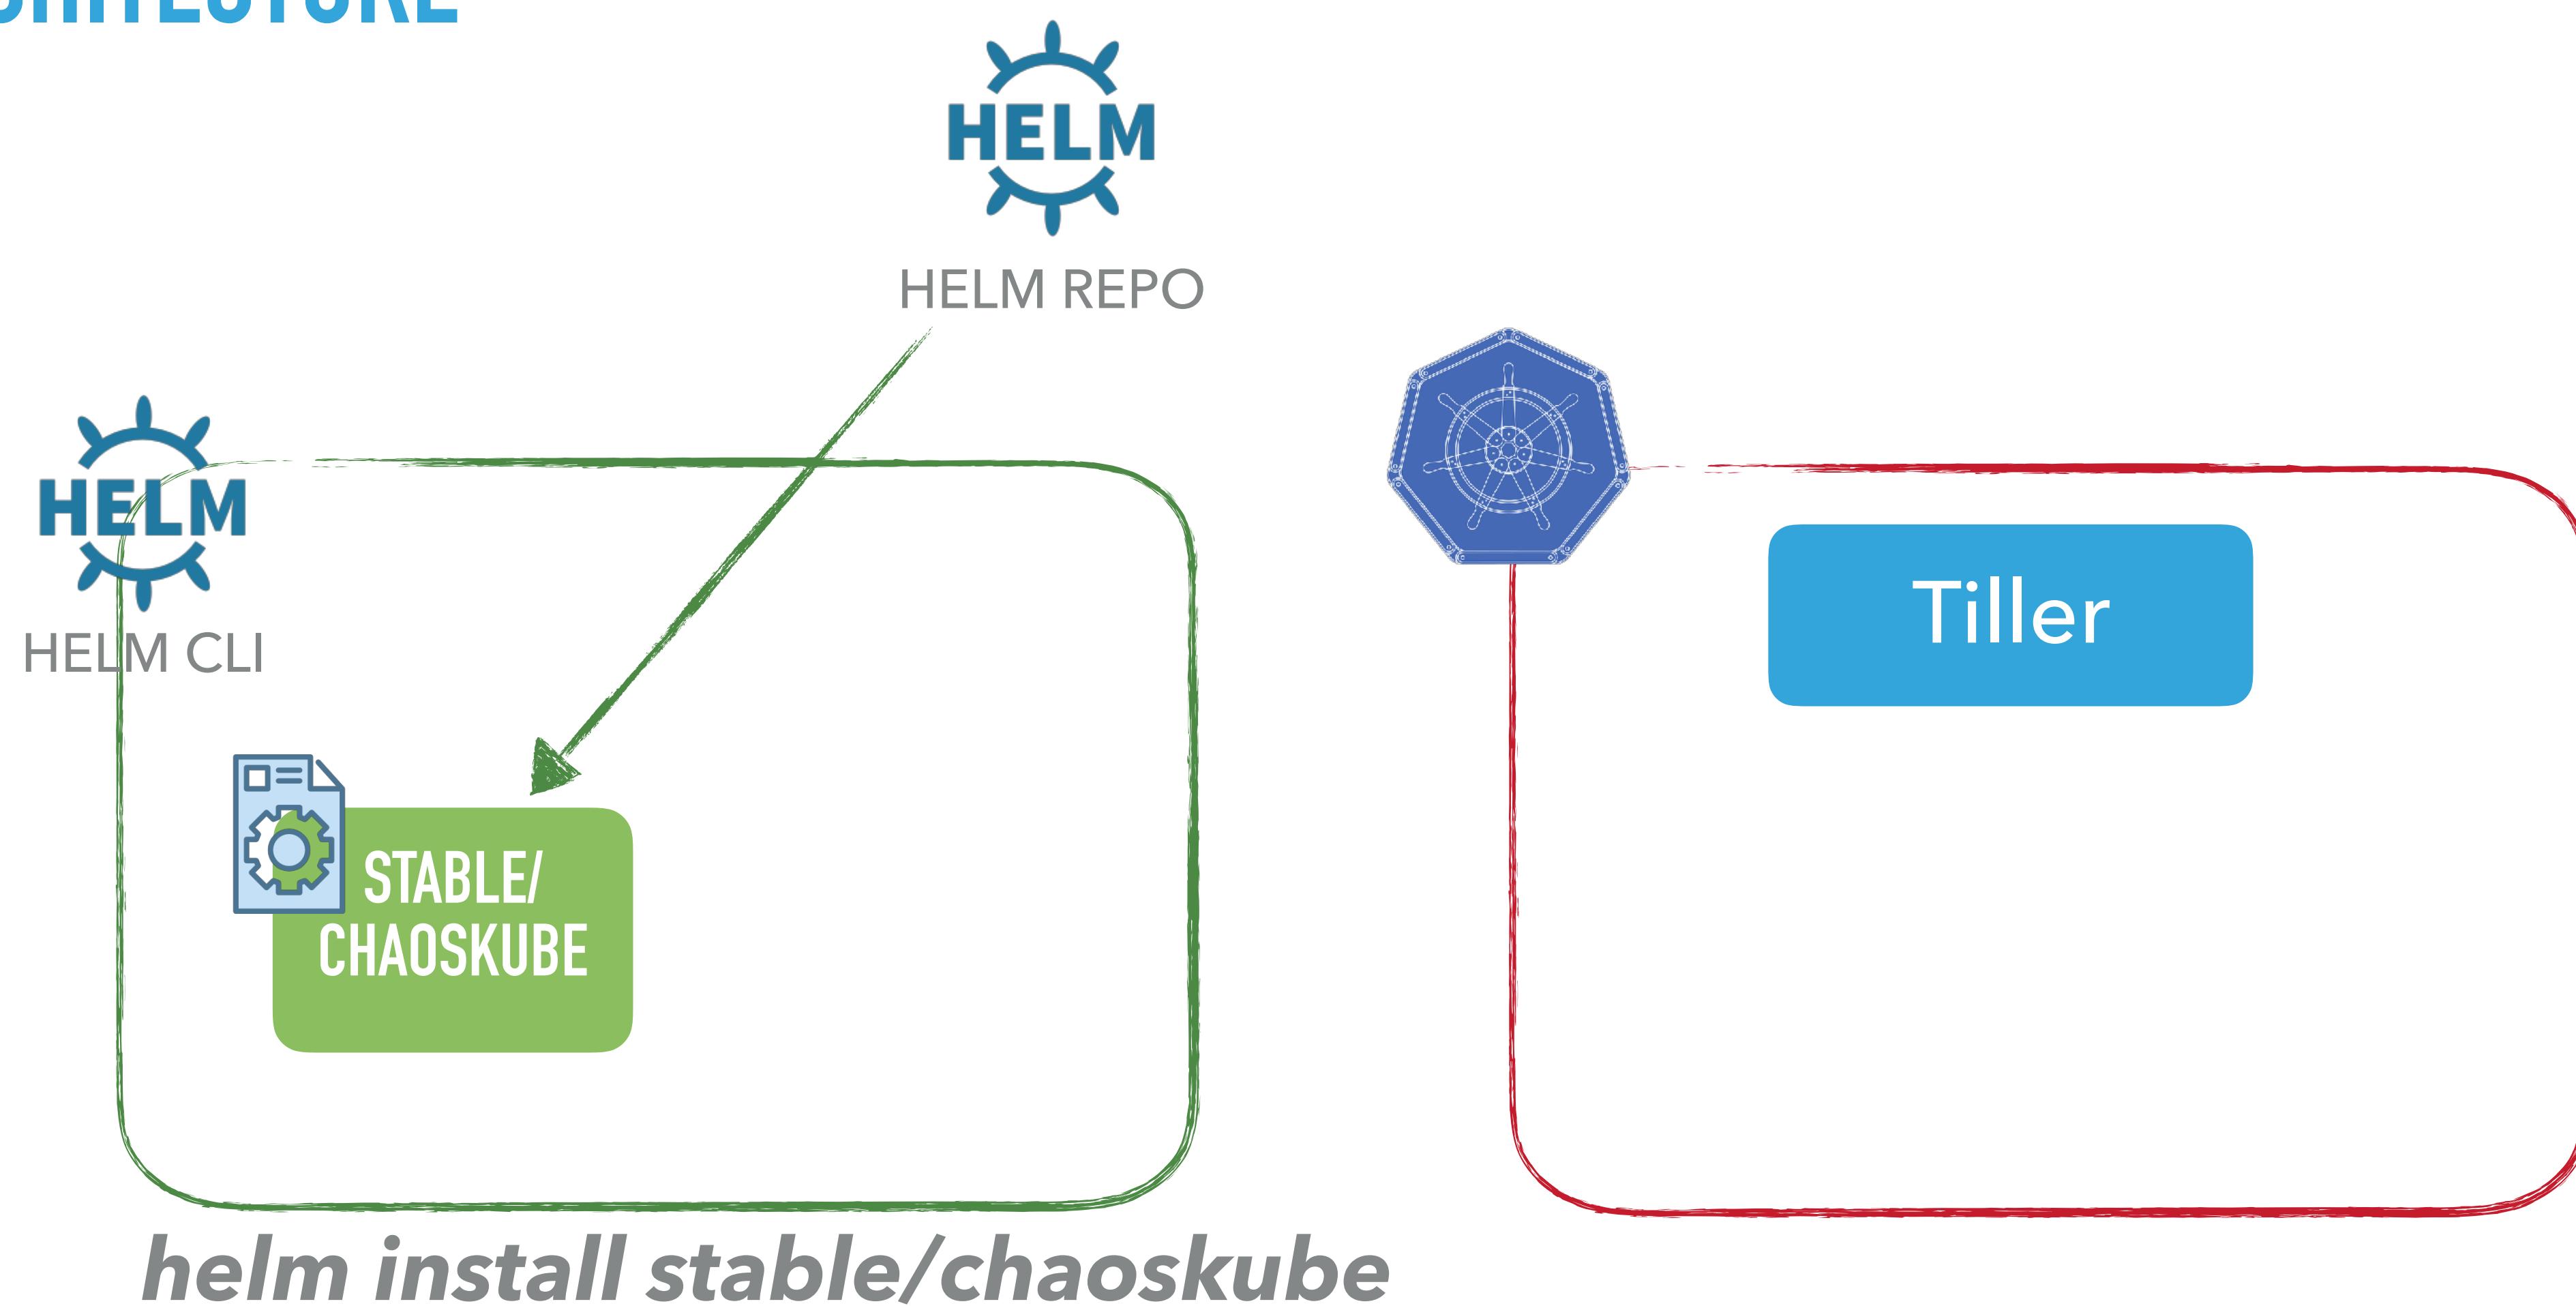


helm repo list

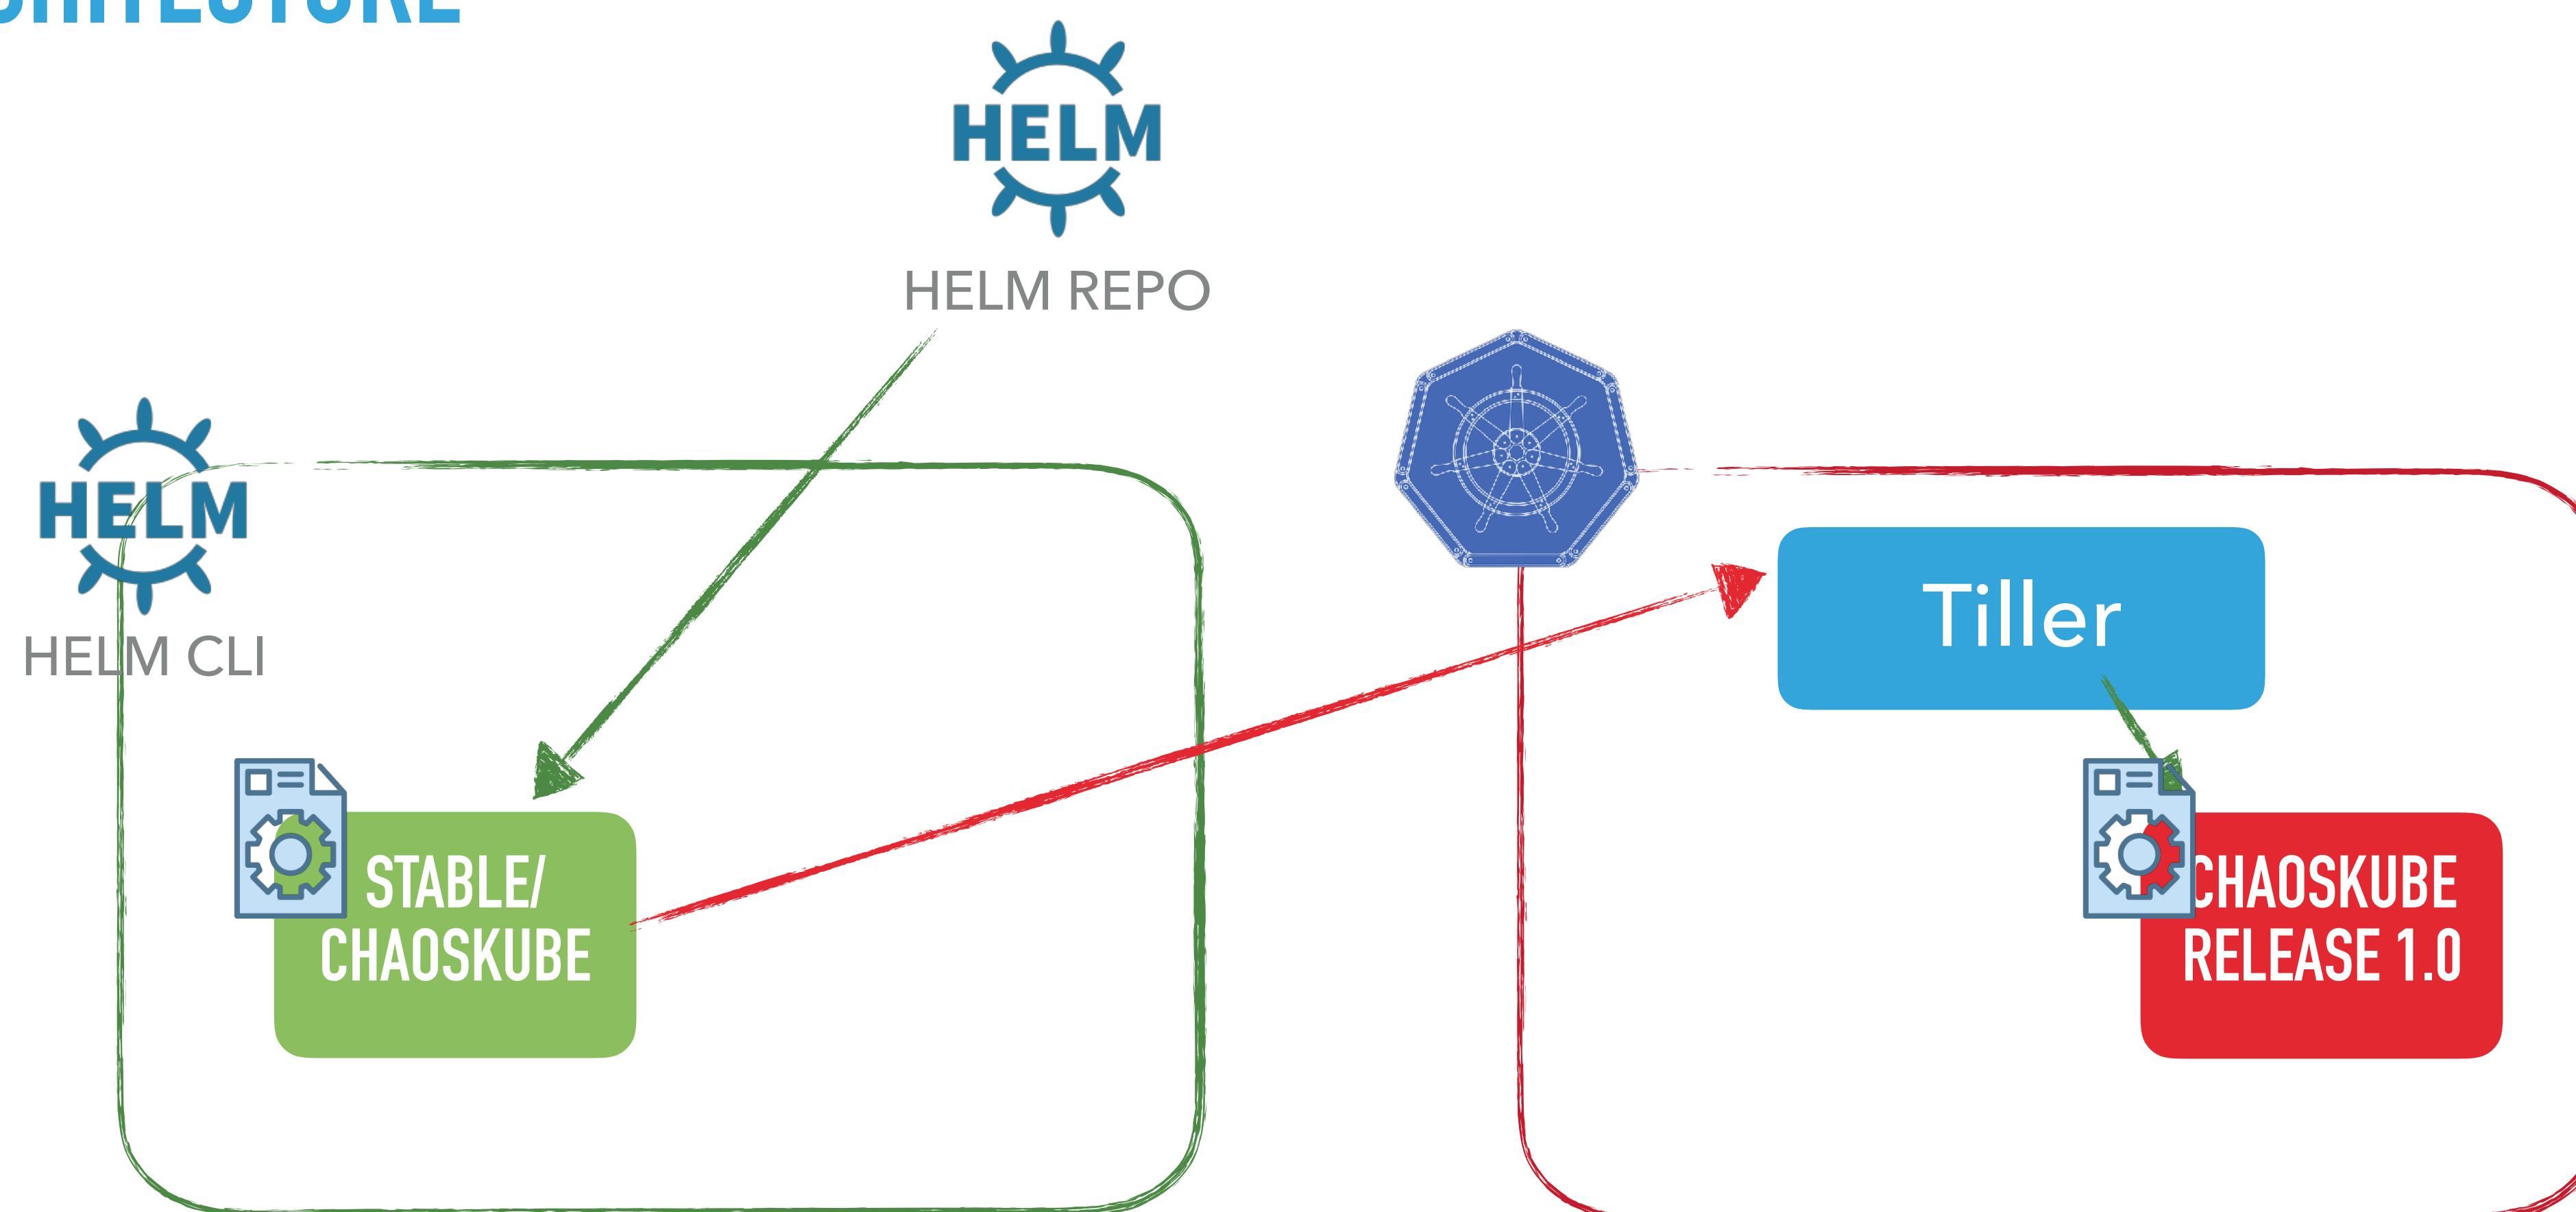
NAME	URL
stable	https://kubernetes-charts.storage.googleapis.com
local	http://127.0.0.1:8879/charts



HELM ARCHITECTURE



HELM ARCHITECTURE



CHARTS STRUCTURE



- - Chart.yaml
 - charts
 - templates
 - NOTES.txt
 - _helpers.tpl
 - deployment.yaml
 - ingress.yaml
 - service.yaml
 - values.yaml

2 directories, 7 files

Chart.yaml

```
apiVersion: v1
appVersion: 0.10.0
description: Chaoskube periodically kills random pods in your cluster.
engine: gotpl
home: https://github.com/linki/chaoskube
maintainers:
- email: linki+kubernetes.io@posteo.de
  name: linki
name: chaoskube
sources:
- https://github.com/linki/chaoskube
version: 0.10.0
```

⚠️ Chart and image version
{not always identical ...}

CHARTS STRUCTURE

```
-- Chart.yaml  
-- OWNERS  
-- README.md  
-- templates  
|   -- NOTES.txt  
|   -- _helpers.tpl  
|   -- clusterrole.yaml  
|   -- clusterrolebinding.yaml  
|   -- deployment.yaml  
|   -- role.yaml  
|   -- rolebinding.yaml  
|   -- serviceaccount.yaml  
-- values.yaml
```

1 directory, 12 files

▶ /templates for your **resource definitions** such as:

- ▶ Service
- ▶ Deployment
- ▶ configMap
- ▶ Secret
- ▶ Daemonset
- ▶ Pod
- ▶ Ingress
- ▶ Job
- ▶ persistentvolume
- ▶ persistentvolumeclaim
- ▶ Service account



CHARTS STRUCTURE

```
-- Chart.yaml  
-- OWNERS  
-- README.md  
-- templates  
| -- NOTES.txt  
| -- _helpers.tpl  
| -- clusterrole.yaml  
| -- clusterrolebinding.yaml  
| -- deployment.yaml  
| -- role.yaml  
| -- rolebinding.yaml  
| -- serviceaccount.yaml  
-- values.yaml
```

1 directory, 12 files

values.yaml

```
# container name  
name: chaoskube  
  
# docker image  
image: quay.io/linki/chaoskube  
  
# docker image tag  
imageTag: v0.10.0  
  
# number of replicas to run  
replicas: 1  
  
# interval between pod terminations  
interval: 10m  
  
# label selector to filter pods by, e.g. app=foo,stage!=prod  
labels:  
  
# annotation selector to filter pods by, e.g.  
chaos.alpha.kubernetes.io/enabled=true  
annotations:  
...
```



CHARTS STRUCTURE



```
-- Chart.yaml
-- OWNERS
-- README.md
-- templates
| -- NOTES.txt
| -- _helpers.tpl
| -- clusterrole.yaml
| -- clusterrolebinding.yaml
| -- deployment.yaml
| -- role.yaml
| -- rolebinding.yaml
| -- serviceaccount.yaml
-- values.yaml
```

1 directory, 12 files

serviceaccount.yaml

```
{-- if .Values.rbac.create --}
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
{{ include "labels.standard" . | indent 4 }}
    name: {{ printf "%s-%s" .Release.Name .Values.name }}
{-- end --}
```

Control Structures:

{ render based on boolean }

values.yaml

```
# container name
name: chaoskube
```

CHARTS STRUCTURE

```
• -- Chart.yaml  
|-- charts  
|-- templates  
|   |-- NOTES.txt  
|   |-- _helpers.tpl  
|   |-- deployment.yaml  
|   |-- ingress.yaml  
|   '-- service.yaml  

```

2 directories, 7 files

_helpers.tpl

```
/*  
Expand the name of the chart.  
*/}  
{-- define "chaoskube.name" --}  
{-- default .Chart.Name .Values.nameOverride |  

```

Custom functions +
set default based on
Environment



CHARTS STRUCTURE



- ▶ A **HELM standard way** of showing information the end user can use in order to access the application.

NOTES.txt

```
chaoskube is running and will kill arbitrary pods every  
{{ .Values.interval }}.
```

You can follow the logs to see what chaoskube does:

```
POD=$(kubectl -n {{ .Release.Namespace }} get pods -  
l='release={{ template "chaoskube.fullname" . }}' --  
output=jsonpath='{.items[0].metadata.name}  
' )
```

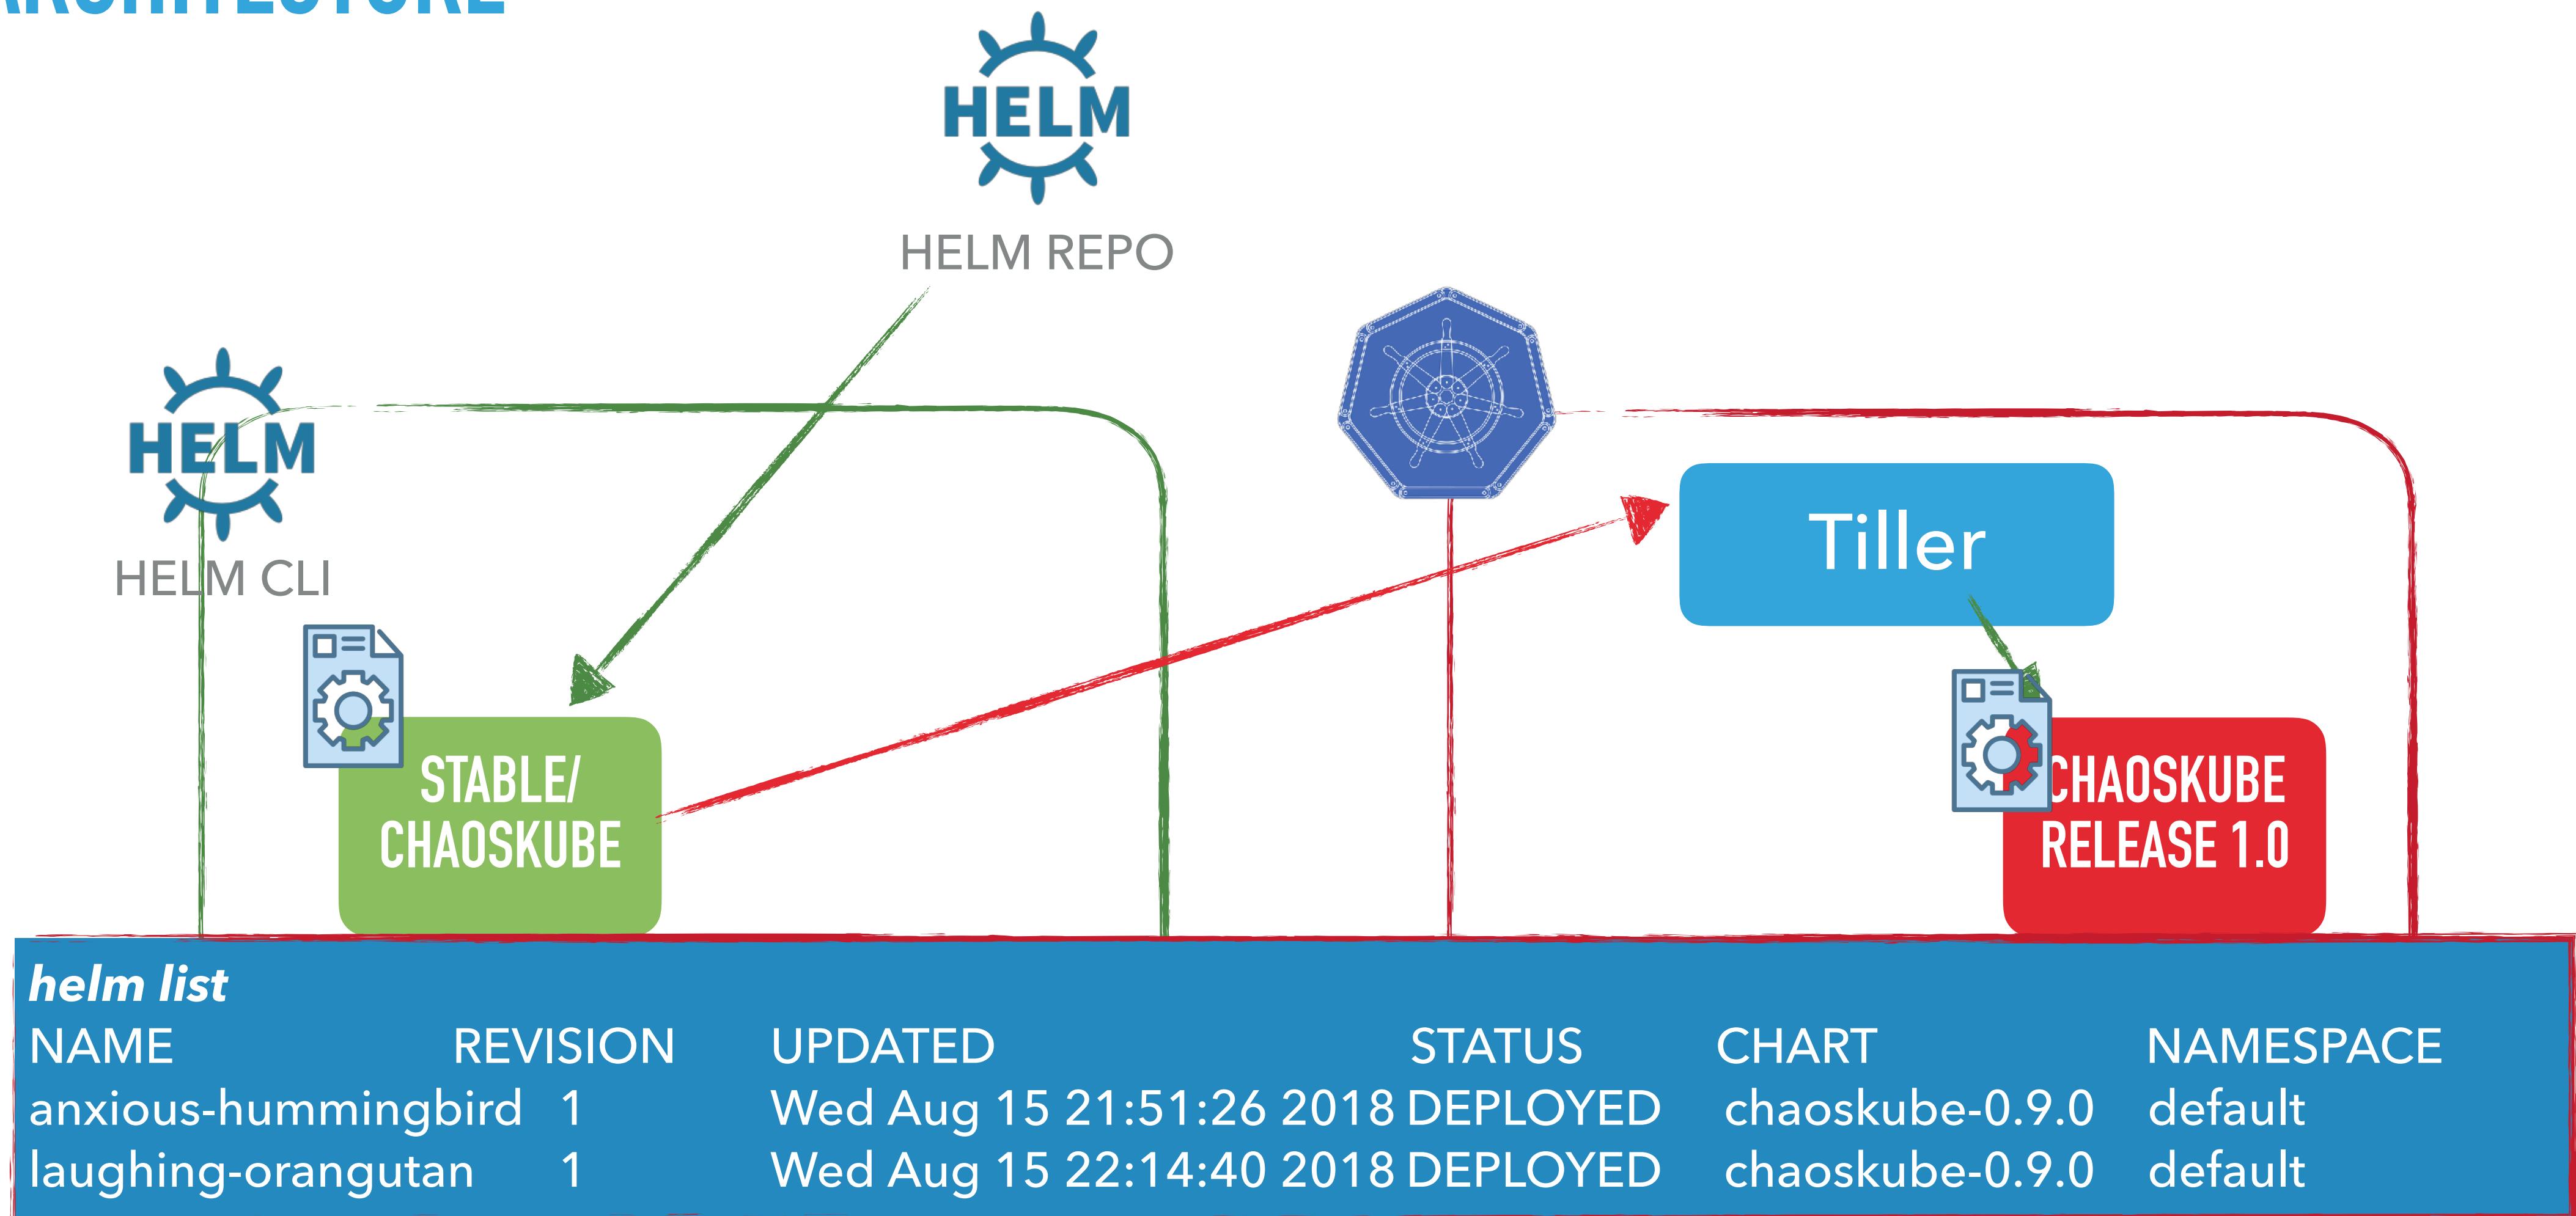
```
kubectl -n {{ .Release.Namespace }} logs -f $POD
```

```
{{ if .Values.dryRun }}
```

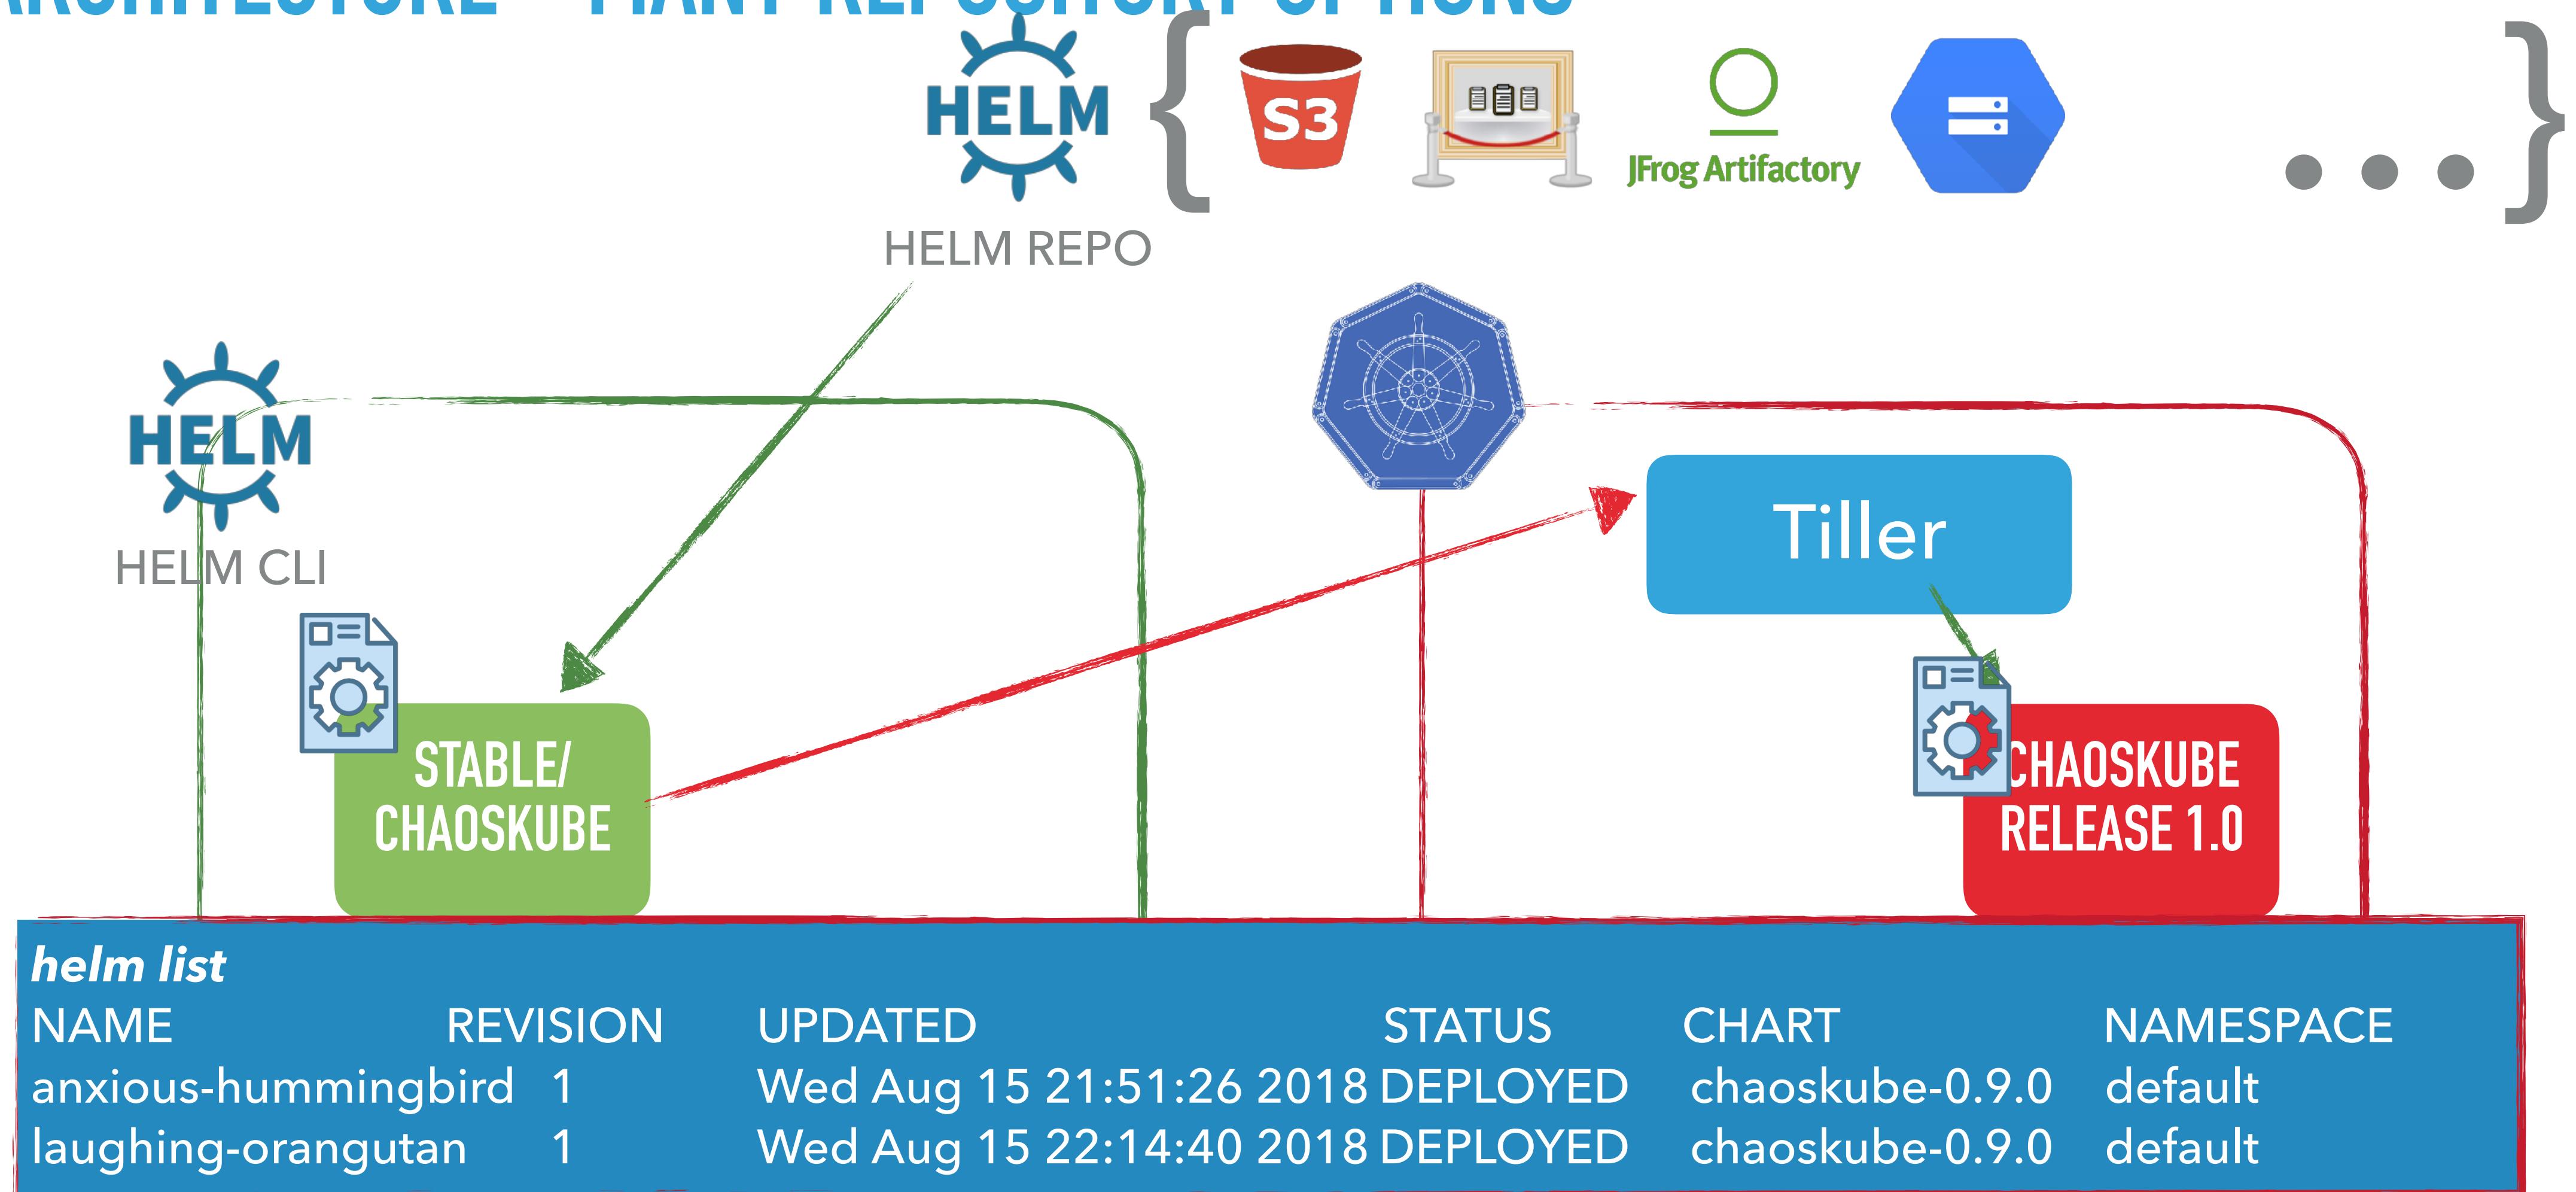
You are running **in** dry-run mode. No pod **is** actually terminated.

```
{{ end -}}
```

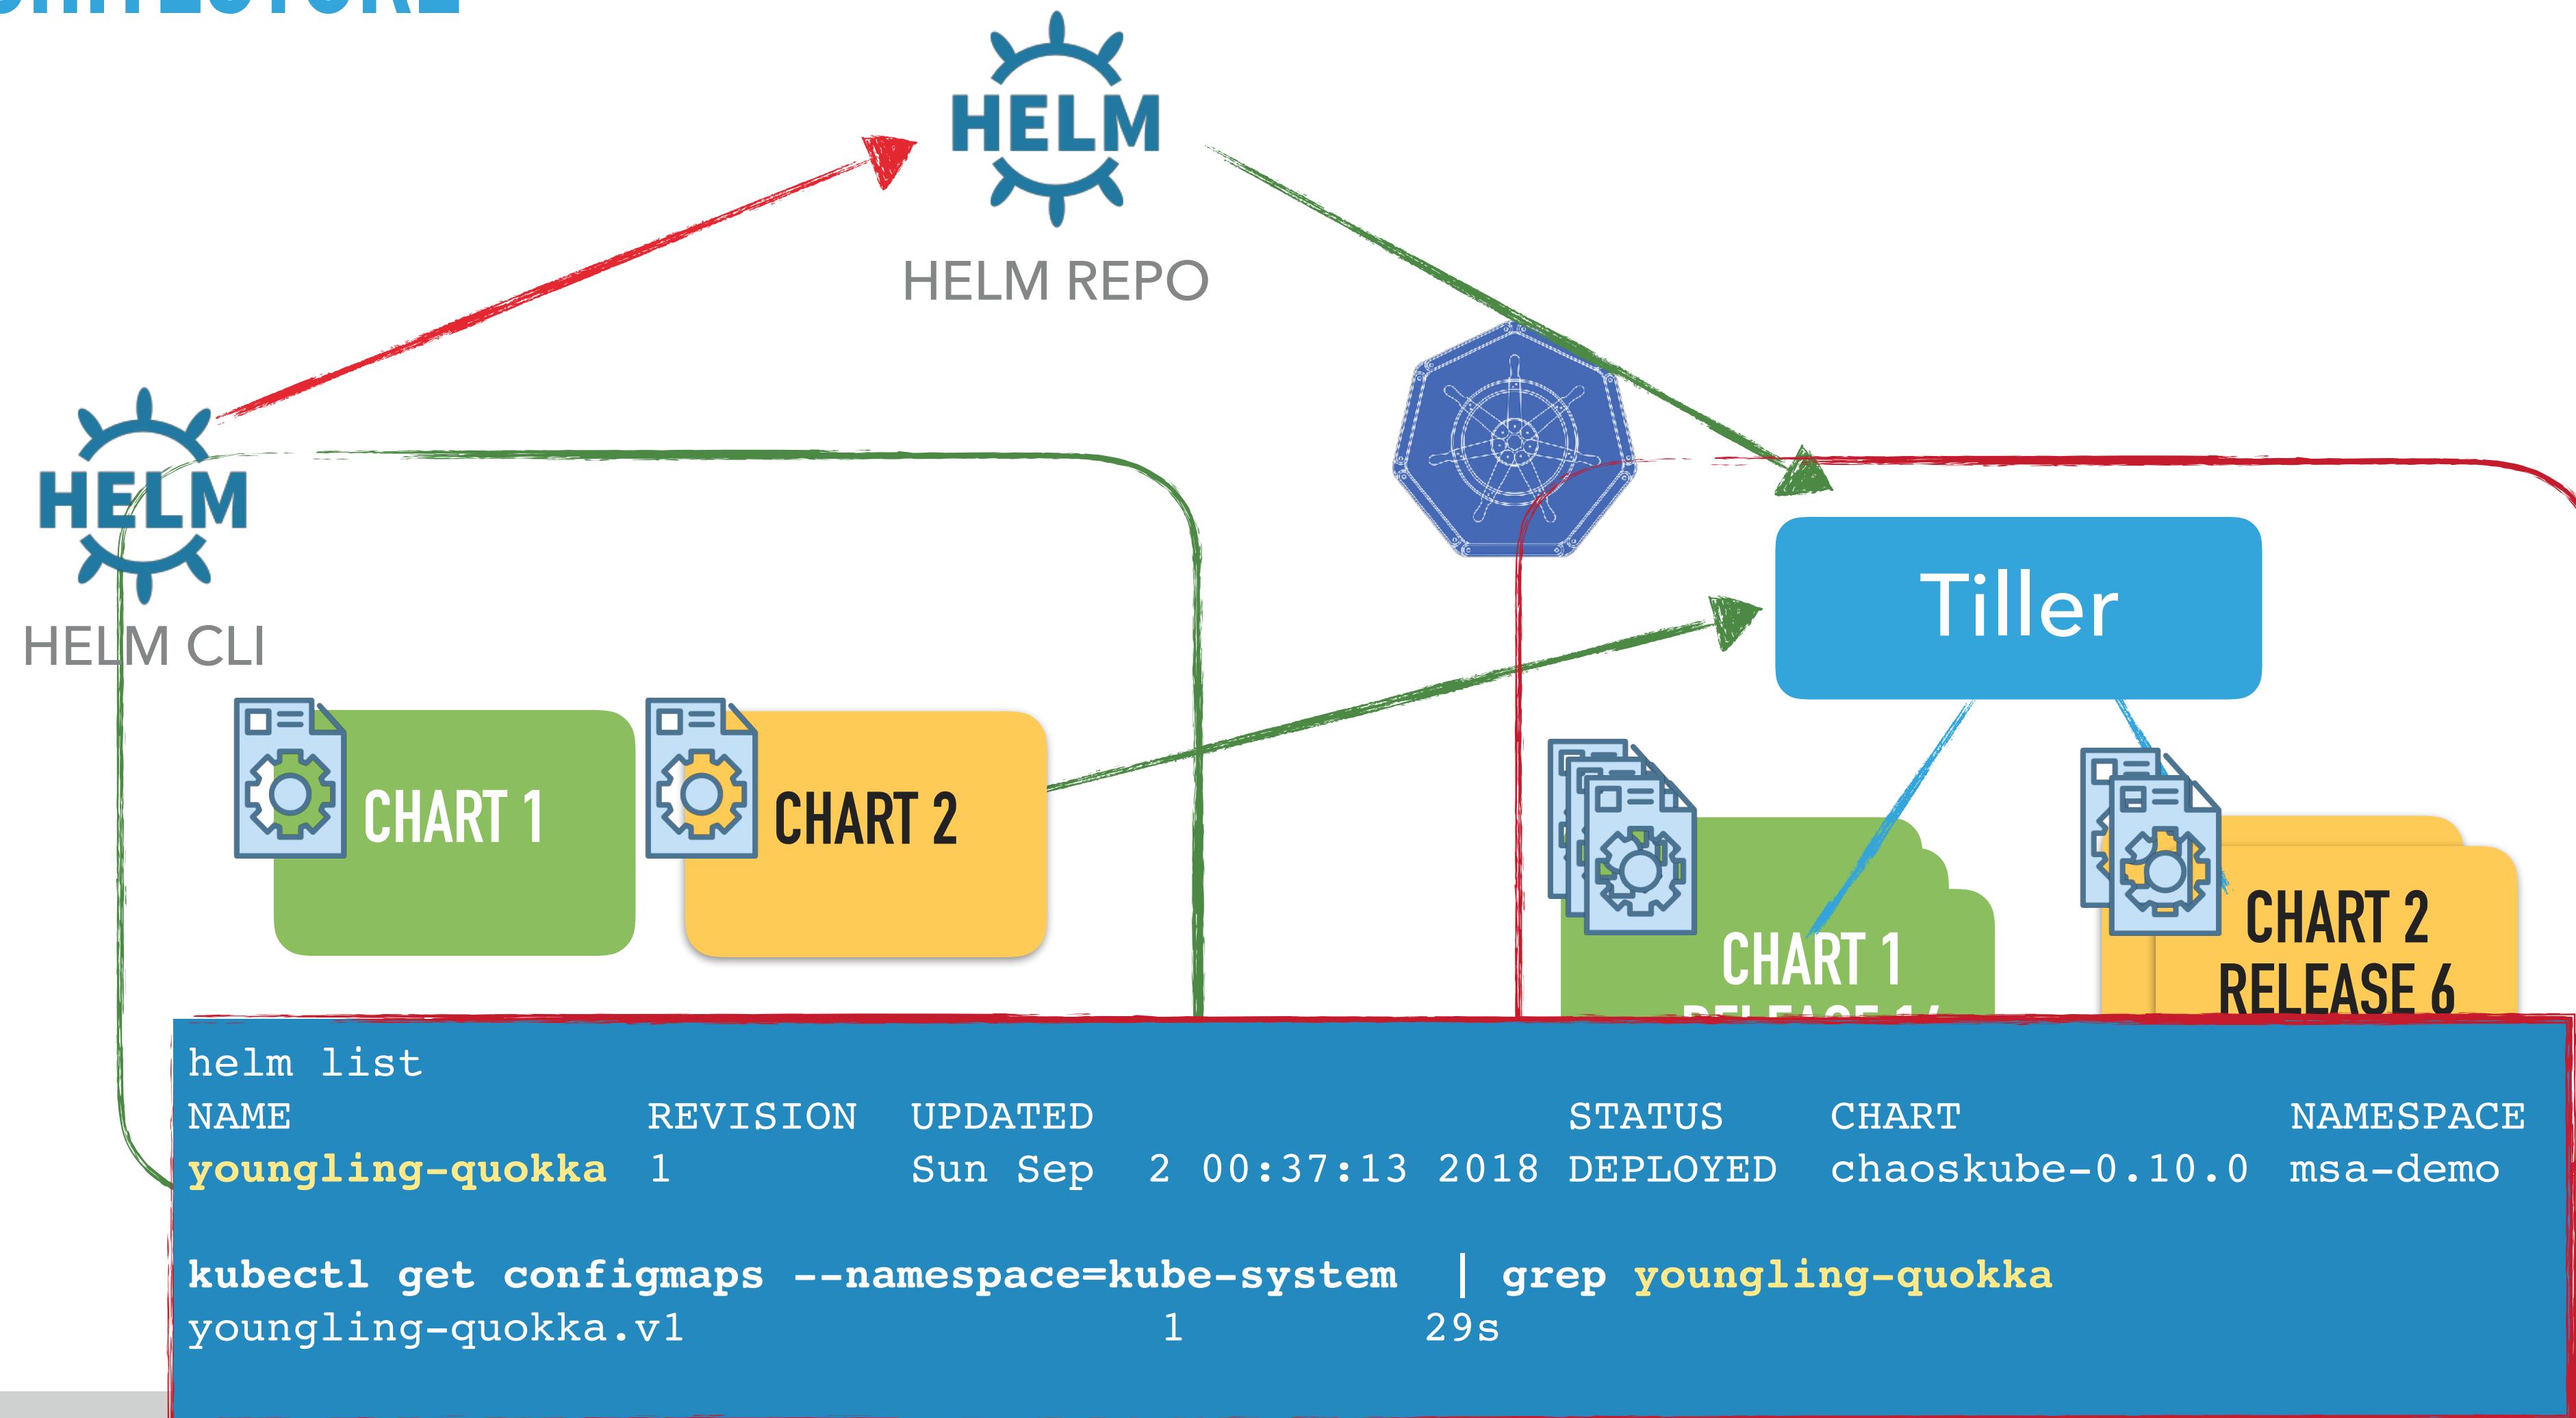
HELM ARCHITECTURE



HELM ARCHITECTURE - MANY REPOSITORY OPTIONS



HELM ARCHITECTURE





“CHARTIFY” (& REUSE) . . .

OUR “COMPLICATED” MICRO SERVICE APPLICATION - DOCKERIZED

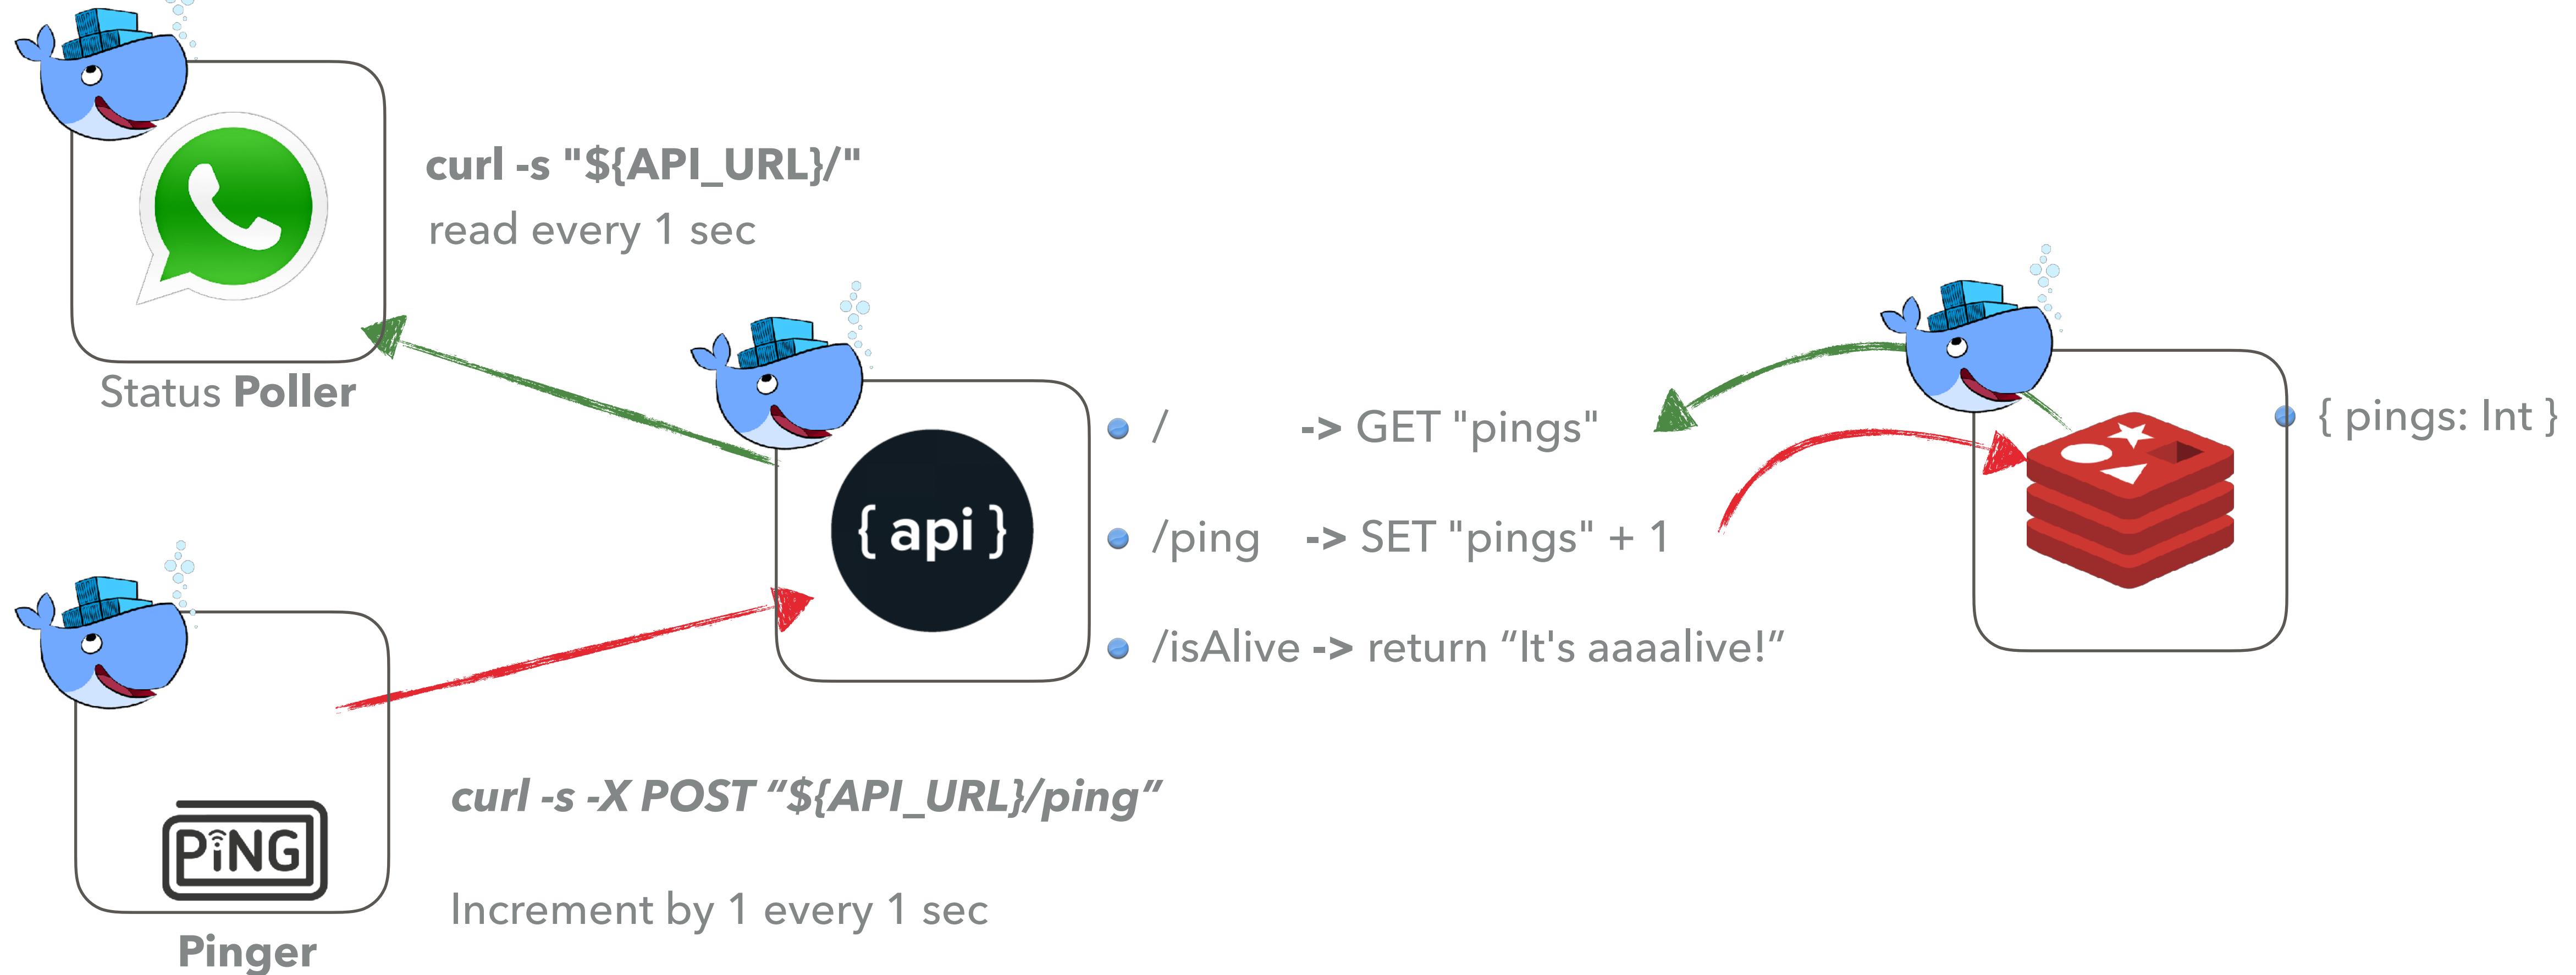


CHART PER COMPONENT

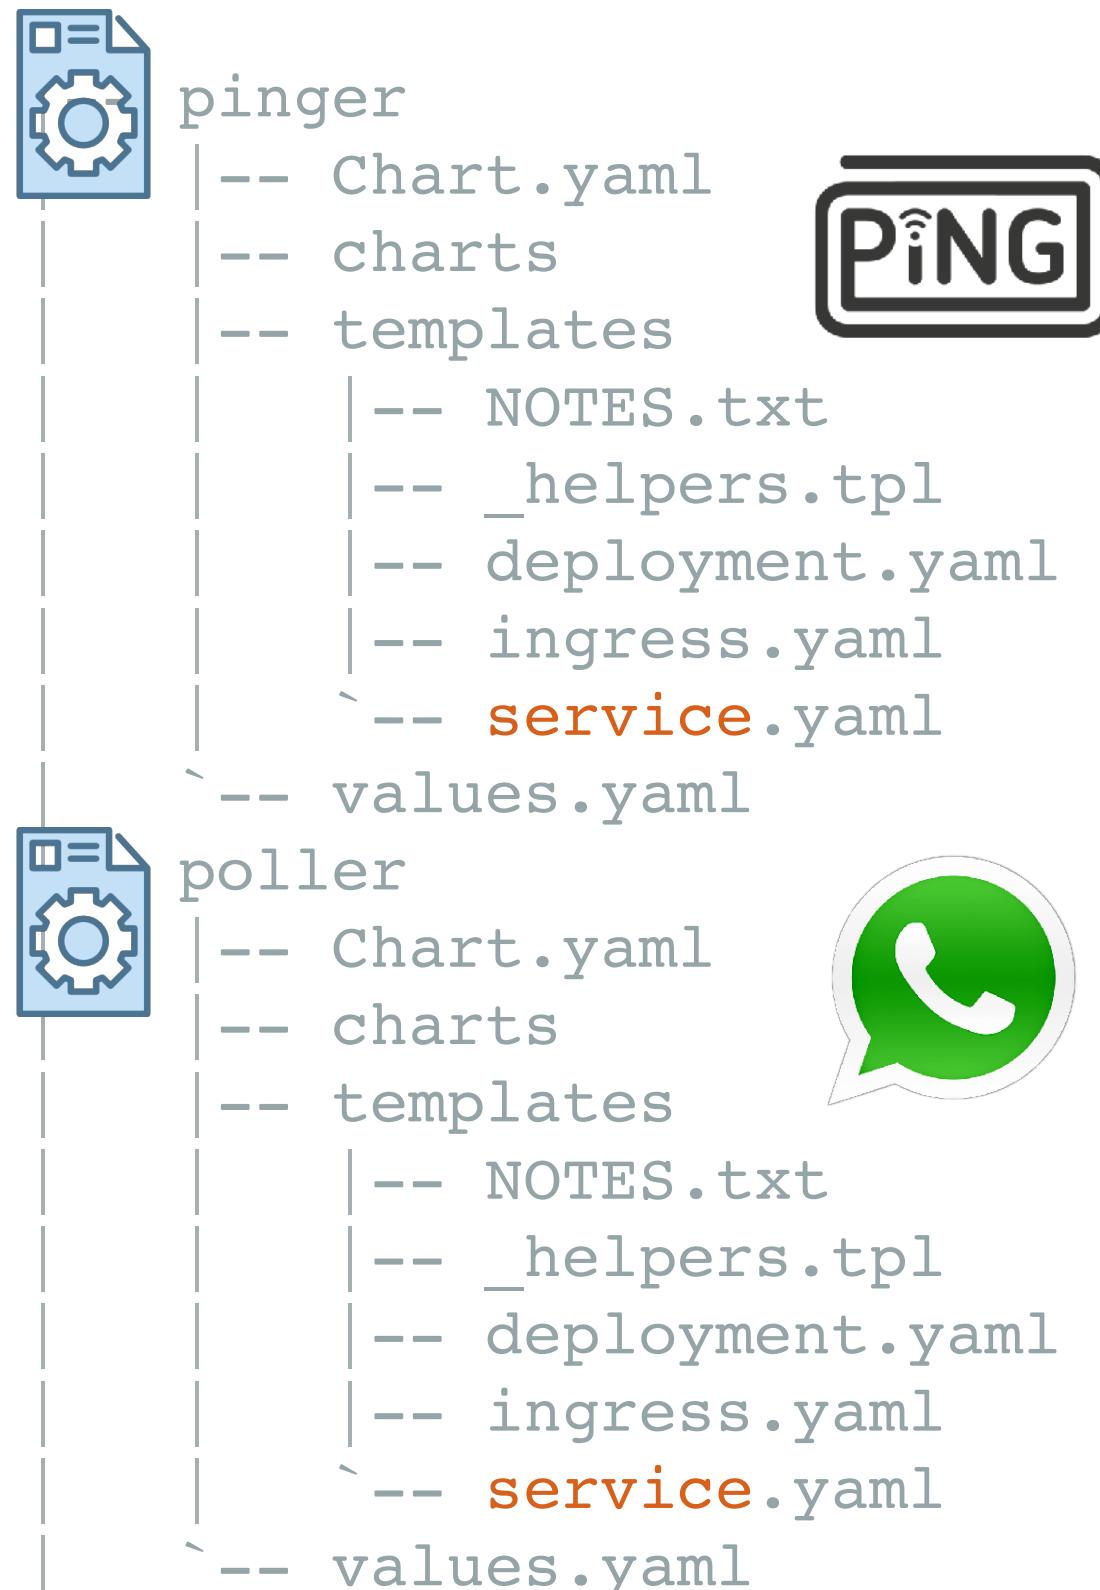
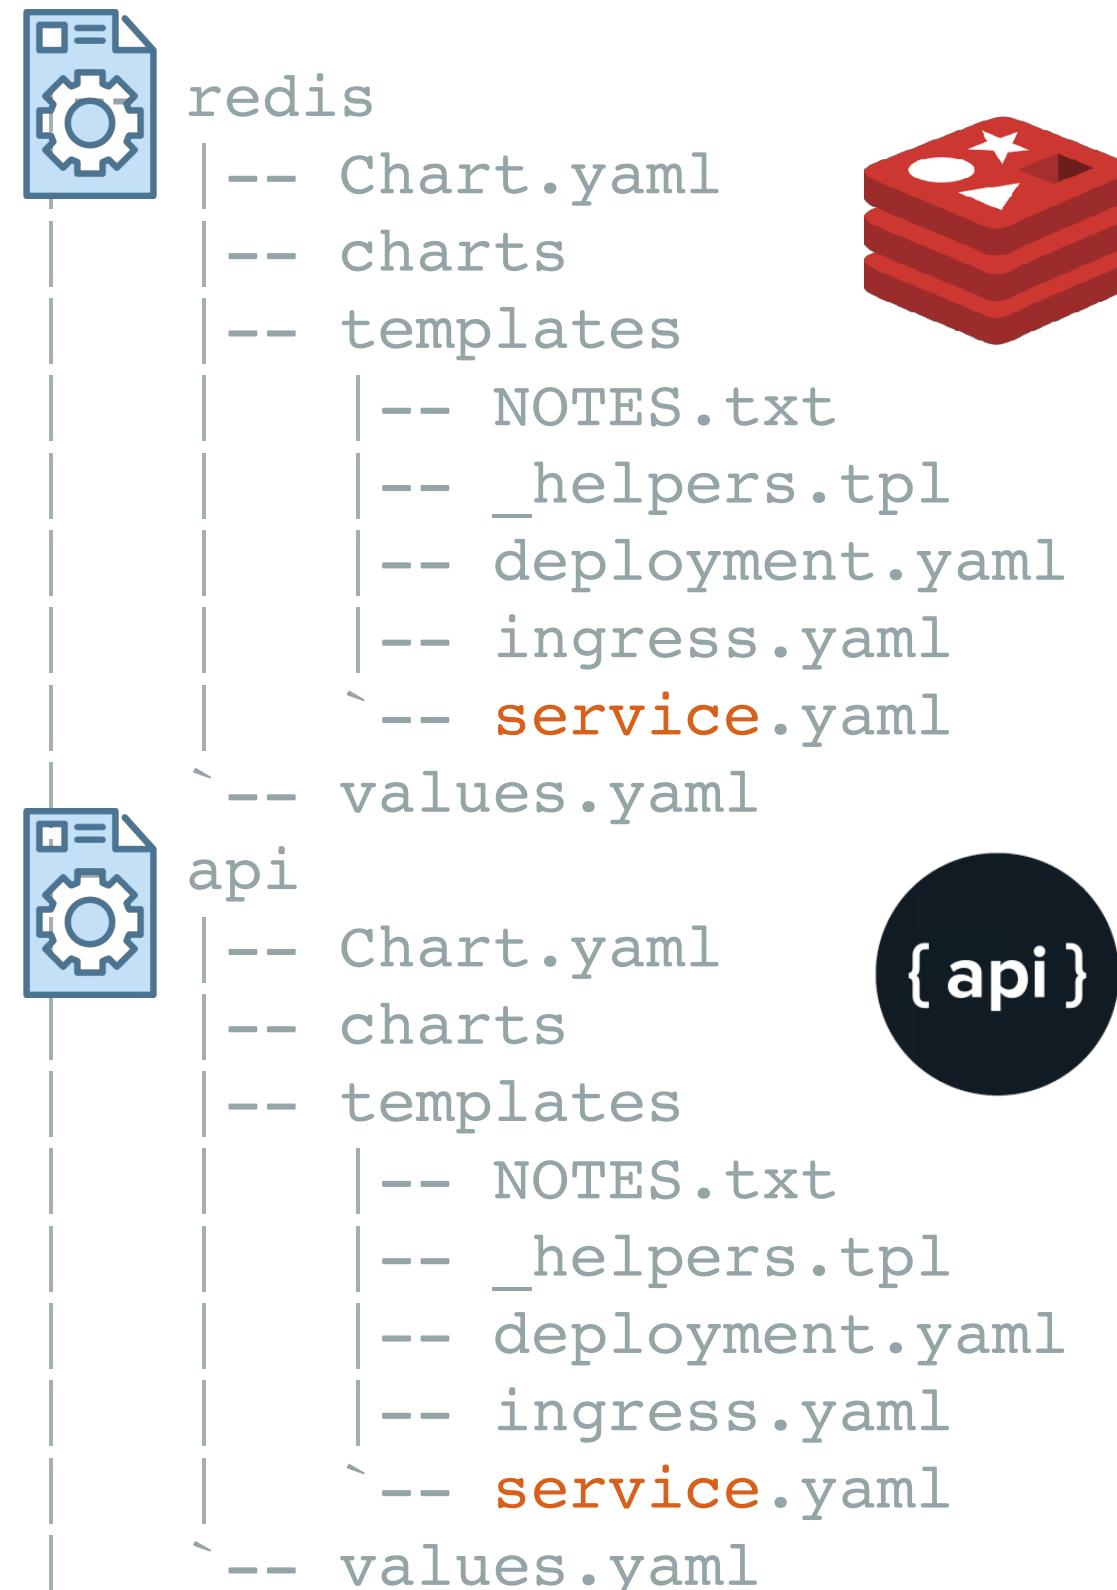
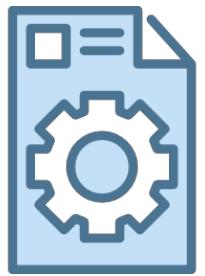


CHART PER COMPONENT



“MSA” CHART

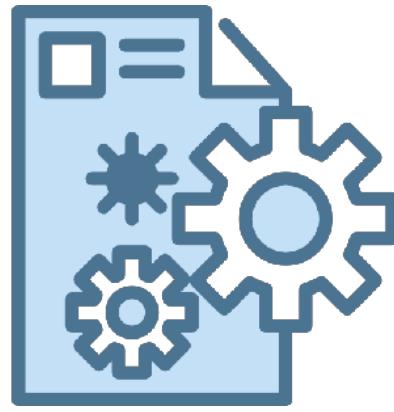


```
•
  |-- Chart.yaml
  |-- charts
  |-- templates
    |-- NOTES.txt
    |-- _helpers.tpl
    |-- api-deployment.yaml
    |-- api-service.yaml
    |-- deployment.yaml
    |-- pinger-deployment.yaml
    |-- poller-deployment.yaml
    '-- service.yaml
  '-- values.yaml
```



INTRODUCTION TO HELM

VALUES



values.yaml

```
api:
  name: api
  replicaCount: 1
  image:
    repository: shelleg/msa-api
    tag: config
    pullPolicy: Always
    containerPort: 8080
  service:
    enabled: true
    type: ClusterIP
    port: 8080
  resources: {}
  nodeSelector: {}
  tolerations: []
  affinity: {}

pinger:
  name: pinger
  replicaCount: 1
  image:
    repository: shelleg/msa-pinger
    tag: latest
    pullPolicy: Always
  service:
    enabled: false
  resources: {}
  nodeSelector: {}
  tolerations: []
  affinity: {}

poller:
  name: poller
  replicaCount: 1
  image:
    repository: shelleg/msa-poller
    tag: latest
    pullPolicy: Always
  service:
    enabled: false
  resources: {}
  nodeSelector: {}
  tolerations: []
  affinity: {}
```



| This works ;) but ...

TEMPLATES



api-deployment.yaml

```
apiVersion: apps/v1beta2
kind: Deployment
metadata:
  name: {{ .Values.api.name }}
  labels:
    app: {{ .Values.api.name }}
    chart: {{ template "msa.chart" . }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  replicas: {{ .Values.api.replicaCount }}
  selector:
    matchLabels:
      app: {{ .Values.api.name }}
      release: {{ .Release.Name }}
  template:
    metadata:
      labels:
        app: {{ .Values.api.name }}
        release: {{ .Release.Name }}
```



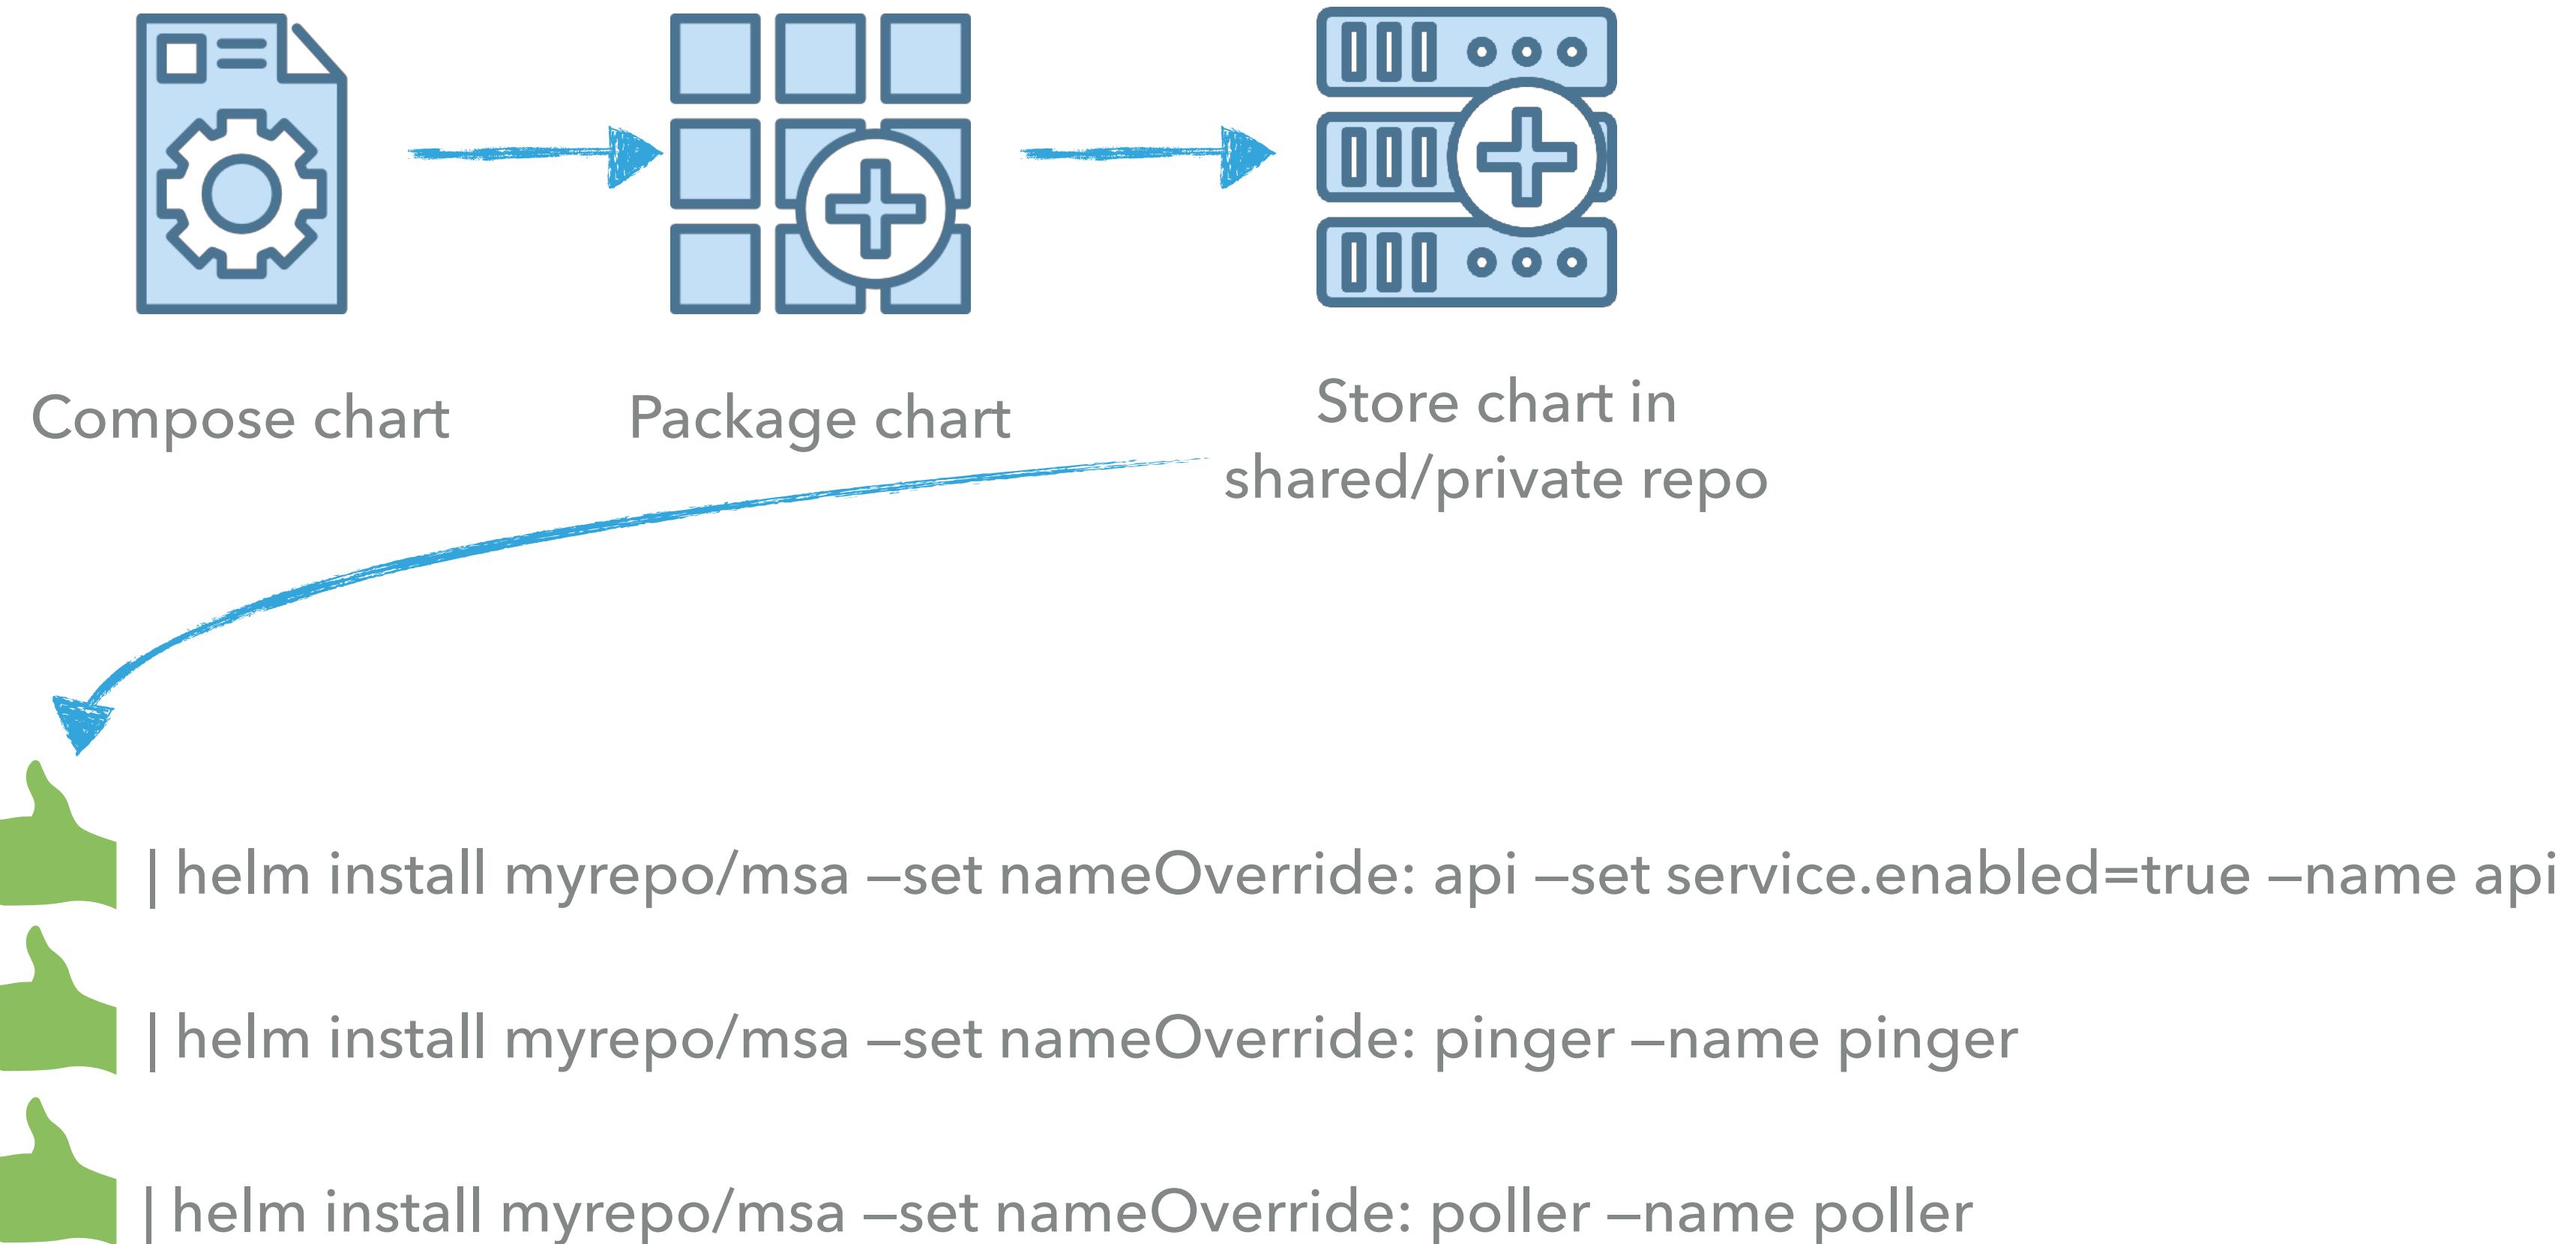
api-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: {{ .Values.api.name }}
  labels:
    app: {{ .Values.api.name }}
    chart: {{ template "msa.chart" . }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  type: {{ .Values.api.service.type }}
  ports:
    - port: {{ .Values.api.service.port }}
      targetPort: http
      protocol: TCP
      name: http
  selector:
    app: {{ .Values.api.name }}
    release: {{ .Release.Name }}
```

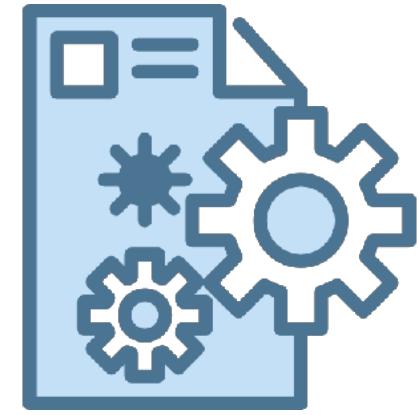


INTRODUCTION TO HELM

the
better
way

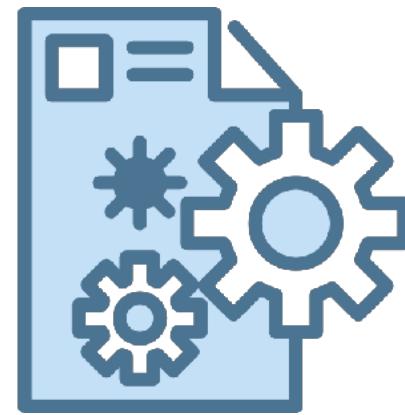


HELPERS & VALUES



- ▶ nameOverride: [e.g. api | pinger | poller]

HELPERS & VALUES

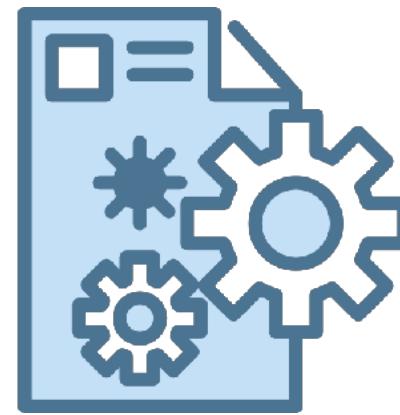


- ▶ nameOverride: [api | pinger | poller]

_helpers.yaml

```
{%- define "chart.name" -%}
{%- default .Chart.Name .Values.nameOverride | trunc 63 | trimSuffix "-" -%}
{%- end -%}
```

HELPERS & VALUES



- ▶ nameOverride: [api | pinger | poller]

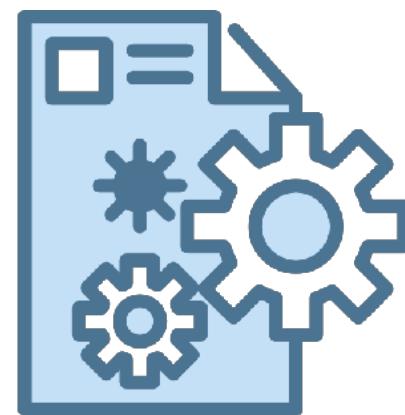
```
{%- define "chart.name" -%}
{{- default .Chart.Name .Values.nameOverride | trunc 63 | trimSuffix "-" -}}
{%- end -%}
```

_helpers.yaml



| Make the chart reusable

HELPERS & VALUES



- ▶ nameOverride: [api | pinger | poller]
- ▶ service.enabled: true

```
{%- define "chart.name" -%}
{{- default .Chart.Name .Values.nameOverride | trunc 63 | trimSuffix "-" -}}
{%- end -%}
```

_helpers.yaml

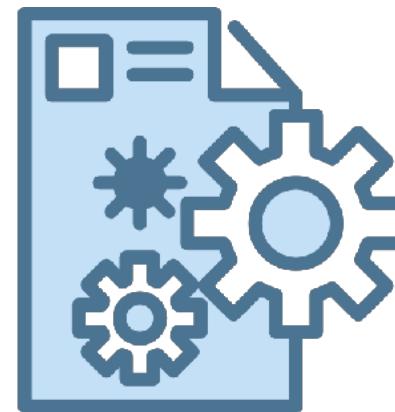
```
{%- if .Values.service.enabled %}
{%- end %}
```

service.yaml



| Make the chart reusable

HELPERS & VALUES



- ▶ nameOverride: [api | pinger | poller]
- ▶ service.enabled: true

```
{%- define "chart.name" -%}
{{- default .Chart.Name .Values.nameOverride | trunc 63 | trimSuffix "-" -}}
{%- end -%}
```

_helpers.yaml



| Make the chart reusable

```
{%- if .Values.service.enabled -%}
{%- end -%}
```

service.yaml



| Make entire RD optional

TEMPLATES



This works ;)
but ...

api-deployment.yaml

```
apiVersion: apps/v1beta2
kind: Deployment
metadata:
  name: {{ .Values.api.name }}
  labels:
    app: {{ .Values.api.name }}
    chart: {{ template "msa.chart" . }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  replicas: {{ .Values.api.replicaCount }}
  selector:
    matchLabels:
      app: {{ .Values.api.name }}
      release: {{ .Release.Name }}
  template:
    metadata:
      labels:
        app: {{ .Values.api.name }}
        release: {{ .Release.Name }}
```



This is D.R.Y !

deployment.yaml

```
apiVersion: apps/v1beta2
kind: Deployment
metadata:
  name: {{ .Values.name }}
  labels:
    app: {{ .Values.name }}
    chart: {{ template "msa.chart" . }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  replicas: {{ .Values.replicaCount }}
  selector:
    matchLabels:
      app: {{ .Values.name }}
      release: {{ .Release.Name }}
  template:
    metadata:
      labels:
        app: {{ .Values.name }}
        release: {{ .Release.Name }}
```

TEMPLATES



This works ;)
but ...

api-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: {{ .Values.api.name }}
  labels:
    app: {{ .Values.api.name }}
    chart: {{ template "msa.chart" . }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  type: {{ .Values.api.service.type }}
  ports:
    - port: {{ .Values.api.service.port }}
      targetPort: http
      protocol: TCP
      name: http
  selector:
    app: {{ .Values.api.name }}
    release: {{ .Release.Name }}
```

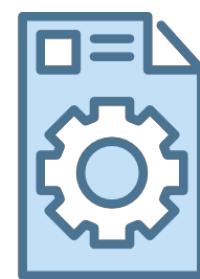


This is D.R.Y !

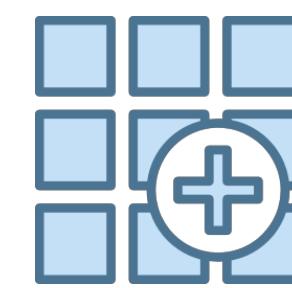
service.yaml

```
{{- if .Values.service.enabled -}}
apiVersion: v1
kind: Service
metadata:
  name: {{ .Values.api.name }}
  labels:
    app: {{ .Values.api.name }}
    chart: {{ template "msa.chart" . }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  type: {{ .Values.api.service.type }}
  ports:
    - port: {{ .Values.api.service.port }}
      targetPort: http
      protocol: TCP
      name: http
...
{{- end -}}
```

“MSA” CHART



```
--- Chart.yaml  
--- charts  
--- templates  
    --- NOTES.txt  
    --- _helpers.tpl  
    --- deployment.yaml  
    --- service.yaml  
--- values.yaml
```

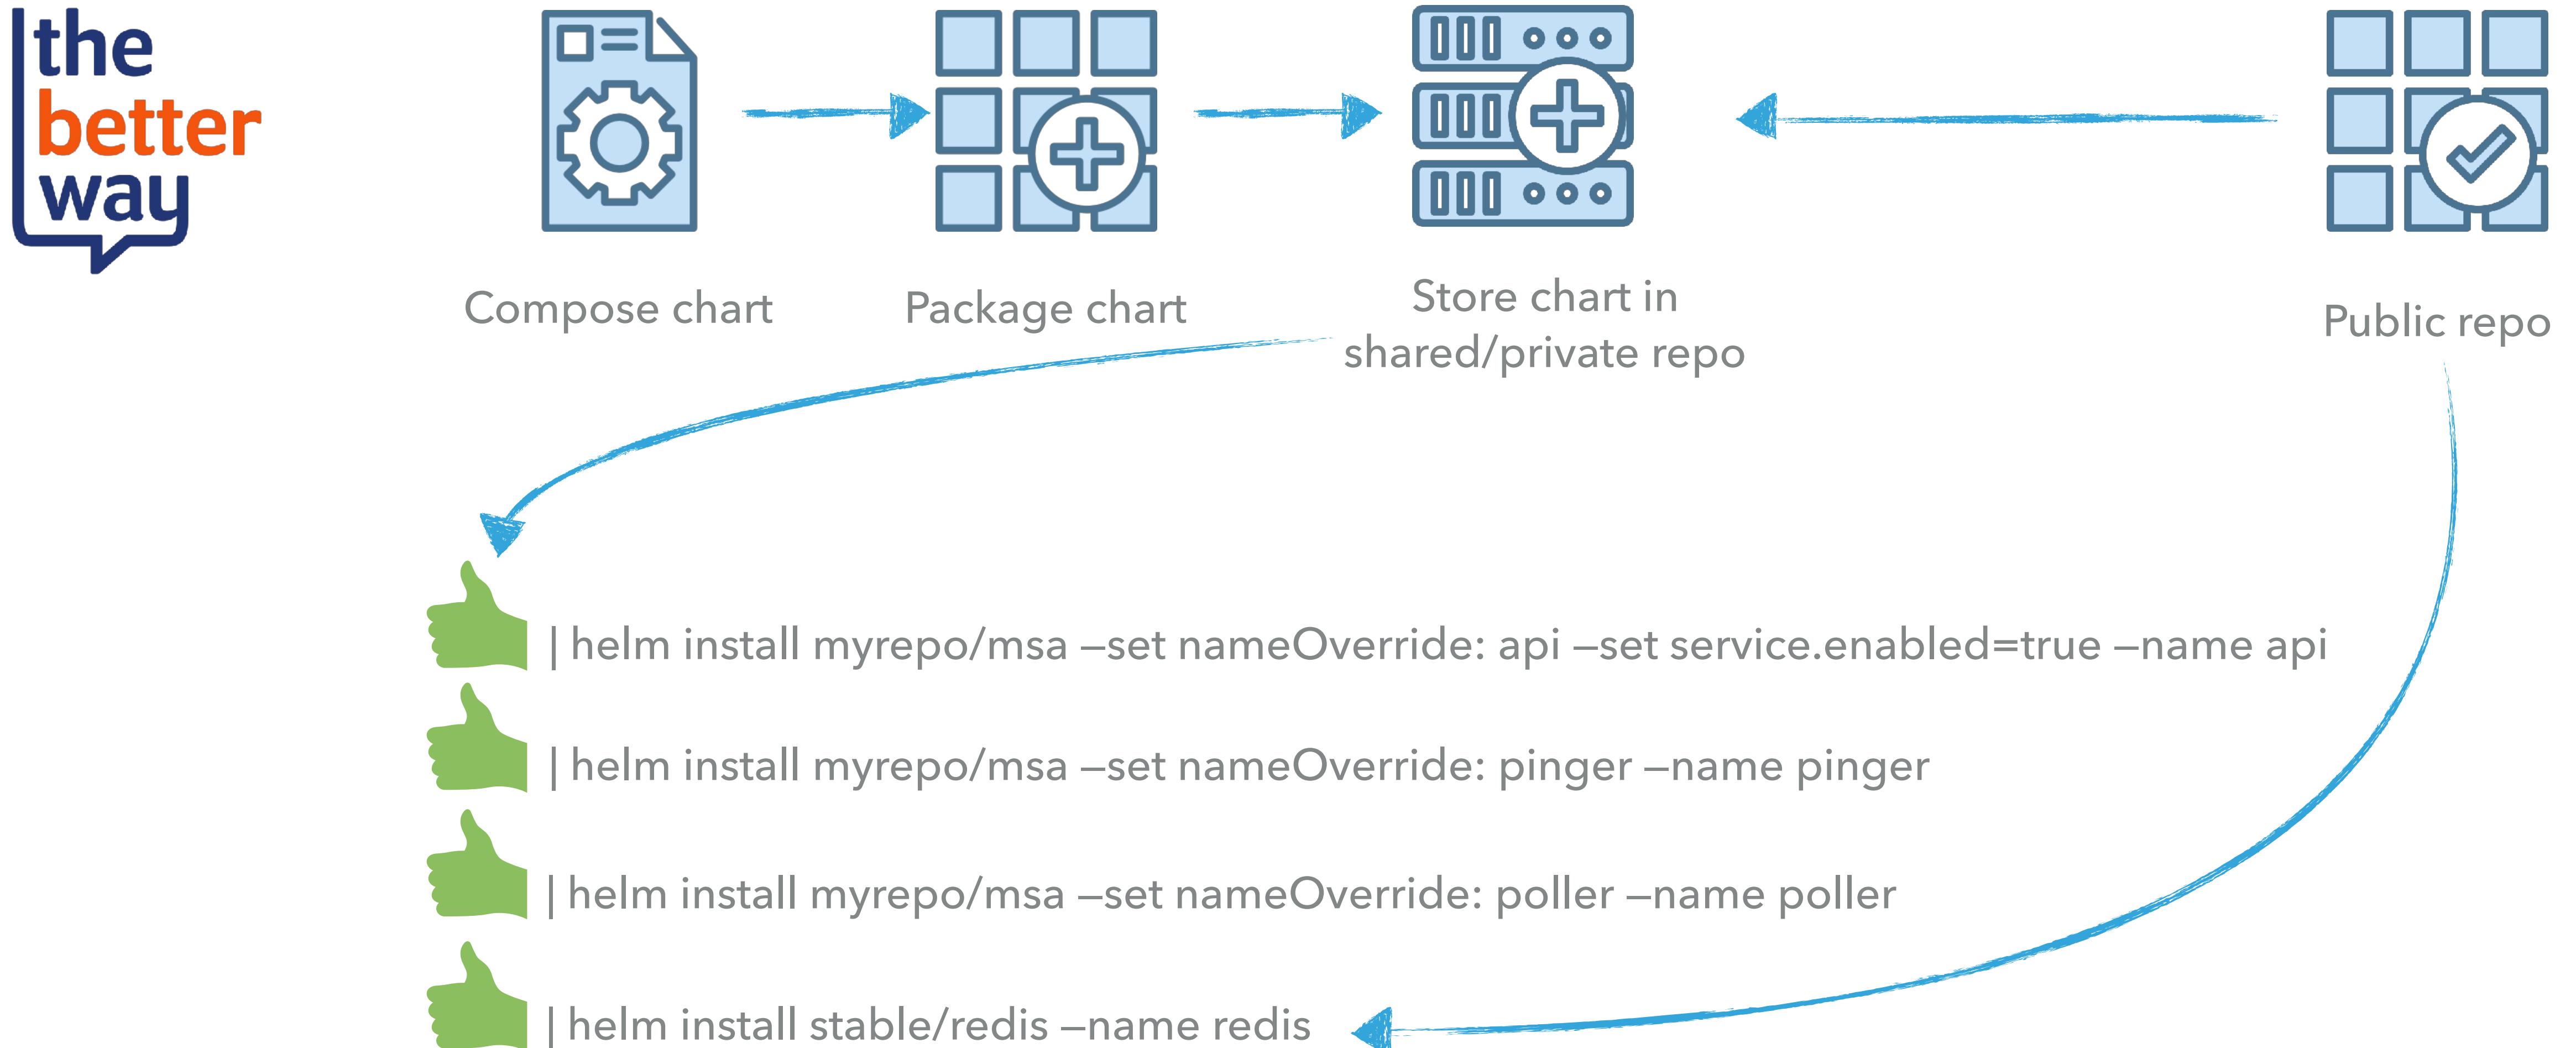


my-chart-repo/msa-0.1.0.gz

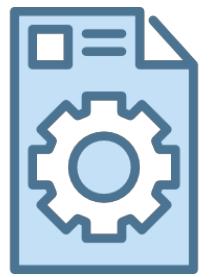
```
apiVersion: v1  
appVersion: "1.0"  
description: msg chart for Kubernetes  
name: api  
version: 0.1.0
```



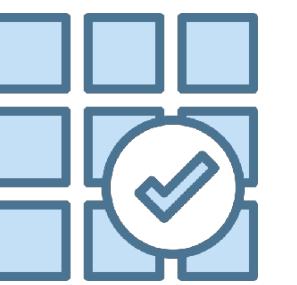
INTRODUCTION TO HELM



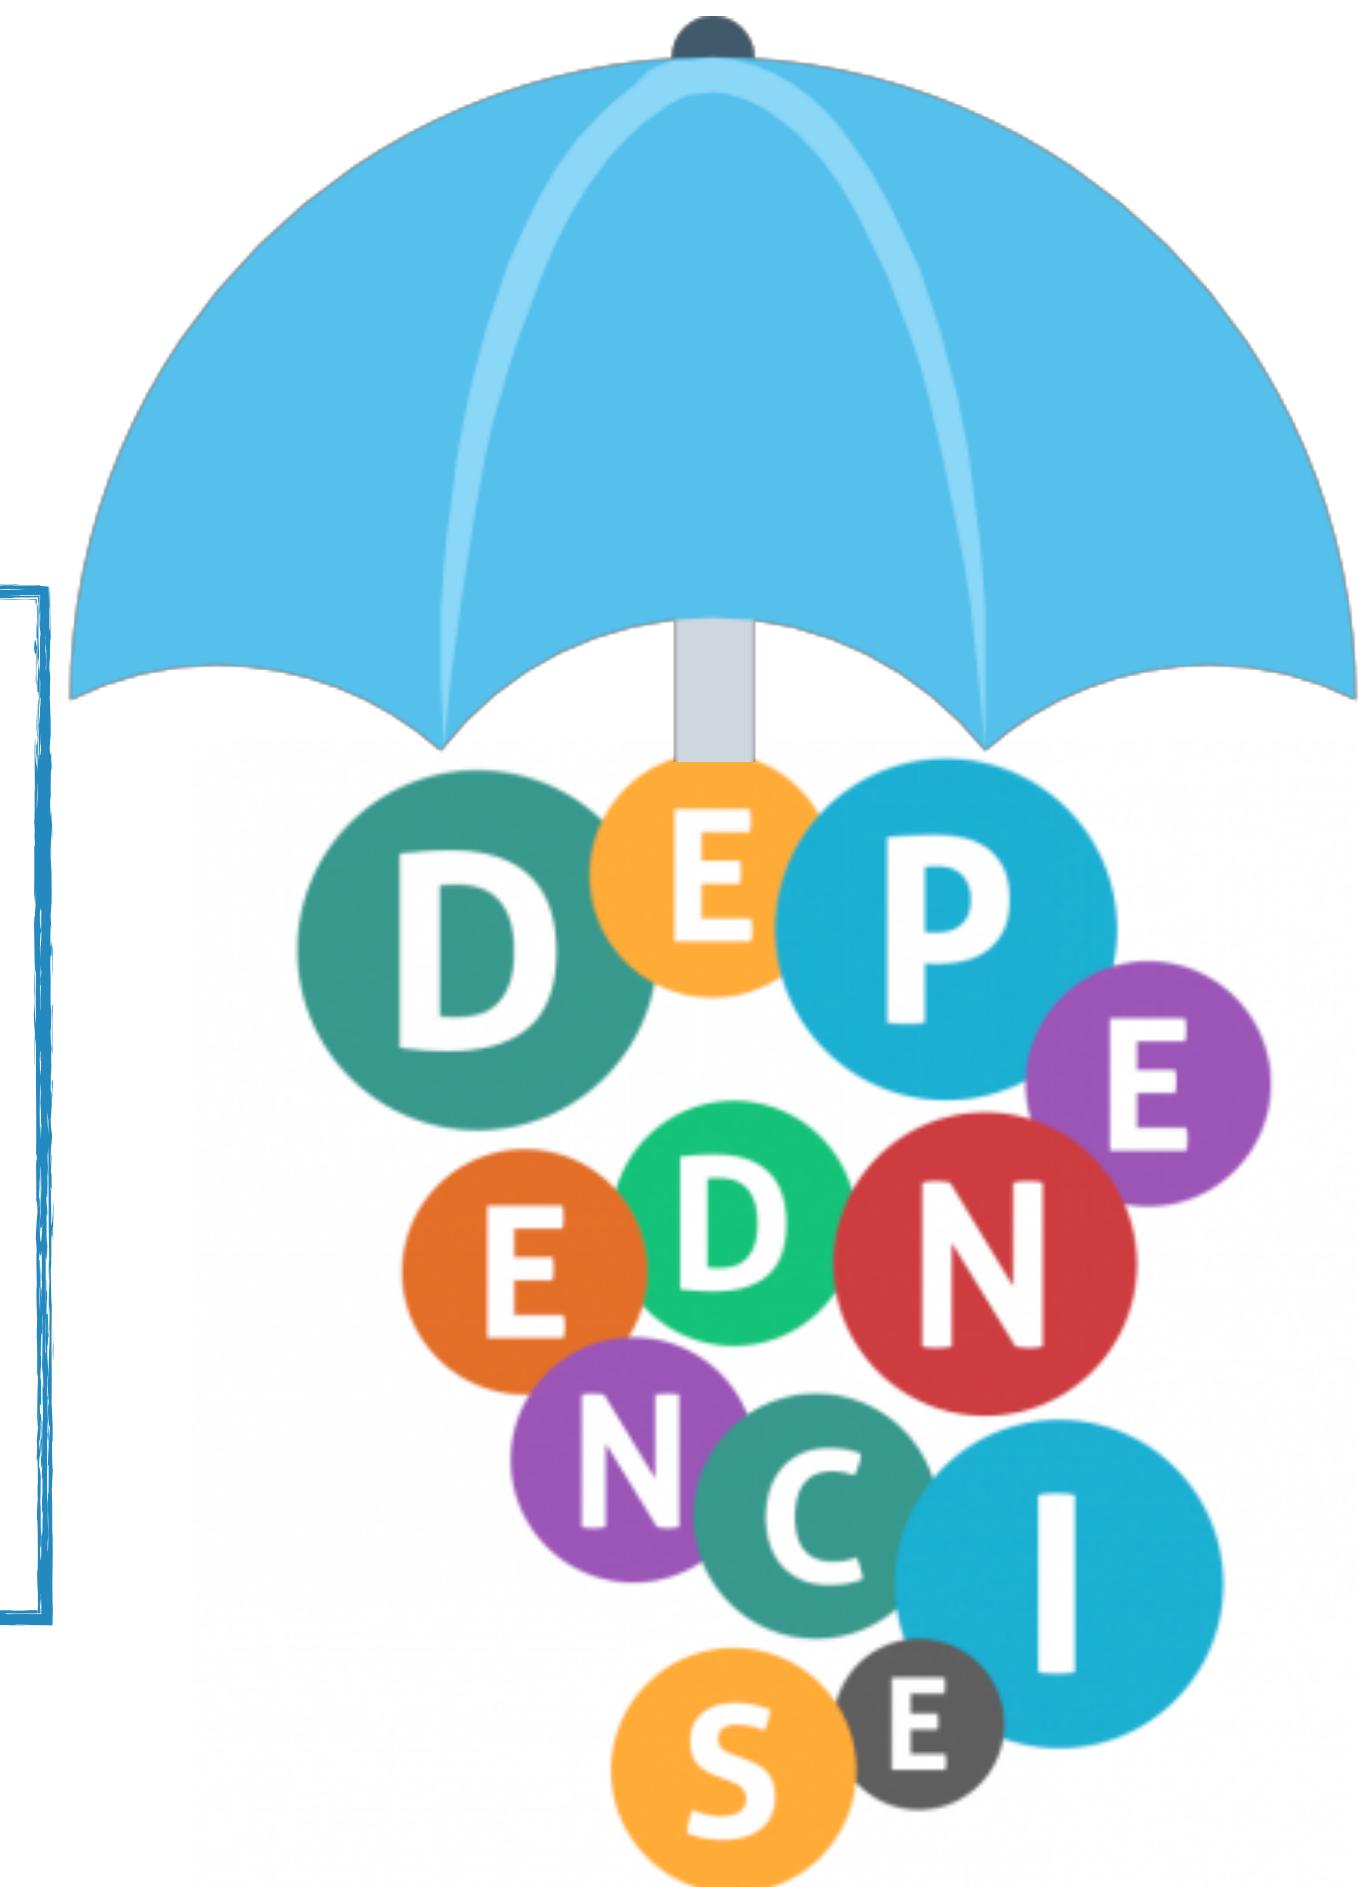
MANAGING DEPENDENCIES - A.K.A UMBRELLA CHART



```
.  
|-- Chart.yaml  
|-- charts  
|-- templates  
|   |-- NOTES.txt  
|   |-- _helpers.tpl  
|   |-- dependencies.yaml  
|-- values.yaml
```



```
dependencies:  
  - name: msa  
    repository: file://../msa  
    version: 0.1.0  
  - name: redis  
    repository: stable/redis  
    version: x.x.x
```





HANDS ON

HOW TO USE THIS LAB ?!



Informational -> you don't have / need to do it ;)



Get your hands dirty -> follow the instructions

MINIKUBE



MINIKUBE

- ▶ Minikube supports Kubernetes features such as:
 - ▶ DNS
 - ▶ NodePorts
 - ▶ ConfigMaps and Secrets
 - ▶ Dashboards
 - ▶ Container Runtime: Docker, rkt and CRI-O
 - ▶ Enabling CNI (Container Network Interface)
 - ▶ Ingress



WHY

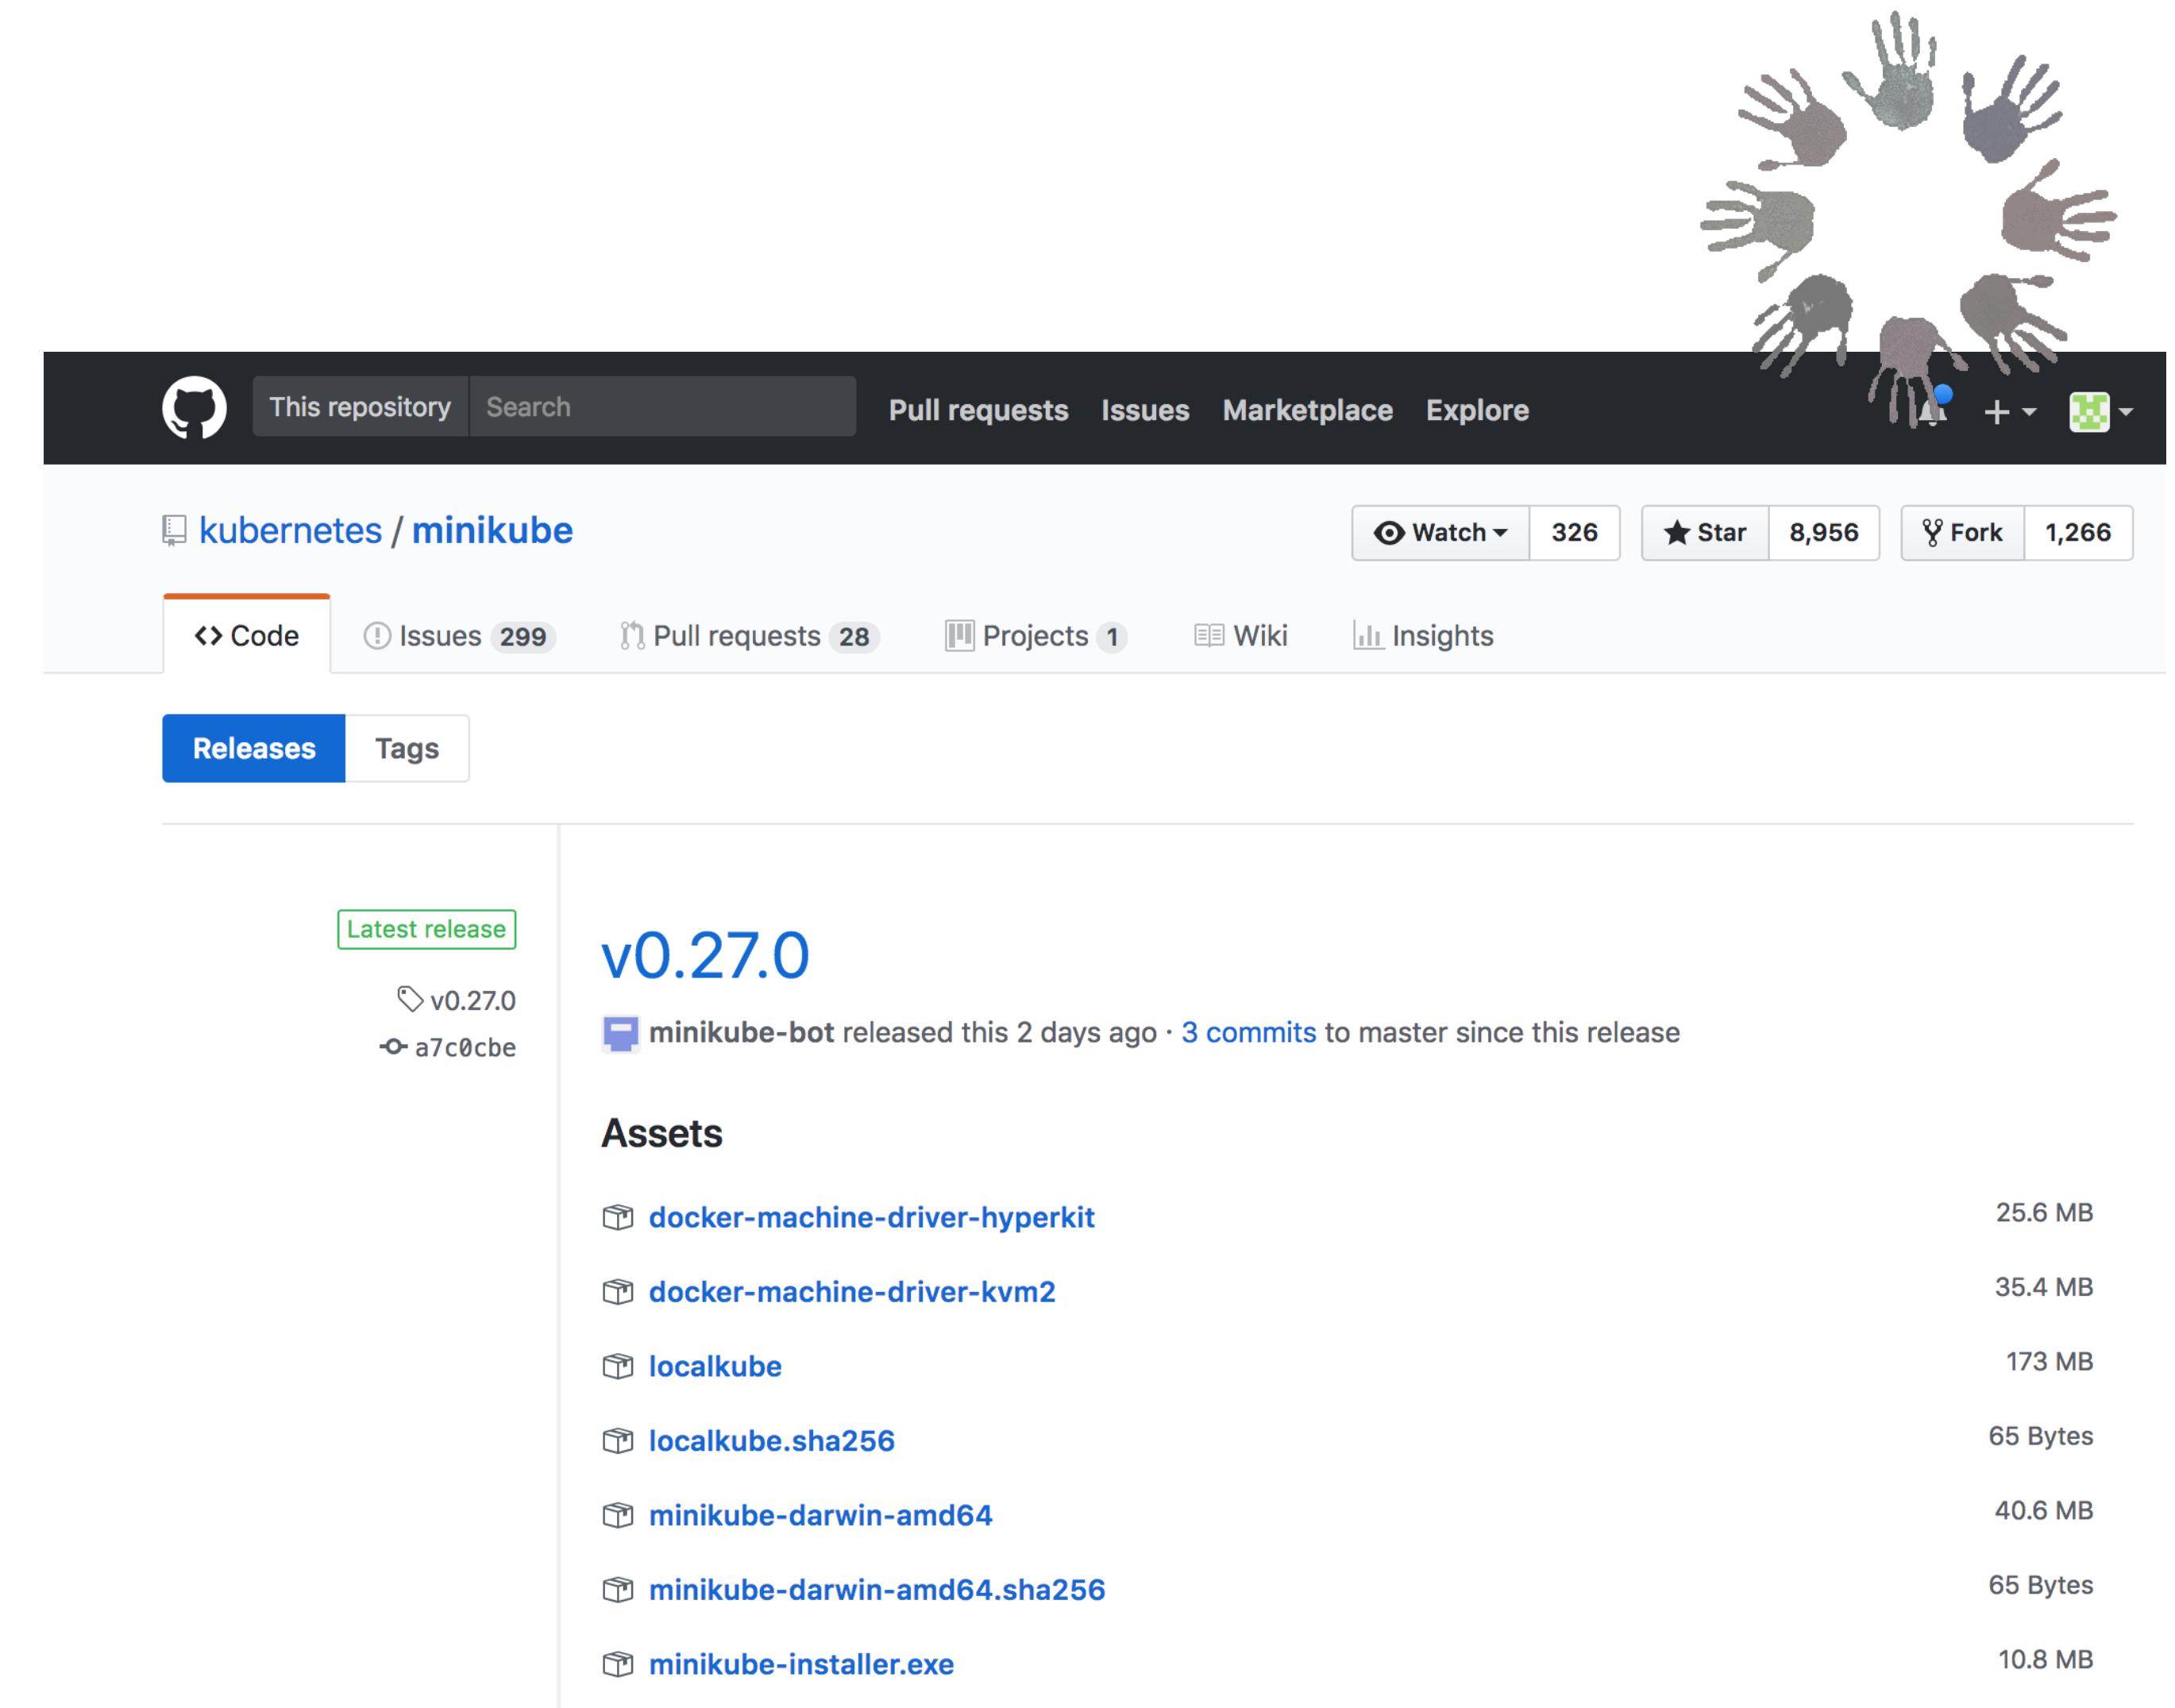
- ▶ Good for POC, QuickStarts and Demo's
- ▶ Gives a taste of Kubernetes and how to work with it, without going through a complex setup



INSTALLING MINIKUBE

▶ Install *kubectl* [link](#)

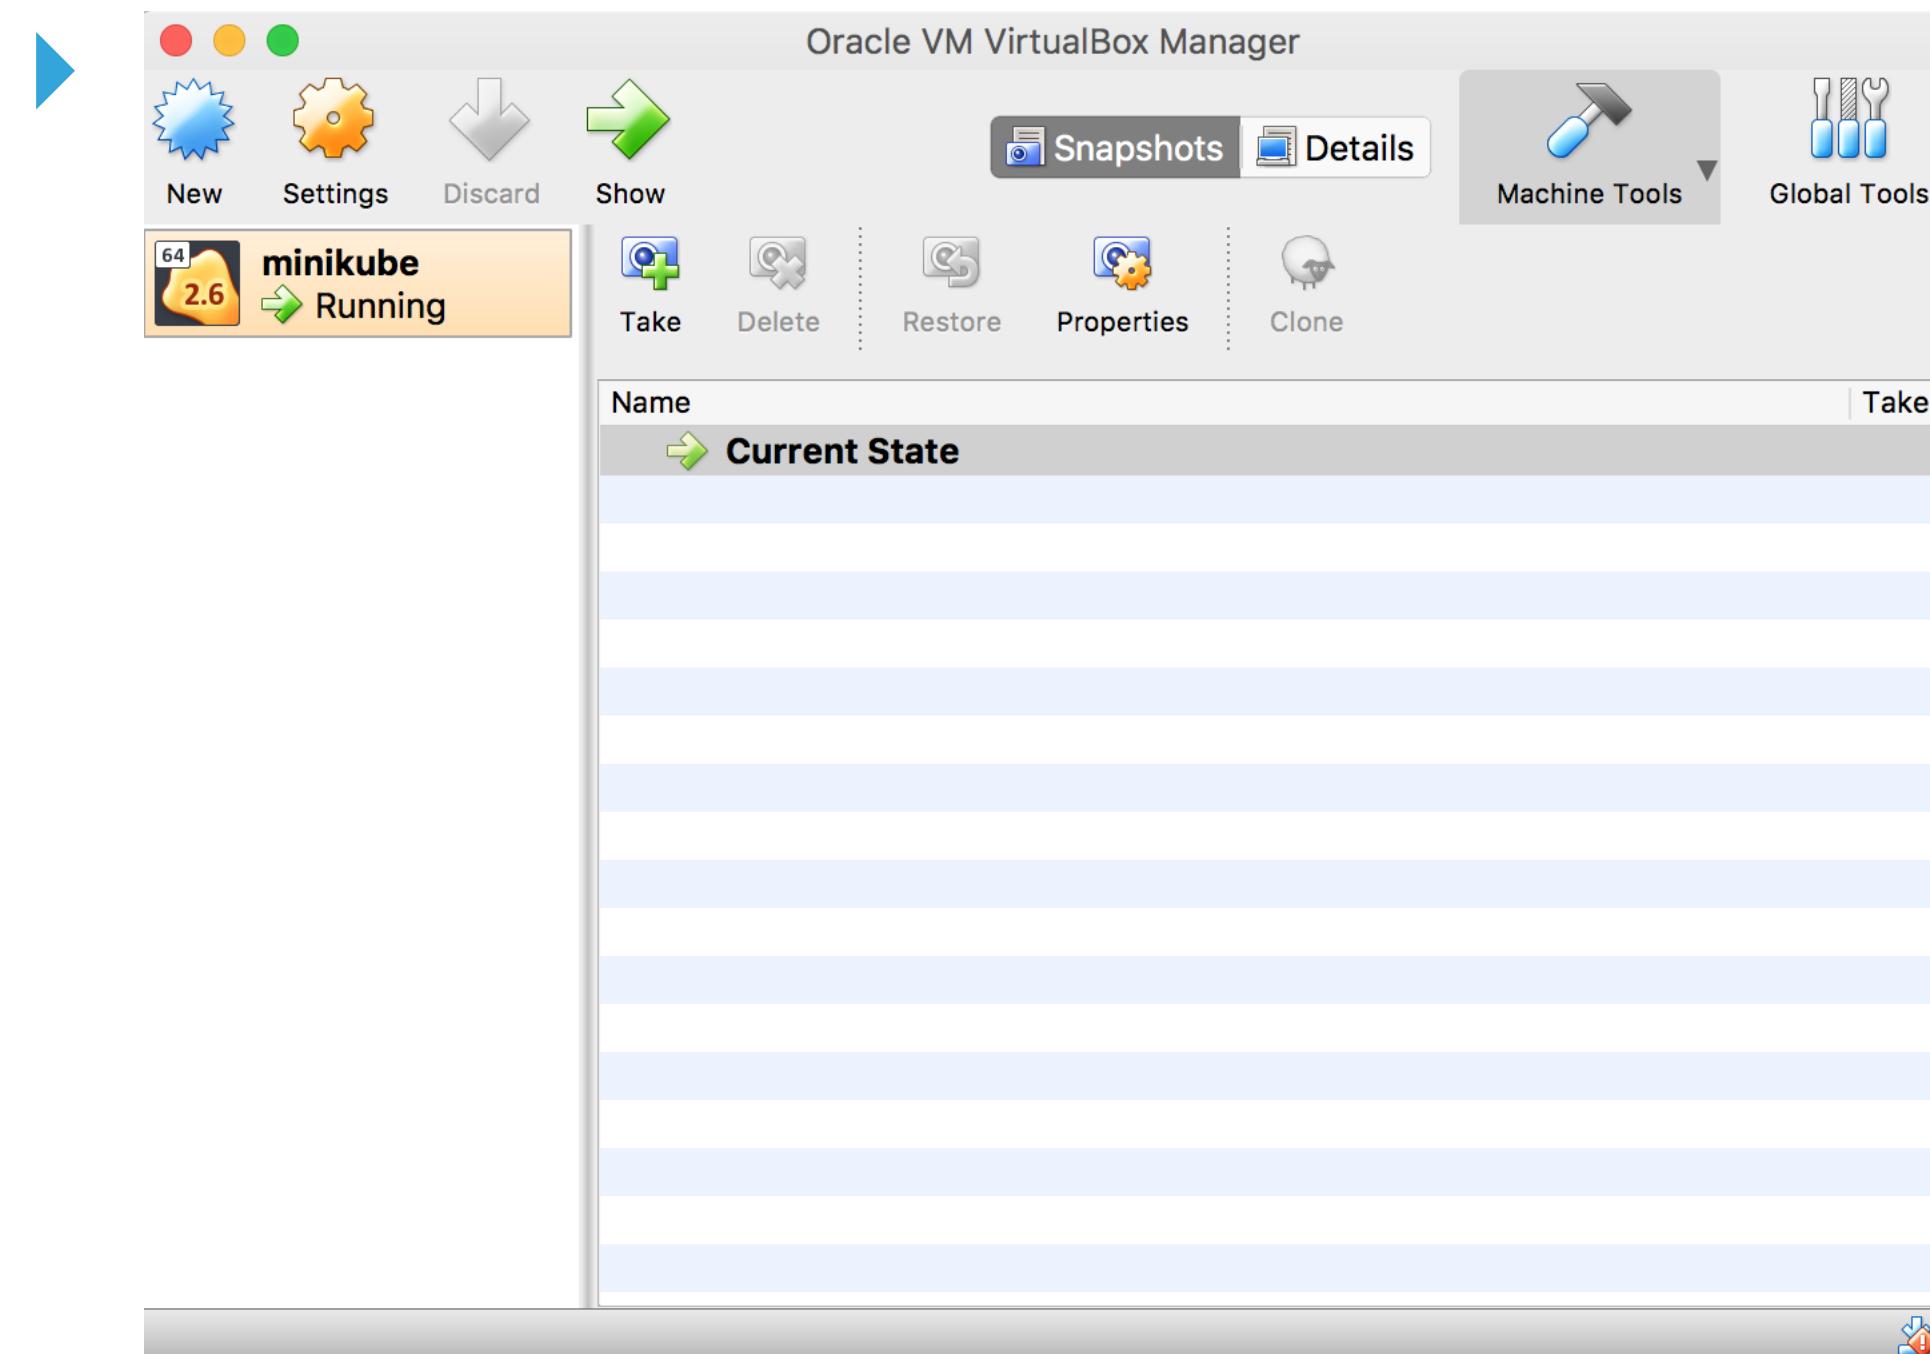
▶ Install *minikube* [link](#)



The image shows a screenshot of a GitHub repository page for the 'kubernetes/minikube' project. The repository has 8,956 stars and 1,266 forks. The 'Releases' tab is selected, showing the 'v0.27.0' release which was released 2 days ago. This release includes 3 commits to the master branch. The release page lists several assets: 'docker-machine-driver-hyperkit' (25.6 MB), 'docker-machine-driver-kvm2' (35.4 MB), 'localkube' (173 MB), 'localkube.sha256' (65 Bytes), 'minikube-darwin-amd64' (40.6 MB), 'minikube-darwin-amd64.sha256' (65 Bytes), and 'minikube-installer.exe' (10.8 MB).

MINIKUBE START

► ***minikube start***



20:31 \$ **minikube start**
Starting local Kubernetes v1.9.0 cluster...
Starting VM...
Getting VM IP address...
Moving files into cluster...
Setting up certs...
Connecting to cluster...
Setting up kubeconfig...
Starting cluster components...
Kubectl is now configured to use the cluster.
Loading cached images from config file.



INTRODUCTION TO HELM

MACOSX VIA HOMEBREW

```
brew update  
brew install kubernetes-cli kubernetes-helm  
brew cask install minikube
```



Screenshot of a GitHub repository page for `kubernetes/minikube`. The page shows the repository's overview with metrics like 326 Watchers, 8,956 Stars, and 1,266 Forks. The `Releases` tab is selected, showing the latest release is `v0.27.0`, released by `minikube-bot` 2 days ago with 3 commits. The `Assets` section lists several files: `docker-machine-driver-hyperkit` (25.6 MB), `docker-machine-driver-kvm2` (35.4 MB), `localkube` (173 MB), `localkube.sha256` (65 Bytes), `minikube-darwin-amd64` (40.6 MB), `minikube-darwin-amd64.sha256` (65 Bytes), and `minikube-installer.exe` (10.8 MB).

INTRODUCTION TO HELM

UBUNTU/DEBIAN



```
sudo apt-get update && sudo apt-get install -y apt-transport-https  
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -  
echo "deb http://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee -a /etc/apt/  
sources.list.d/kubernetes.list
```

```
sudo apt-get update
```

```
sudo apt-get install -y kubectl
```

```
curl https://raw.githubusercontent.com/kubernetes/helm/master/scripts/get | bash  
curl -Lo minikube https://storage.googleapis.com/minikube/releases/v0.30.0/minikube-linux-amd64 && chmod +x minikube && sudo cp minikube /usr/local/bin/ && rm minikube
```

INTRODUCTION TO HELM

CENTOS/RHEL/FEDORA

```
cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86\_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF
```

```
yum install -y kubectl
```

```
curl https://raw.githubusercontent.com/kubernetes/helm/master/scripts/get | bash
curl -Lo minikube https://storage.googleapis.com/minikube/releases/v0.30.0/minikube-linux-amd64 && chmod +x minikube && sudo cp minikube /usr/local/bin/ && rm minikube
```



WINDOWS

Use the [Chocolaty](#) package manager:

```
choco install kubernetes-cli kubernetes-helm
```

Download Minikube from: <https://github.com/kubernetes/minikube/releases>

Download the minikube-windows-amd64.exe file,
rename it to minikube.exe and add it to your path.





INSTALL HELM ON KUBERNETES

```
helm init --service-account helm --dry-run --debug > helm-init.yml
kubectl create -f helm-rbac-cluster-wide.yml
kubectl create -f helm-init.yml
```



```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: helm
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  name: helm
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- kind: ServiceAccount
  name: helm
  namespace: kube-system
```

INTRODUCTION TO HELM

CONFIGURE HELM REPOSITORIES

helm repo add stable https://kubernetes-charts.storage.googleapis.com



Screenshot of the GitHub repository for `helm / charts`. The repository has 243 issues, 5,186 stars, and 4,607 forks. The `charts / stable` branch is selected. The commit history shows recent activity from `okgolove` and `k8s-ci-robot`, including changes to the `NOTES` file and various chart updates. The commits are listed below:

Commit	Message	Time Ago
acs-engine-autoscaler	Enrich deploy. template for acs-engine-autoscaler (#5662)	5 months ago
aerospike	[stable/aerospike] Add cmd and args options to Aerospike config (#3856)	8 months ago
anchore-engine	add brady todhunter as review/maintainer to anchore-engine (#8514)	5 days ago
apm-server	[stable/apm-server] Elastic APM Server (#6058)	4 months ago
ark	[stable/ark] Quote configuration parameter backupStorageProvider.conf...	8 days ago
artifactory-ha	Deprecate JFrog charts (moved to https://github.com/jfrog/charts) (#7627)	a month ago
artifactory	Deprecate JFrog charts (moved to https://github.com/jfrog/charts) (#7627)	a month ago
auditbeat	upgrade auditbeat (#8277)	12 days ago
aws-cluster-autoscaler	Typo fix in aws-cluster-autoscaler/README.md (#4297)	7 months ago

INSTALL HELM-S3 PLUGIN

helm plugin install https://github.com/hypnoglow/helm-s3.git



- > If you are planning on pushing to your s3 based helm repo
- > get your own repo and setup the required permissions.
- > In our case you will be using it for read / download purposes which is already public / open source.

ADD THE MSA-CHARTS S3 HELM REPOSITORY

helm repo add msa-charts s3://msa-charts/



Type a prefix and press Enter to search. Press ESC to clear.

Upload Create folder Actions EU (Ireland) ▾

Name	Last modified	Size	Storage class
<input type="checkbox"/> index.yaml	Oct 18, 2018 11:00:19 AM GMT+0300	4.0 KB	Standard
<input type="checkbox"/> msa-api-0.1.0.tgz	Oct 14, 2018 10:38:10 AM GMT+0300	2.6 KB	Standard
<input type="checkbox"/> msa-api-0.1.10.tgz	Oct 17, 2018 11:52:21 AM GMT+0300	2.3 KB	Standard
<input type="checkbox"/> msa-api-0.1.13.tgz	Oct 17, 2018 3:43:16 PM GMT+0300	2.3 KB	Standard
<input type="checkbox"/> msa-api-0.1.14.tgz	Oct 17, 2018 3:49:12 PM	2.3 KB	Standard

s3 plugin updates
the index.yaml on
each version
upload

DEPLOY YOUR UMBRELLA ...

```
kubectl create namespace msa-umbrella  
helm dep build ./helm/msa-umbrella  
helm install --name v1 --namespace msa-umbrella ./helm/msa-umbrella
```



DEPLOY YOUR UMBRELLA ...

```
kubectl create namespace msa-umbrella  
helm dep build ./helm/msa-umbrella  
helm install --name v1 --namespace msa-umbrella ./helm/msa-umbrella
```



CHECK YOUR UMBRELLA ...

helm list ***# view helm releases which are installed***

kubectl -n msa-umbrella get po,svc,deploy ***# view msa-umbrella pods, services & deployments***



RECAP

- ▶ Helm is awesome -> (completes Kubernetes)
- ▶ Application centric solution
- ▶ Flexible templating language
- ▶ Well integrated into workflows such as GitKube, Scafold and others

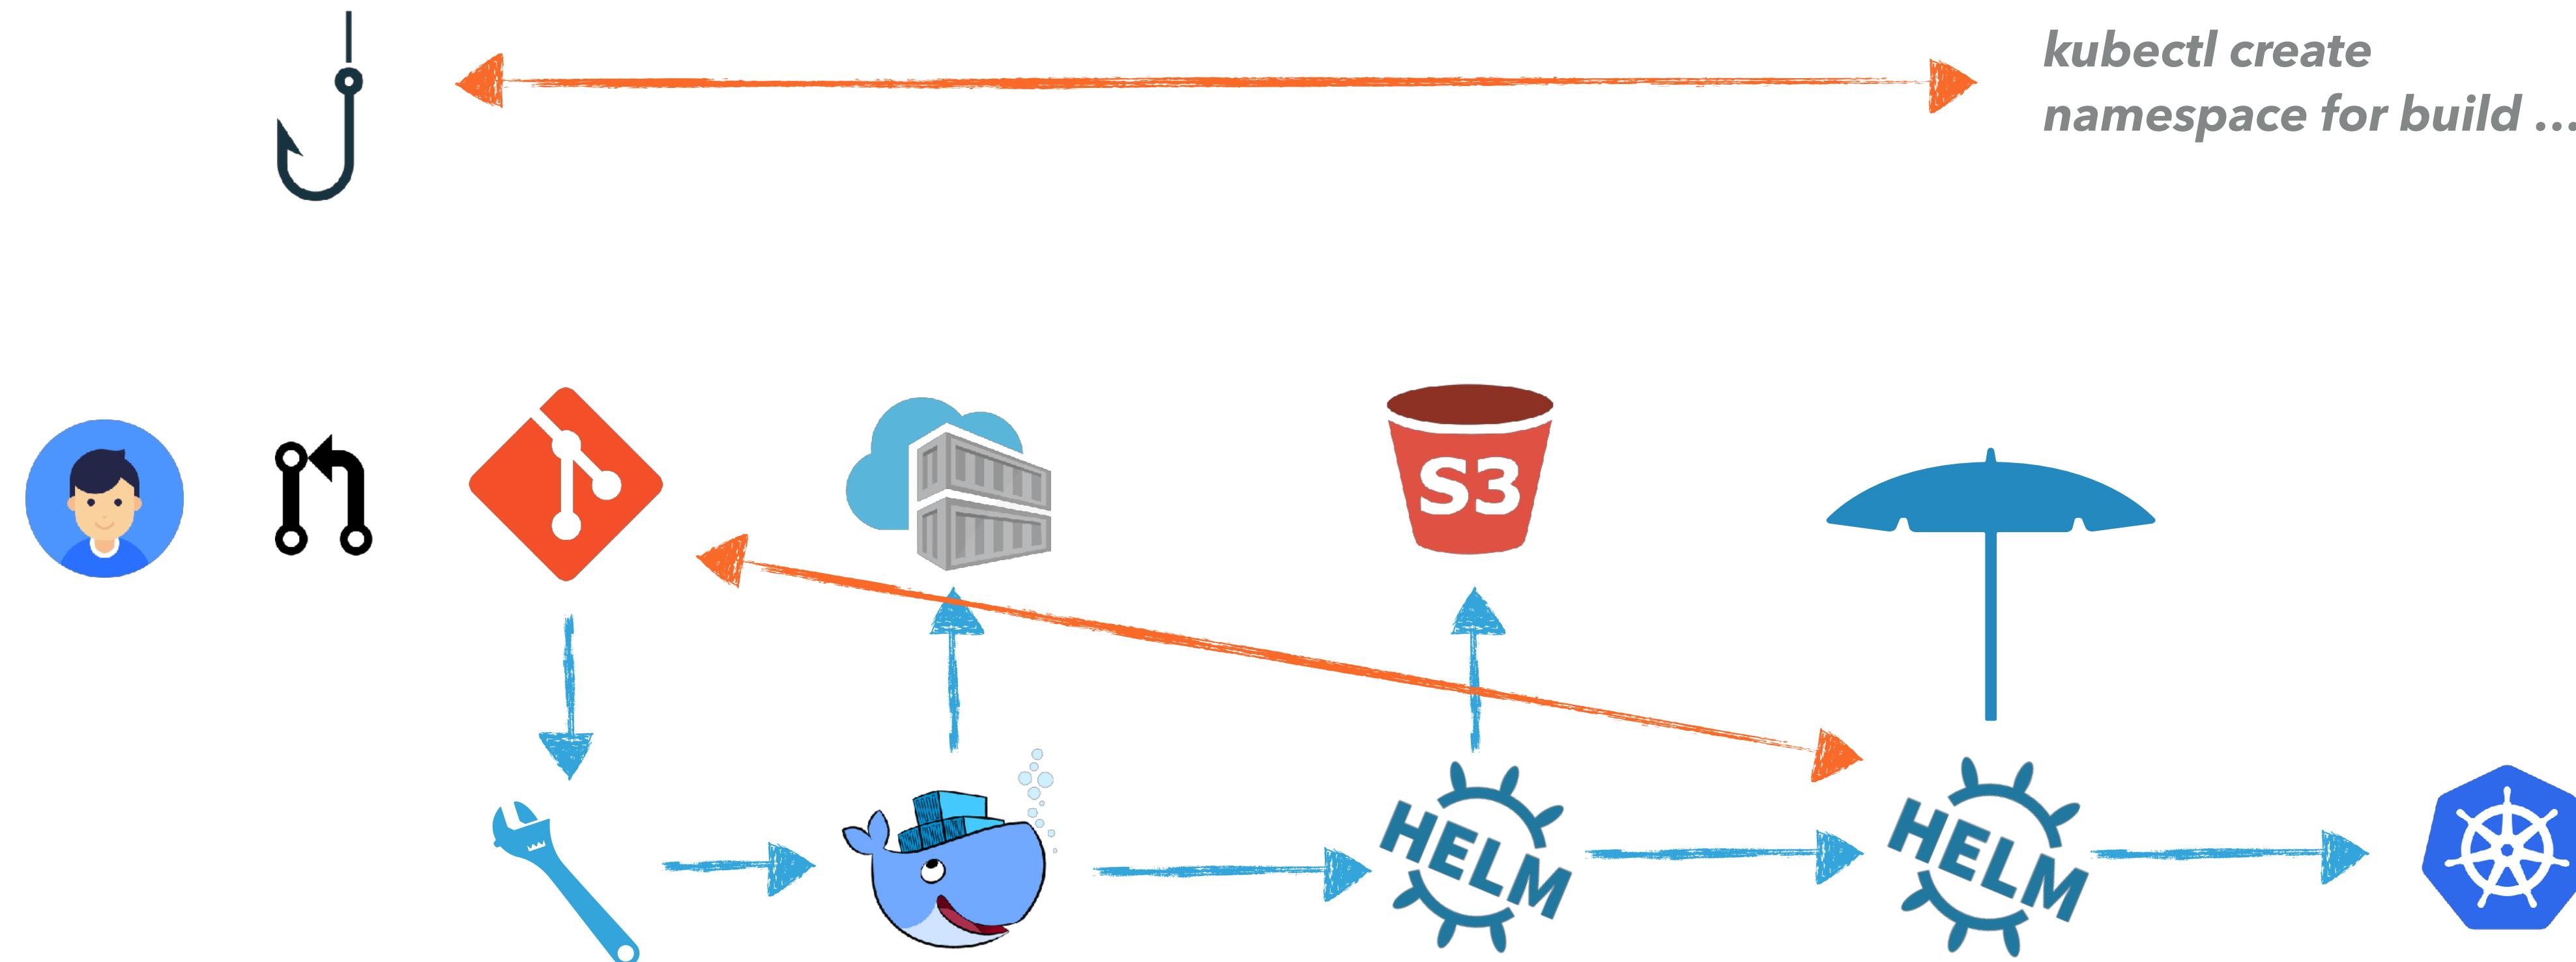
USE [DON'T OVER USE] YOUR IMAGINATION

- ▶ Custom Resource Definitions { e.g istio }
- ▶ Common monitoring RD's { e.g prometheus }  | {{ if .Values.monitoring.enabled }}
- ▶ Common logging RD's { e.g fluentd }  | {{ if .Values.logging.enabled }}
- ▶ Ingresses (yes / no)
- ▶ Services (type: NodePort, ClusterIp, LoadBalancer)  | {{ if .Values.service.nodelp }}
- ▶ Provider specific configs { on-prem vs. cloud }
- ▶



So Whats Next ? ...

CONTINUES INTEGRATION PROCESS [JENKINS || GITLAB || “FOOBAR”]



INTRODUCTION TO HELM

JENKINS PIPELINE

← 2 search ⚡ ADMIN | LOG OUT

Jenkins > msa-api > ENABLE AUTO REFRESH

Pipeline msa-api

i Status + add description
refresh Changes DISABLE PROJECT
play Build with Parameters
trash Delete Pipeline
gear Configure
magnifying-glass Full Stage View
globe Open Blue Ocean
refresh Rename
shield Embeddable Build Status
chat Build Review
question-mark Pipeline Syntax

↻ Recent Changes

Stage View

Average stage times:
(Average full run time: ~1min 12s)

Setup	Checkout	Build	Prepare Test Env	System Tests	Deploy	Deployment env. setup	HELM package	HELM dependency update	HELM deploy
22ms	8s	12s	16ms	22ms	67ms	49s	8s	5s	2min 55s
Oct 18 10:59	2 commits	21ms	7s	26s	17ms	79ms	97ms	18s	10s
Oct 17 15:48	2 commits	17ms	7s	4s	14ms	16ms	57ms	13s	8s

#15 Oct 18, 2018 7:59 AM #14 Oct 17, 2018 12:48 PM

find trend

Build History

#	Date	Commits
#15	Oct 18, 2018 7:59 AM	2 commits
#14	Oct 17, 2018 12:48 PM	2 commits

PIPELINE LIBRARY - NOW SUPPORTS HELM LIFECYCLE



Jump to... / Pull requests Issues Marketplace Explore

tikalk / tikal-pipelib Unwatch 15 Star 1 Fork 0

Code Issues 0 Pull requests 0 Projects 1 Wiki Insights Settings

Tikal Jenkins Pipeline Library Edit

Manage topics

303 commits 3 branches 0 releases 4 contributors

Branch: master New pull request Create new file Upload files Find file Clone or download

Tom Gur fixed helm upgrade command Latest commit 712997e 6 days ago

src fixed helm upgrade command 6 days ago

vars AWS configure 16 days ago

.gitignore Initiate tikal-pipelib a year ago

README.md rename and replace tpl_ with tpl prefix a year ago

This screenshot shows a GitHub repository page for 'tikalk / tikal-pipelib'. The page title is 'Tikal Jenkins Pipeline Library'. The repository has 303 commits, 3 branches, 0 releases, and 4 contributors. The 'Code' tab is selected. At the bottom, there's a list of recent commits:

- Tom Gur fixed helm upgrade command (Latest commit 712997e, 6 days ago)
- src fixed helm upgrade command (6 days ago)
- vars AWS configure (16 days ago)
- .gitignore Initiate tikal-pipelib (a year ago)
- README.md rename and replace tpl_ with tpl prefix (a year ago)

CREDITS

- ▶ Mark Kirshner & Anatoly Rabkin
- ▶ Helm documentation -> <https://docs.helm.sh/>
- ▶ ChaosKube helm chart -> <https://github.com/helm/charts/tree/master/stable/chaoskube>
- ▶ Icons -> <https://github.com/octo-technology/kubernetes-icons>
- ▶ Kubernetes docs -> <https://kubernetes.io/docs/home/?path=users>



Beautiful